

Handoff to Chapter 2. Chapter 2 receives: 4 problems, 5 objectives, 5 research questions, and the RGAIG concept. It must ground these in the literature and derive the 6 gaps that justify the methodology.

Limitations of the Study

This section articulates the boundary conditions of the research. Table 1 identifies six dimensions along which the study’s scope is explicitly bounded, together with the justification for each boundary. These limitations are revisited in Chapter 5 with corresponding mitigation strategies and future research directions.

These boundaries are not weaknesses to be apologised for but deliberate scope decisions that ensure the research remains tractable, reproducible, and falsifiable within the doctoral timeframe. The framework architecture is designed to accommodate broader geographic, temporal, and domain coverage in post-doctoral extensions.

Chapter Summary

This chapter established the dissertation’s problem context, identified six research gaps from a 156-study evidence base, and translated those gaps into the study’s objectives, research questions, hypotheses, and intended contributions. It also defined the scope and boundary conditions of the research so that the claims taken forward into the later chapters remain aligned with the available evidence.

The key elements established here are the problem statement, six literature-grounded gaps, the objective and question set, the five operational hypotheses, the bounded ASD scope, and the intended contributions. Those elements are now carried forward directly into Chapters 2–5 without an additional recap inventory table.

Link to Chapter 2. Chapter 2 examines the theoretical and empirical foundations that support this argument, systematically documenting the research gaps through a literature review spanning ASD/EEG diagnostics, agentic AI, and RAG-based explainability.

This thesis argues that domain fragmentation in biomedical AI diagnostics is an unnecessary and costly design failure—not a technical inevitability—and that a unified, agentic, and explainable framework can replace it with a platform that

is more accurate, more transparent, and less expensive than the fragmented status quo.

Appendix A

Chapter 1 — Introduction Supplementary Material

Appendix: Chapter 1 — Introduction

Supplementary Material

This appendix provides supplementary detail for Chapter 1. Each item traces to a specific chapter objective or problem statement.

Limitations of the Introductory Framing

Table 1.34 documents eight limitations of the Chapter 1 framing, each mapped to a mitigation strategy. These constraints are revisited in Chapter 5 where empirical evidence either confirms or narrows them.

Table 1.34. Limitations of the Introductory Framing. Each limitation traces to a specific scope boundary established in Section 1.8 and is revisited in Chapter 5.

#	Limitation	Scope	Mitigation
1	Domain scope	ASD only; excludes cardiovascular, dermatology, genomics	Selected for EEG suitability and data availability
2	Literature window	2014–2025; pre-2014 cited selectively	RAG pipeline partially addresses recency
3	No prospective data	Retrospective public datasets only	Prospective validation prioritised for future work
4	Cost benchmarks	Published benchmarks, not hospital audits	3-scenario sensitivity analysis applied
5	Regulatory scope	FDA + EU AI Act + HIPAA only	Per-jurisdiction mapping needed for deployment
6	Prevalence estimates	WHO aggregates; regional variation not captured	Underdiagnosis acknowledged; best-available data used
7	Market projections	Industry analyst forecasts; inherent uncertainty	Regulatory environment may accelerate or constrain
8	Causal claims	Cross-sectional design limits causal inference	Mediation model tested empirically (Ch. 4)

Note. Limitations 1–3 constrain external validity; 4–7 constrain scope claims; 8 constrains causal interpretation. Each is acknowledged proactively rather than discovered by the examiner.

Chapter 1 Objective Traceability

Table 1.35 maps each Chapter 1 section to the thesis objective it serves, ensuring every section contributes to the overall argument.

Table 1.35. Chapter 1 Section-to-Objective Traceability. Each section maps to at least one thesis objective (O1–O5), ensuring no section is orphaned from the research argument.

Sec.	Section Title	Objective	How It Contributes
1.1	Background of the Study	O1, O2	Establishes ASD context and diagnostic gap
1.2	Problem Statement	O1–O5	Defines P1–P4 that motivate all objectives
1.3	Research Objectives	O1–O5	States the five main objectives
1.4	Operationalisation (A–D)	O1–O5	Maps A1–D3 sub-objectives to Parts
1.5	Research Questions	O1–O5	RQ1–RQ5 operationalise objectives
1.6	Research Structure	O1–O5	A–D framework organises evaluation
1.7	Significance	O4, O5	Academic, clinical, policy impact
1.8	Scope	O1–O3	Boundaries: ASD-only, public data
1.9	Dissertation Structure	All	Chapter-by-chapter handoff chain
1.10	Output Card	All	Deliverables to Chapter 2
1.11	Limitations	O1–O5	Known constraints stated upfront

Problem-to-Gap-to-Objective Chain

- **P1** (Workflow fragmentation) → G1 (No unified pipeline) → O2 (Technical validation) → H1 (ASD accuracy $\geq 85\%$)
- **P2** (Black-box opacity) → G3 (No operationalised explainability) → O2 (Technical validation) → H3 (SHAP features clinically valid)
- **P3** (Screening bottleneck) → G5 (AI not embedded in workflows) → O3 (Clinical adoption) → H4 (Workflow -94.6% manual)
- **P4** (Governance gap) → G4, G7 (No governance scoring) → O4 (Governance) → H5 (RAG expert $\alpha \geq 0.75$)

This chain ensures every problem identified in Chapter 1 is addressed by a specific hypothesis tested in Chapter 4 and interpreted in Chapter 5.

Chapter 2

Review of Literature

Decision framework. The literature confirms that no reviewed study produces an adopt/pilot/defer deployment verdict. This gap motivates the decision framework developed in Chapter 3 and applied in Chapter 5.

Key Sections Covered

This chapter covered four literature blocks: ASD/EEG diagnostics, agentic AI and orchestration, RAG-based explainability, and the theory base used to interpret adoption, governance, and organisational value.

Interpretation. ASD models can achieve high accuracy in isolation, but the literature still does not show a deployment-ready architecture that combines explainability, workflow automation, and governance. Governance is a missing architectural layer, not a post-hoc policy add-on.

Objective Achievement Matrix

All eight research objectives are justified by the literature reviewed in this chapter. The review establishes the need for standardised preprocessing, ASD performance thresholds, explainability support, governance architecture, validation discipline, and bounded business relevance. Methodological operationalisation follows in Chapter 3 and empirical adjudication in Chapter 4.

Link to Chapter 3. The six gaps (G1–G6) map to research objectives O1–O8 via the alignment matrix in Chapter 1 (Table 1.30). Chapter 3 develops the pragmatist, design-science methodology to operationalise these gaps into a testable research design.

This chapter argues that the existing ASD diagnostic literature is individually impressive but collectively insufficient, and that the six gaps identified are documented consequences of how the field developed rather than matters of opinion. An explainable, automated, and governance-aware framework is therefore a necessary response to the deployment demands the literature does not yet solve.

Appendix B

Chapter 2 — Literature Review Supplementary Material

Appendix: Chapter 2 — Literature Review

Supplementary Material

This appendix provides supplementary detail for Chapter 2. Each item traces to a specific research gap or theoretical foundation.

Gap-to-Problem Traceability

- **G1** (No unified ASD pipeline) $\leftarrow$ P1 (Workflow fragmentation): 138/156 reviewed studies address single domains only.
- **G2** (No clinical explainability) $\leftarrow$ P2 (Black-box opacity): 68% clinician rejection rate for opaque AI.
- **G3** (No operationalised XAI) $\leftarrow$ P2: SHAP/LIME exist but are not embedded in deployment architectures.
- **G4** (No governance scoring) $\leftarrow$ P4 (Governance gap): 6.4% of studies embed operational governance.
- **G5** (No workflow automation) $\leftarrow$ P3 (Screening bottleneck): 0/156 studies automate end-to-end pipeline.
- **G6** (No regulatory transparency) $\leftarrow$ P4: FDA SaMD, EU AI Act alignment absent from reviewed architectures.

ASD Detection Literature Trends

Three observations from the 156-study systematic review:

1. **Foundation model paradigm shift:** Pre-training on 100K+ hours of unlabelled EEG enables few-shot ASD transfer (EEG-X, 96.3%), circumventing the small-sample constraint.
2. **Graph-DL convergence:** Graph methods (DyGCN, 96.1%) match deep learning (ViT, 97.2%), confirming that functional connectivity modelling is as informative as implicit feature learning.
3. **Persistent structural gaps:** None of the top-performing studies provides cross-dataset validation *or* clinical explainability — the same G1 and G2 gaps that motivated this dissertation.

Theoretical Foundation Summary

Nine theories collectively ground the RGAIG framework:

- **Design Science Research (DSR):** Governs artifact creation (Ch. 3 methodology).
- **TAM + UTAUT:** Predict clinician adoption based on usefulness and ease of use.
- **Diffusion of Innovation:** Explains adoption dynamics across hospital settings.
- **Resource-Based View:** Justifies organisational value of unified AI infrastructure.
- **Dynamic Capabilities:** Explains how organisations adapt the framework over time.
- **Sociotechnical Systems:** Mandates human-in-the-loop design.
- **Change Management (Kotter):** Sequences implementation steps.
- **Institutional Theory:** Explains regulatory and normative pressures.
- **Trust in Automation:** Calibrates appropriate clinician reliance on AI output.

No single theory is sufficient; their combination ensures the RGAIG framework addresses technical accuracy, human trust, organisational adoption, and regulatory compliance simultaneously.

Chapter 2 Objective Traceability

- **O1** (B2C trust): Literature confirms 68% clinician rejection → trust is the primary adoption barrier.
- **O2** (Technical validation): ASD accuracy benchmarks (75–97%) establish the performance baseline.
- **O3** (Clinical adoption): TAM/UTAUT evidence confirms adoption requires usefulness + ease of use + explainability.
- **O4** (Governance): Regulatory landscape review (FDA, EU AI Act, HIPAA) identifies compliance requirements.
- **O5** (Deployment verdict): No reviewed study produces an adopt/pilot/defer verdict → this gap motivates Ch. 5.

Chapter 3

Research Methods

Table 3.84. Chapter 3: Key Sections Covered

Section	Content
Research design	Pragmatist philosophy; Design Science Research (DSR); twelve design decisions documented
Data sources	Six public benchmark datasets (305 GB); ASD domains
Preprocessing pipeline	Eight-stage pipeline: format harmonisation through z-score normalisation
Model design	Six DL architectures + five classical baselines; VotingClassifier ensemble; GAN+ResNet-18
RAG explainability	Ten-stage agentic RAG pipeline; ChromaDB 1.17 GB; Llama-3.2 local inference
Quality taxonomy	525-module unified taxonomy: 8-phase ML/DL (111), 13-phase GenAI QA (220), 10 governance pillars (194)
Validation strategy	Stratified k -fold CV; ablation; bootstrap CI (1,000 resamples); 12-expert panel
Ethics and governance	Seven-risk matrix; three-pillar governance; IRB exemption; HIPAA-safe local processing

Note. DSR = design science research; DL = deep learning; GAN = generative adversarial network; RAG = retrieval-augmented generation; CV = cross-validation; CI = confidence interval.

The methodology ensures rigor and reproducibility by combining technical evaluation with governance-aware validation. Every analytical choice—from model selection to validation thresholds—was specified before analysis to reduce post-hoc rationalisation. Chapter 4 then operationalises these methods empirically.

Chapter 4 presents the data analysis and empirical findings that operationalise the methods described here, testing all five hypotheses and reporting ASD-focused classification results, ablation studies, expert panel evaluation, and governance framework validation.

This chapter argues that the research gaps identified in Chapter 2 require a methodology that treats the framework artifact, its empirical evaluation, and its governance architecture as inseparable components—a design science strategy that no single existing paradigm provides.

Appendix C

Chapter 3 — Research Methods Supplementary Material

Appendix: Chapter 3 — Research Methods

Supplementary Material

This appendix contains supplementary tables, figures, and detailed analyses that support Chapter 3 but were moved from the main text to maintain narrative focus. Each item is cross-referenced from the main chapter.




Figure 3.22. Appendix Task Flow: Supplementary Material for Chapter 3 (Research Methods)

EEG Signal Processing Pipeline

The EEG-RAG framework [54] implements a seven-stage signal processing pipeline that transforms raw multi-channel EEG recordings into classification-ready feature vectors. Figure 3.23 illustrates this pipeline, and Table 3.85 documents the filtering parameters with their neurophysiological justification. The pipeline design reflects two constraints:

1. **Preserve clinically relevant content:** Retain oscillatory content in the 0.5–45 Hz range while removing non-neural contamination such as DC drift, muscle artifacts, and power-line interference.
2. **Produce classification-ready segments:** Generate fixed-length, normalised segments suitable for both classical feature extraction and deep learning input.




Figure 3.23. Seven-Stage EEG Signal Processing Pipeline from the EEG-RAG Framework

Note. ICA = independent component analysis; IIR = infinite impulse response; Hz = hertz; s = seconds. The pipeline produces fixed-length, normalised feature vectors suitable for both VotingClassifier and deep learning architectures. *Source:* [54].

Table 3.85 presents the filtering parameters with their neurophysiological justification.

Appendix C Objective. Provide detailed survey instruments, EEG pipeline flowcharts, clinical assessment forms, and methodology tables that support Chapter 3 without cluttering the main text.

Table 3.85. EEG Filtering Parameters: Type, Specifications, and Neurophysiological Justification

Filter	Type	Parameters	Purpose
High-pass	Butterworth 4th order	$f_c = 0.5$ Hz	Remove DC drift and slow electrode potential changes that obscure neural oscillations
Low-pass	Butterworth 4th order	$f_c = 45$ Hz	Remove electromyographic (muscle) artifacts above the gamma band upper limit
Notch	IIR	$f_0 = 50/60$ Hz; $Q = 30$	Remove power-line interference (50 Hz in Europe/Asia; 60 Hz in North America)
Band-pass	Butterworth 4th order	$f_L = 0.5$ Hz; $f_H = 45$ Hz	Combined high-pass and low-pass; retains all five canonical EEG bands (δ through γ)

Note. f_c = cutoff frequency; f_0 = centre frequency; Q = quality factor; f_L = lower cutoff; f_H = upper cutoff; IIR = infinite impulse response. Butterworth design chosen for maximally flat passband response; 4th-order provides 24 dB/octave roll-off. Parameters follow [54].

Table 3.86. Appendix C Roadmap: Contents and Purpose.

Content	What It Contains
Survey instruments	S1–S4 full item tables (20 items each)
Clinical forms	ASD, PD, Epilepsy assessment forms
EEG flowcharts	Phase 1–5 pipeline diagrams
Variable tables	Questionnaire variable mapping, research variables
Planning tables	Chapter alignment, SMART matrix, decision flow

Archived Comparative Preprocessing Note

The earlier dual-track preprocessing diagram combined the active ASD EEG workflow with an archived non-ASD imaging branch. That comparative figure is not used in the ASD-only thesis build and is therefore omitted here. The active preprocessing logic retained in this thesis is the EEG-based ASD pipeline described in Chapter 3, including signal quality checks, band-pass filtering, ICA-based artifact removal, epoching, normalisation, and feature extraction.

Note. The omitted comparative branch previously illustrated non-ASD image processing and is intentionally excluded from the ASD build to avoid ASD-focused ambiguity.

Comprehensive EEG Feature Extraction Taxonomy

Table 3.87 presents the five-category feature extraction taxonomy used in the ASD EEG pipeline. Each category captures a distinct aspect of neural dynamics relevant to ASD screening and interpretability:

- **Time-domain features:** Characterise signal morphology through statistical descriptors and Hjorth parameters.
- **Frequency-domain features:** Capture oscillatory power distributions across the five canonical EEG bands.
- **Time-frequency features:** Reveal non-stationary spectral changes via wavelet and short-time Fourier transforms.
- **Connectivity features:** Quantify inter-channel synchronisation using coherence and phase-locking metrics.
- **Nonlinear features:** Measure signal complexity and chaotic behaviour through entropy and fractal dynamics.

Together, these five categories produce the 127 features per channel documented in Table 3.

Table 3.87 categorises comprehensive EEG Feature Extraction Categories: Methods, Examples, and Clinical Relevance.

Archived Wearable Extension Note

Earlier supplementary drafts included a wearable multi-sensor extension and edge-deployment pathway. That material is not part of the ASD-only thesis build because the evaluated artifact and evidence base in this version are limited to the ASD EEG workflow and its associated questionnaire, governance, and explainability layers.

Note. The archived wearable extension is intentionally omitted here to keep Appendix C aligned with the ASD-only scope of the thesis.

Table 3.87. Comprehensive EEG Feature Extraction Categories: Methods, Examples, and Clinical Relevance

Category	Feature Family	Specific Features	Clinical Relevance
Time domain	Statistical, Hjorth params	Mean, variance, skewness, kurtosis, zero-crossing rate, line length, RMS, peak-to-peak; activity, mobility, complexity	Signal morphology and behavioural-state variation relevant to ASD screening
Frequency domain	Band powers, spectral ratios	Delta (0.5–4 Hz), theta (4–8 Hz), alpha (8–13 Hz), beta (13–30 Hz), gamma (30–45 Hz); relative powers; θ/β , α/β ratios	Oscillatory-state characterisation and ASD-relevant band-power differences
Time–frequency	Wavelets, STFT	CWT scalograms (Morlet wavelet, 64 frequency bins), STFT spectrograms (Hann window, 256 samples)	Non-stationary spectral dynamics and short-lived ASD-relevant shifts in EEG structure
Connectivity	Coherence, phase synchrony	Inter-channel coherence, phase-locking value (PLV), weighted phase-lag index (wPLI), mutual information	Inter-regional neural synchronisation and ASD-relevant connectivity differences
Nonlinear	Entropy, fractal dynamics	Shannon entropy, sample entropy ($m=2$, $r=0.2\sigma$), Lyapunov exponent, Hurst exponent, correlation dimension	Signal complexity, regularity, and developmental EEG variability relevant to ASD screening

Note. CWT = continuous wavelet transform; STFT = short-time Fourier transform; PLV = phase-locking value; wPLI = weighted phase-lag index; RMS = root mean square; PSD = power spectral density; ASD = autism spectrum disorder. Feature taxonomy follows the ASD EEG implementation documented in this thesis.

Data Risk Assessment Matrix

Figure 3.24 presents a risk assessment matrix for the nine data-related threats identified during the methodological design. Each risk is scored on two dimensions—likelihood (how probable the risk is given the study’s design safeguards) and impact (the severity of consequence if the risk materialises)—producing a composite risk level that informs the prioritisation of mitigation efforts.




Figure 3.24. Data Risk Assessment Matrix: Nine Risks by Composite Level

Note. Composite level = likelihood $\times$ impact. Three risks (R2, R4, R9) reside in the high-risk zone; none reach critical. SMOTE = Synthetic Minority Over-sampling Technique; CV = cross-validation; ICA = independent component analysis.

Table 3.88 documents the 12 MCP tools available to agents, organised by functional category. Each tool is exposed via the JSON-RPC 2.0 interface and subject to the access control and audit logging described in Table 3.79.

Table 3.88. NeuroMCP Tool Registry: 12 Tools Organised by Functional Category

Category	Tool Name	Function	Access
Data ingest	load_eeg	Load and validate raw EEG files (EDF/BDF/GDF)	All agents
Pre-process	filter_signal	Apply band-pass and notch filters	All EEG agents
	extract_features	Compute 127-feature vector per channel	All EEG agents
Inference	classify	Run VotingClassifier or ResNet-18 inference	All agents
	explain_shap	Generate SHAP feature importance values	All agents
RAG	retrieve_context	Query FAISS vector store for clinical literature	All agents
	generate_narrative	Produce evidence-grounded clinical explanation via SLM	All agents
Governance	check_guardrails	Validate input/output against safety rules (Sec. 3)	All agents
	log_audit	Write immutable audit record of prediction and explanation	System (auto)
	redact_pii	Detect and mask PII	System (auto)

Note. MCP = Model Control Protocol; EDF = European Data Format; BDF = BioSemi Data Format; GDF = General Data Format; PBS = peripheral blood smear; GAN = generative adversarial network; SHAP = SHapley Additive exPlanations; FAISS = Facebook AI Similarity Search; SLM = small language model; PII = personally identifiable information. Tool registry follows [54].

PII Protection and Security Architecture

The PII Protection flow (Figure 3.79) implements a seven-stage privacy pipeline. This subsection documents the specific detection methods, encryption standards, and regulatory alignment that ensure patient data protection throughout the pipeline.

Table 3.89 documents the six-stage PII protection pipeline ensuring patient data protection throughout the RGAIG system.

13-Phase ASD GenAI QA Framework overview table

Table 3.90 presents the thirteen-phase GenAI quality-assurance framework used to assess the ASD build, detailing each phase’s focus area, module count, key metrics, and sign-off gates.

Table 3.89. PII Protection Pipeline: Six-Stage Implementation

#	Process	Technology	Verification
1	Detect: identify PII entities in input text	Regex patterns (SSN, DOB, email, phone) + spaCy NER	Recall >99% on synthetic PII test set
2	Classify: categorise entities by sensitivity	PHI (HIPAA-defined), PII (general), sensitive (clinical notes)	Three-tier classification
3	Redact: replace entities with type tokens	“John Smith” → “[PERSON]”; “123-45-6789” → “[SSN]”	Redacted text reviewed by audit layer
4	Encrypt: protect data at rest and in transit	AES-256-GCM at rest; TLS 1.3 in transit	FIPS 140-2 compliant; 90-day key rotation
5	Access control: restrict by role and purpose	RBAC with three roles; GDPR Art. 5(1)(b) purpose limitation	Access denied without valid role + purpose
6	Audit: immutable access and processing log	Append-only SQLite log; tamper-evident hash chain	HIPAA §164.312(b) compliance

Note. PHI = protected health information; PII = personally identifiable information; NER = named entity recognition; RBAC = role-based access control; GCM = Galois/Counter Mode; FIPS = Federal Information Processing Standards.

Table 3.90. Thirteen-Phase ASD GenAI Quality Assurance Framework

Phase	Focus Area	Moc	Key Metrics	Sign-Off Gate
1	Knowledge & Data	16	Source inventory, authority scoring, coverage gaps, PHI/PII scan, ingestion completeness	All sources registered; 0 unhandled PHI
2	Representation & Retrieval	17	Chunk integrity, embedding quality, Recall@K, Precision@K, MMR, Table-RAG routing	Recall@K ≥ baseline; no index corruption
3	Generation & Reasoning	17	Citation coverage, claim verification, hallucination rate, numeric integrity, answer relevancy	Unsupported claims ≤3%; faithfulness ≥0.85
4	Decision Policy	17	Intent classification, risk-tier assignment, confidence calibration, escalation accuracy	False answers ≤ limit; 0 unsafe passes
5	Agent Behaviour	17	Task completion, tool selection, loop detection, error recovery, autonomy bounds	Goal completion ≥ target; 0 unsafe actions

6	Agent-to-Agent	17	Role adherence, message protocol, conflict resolution, coordination, failure isolation	$\geq 95\%$ role adherence; 0 cascade failures
7	Model Control Protocol	17	Policy coverage, tool access enforcement, bypass detection, prompt override resistance	0 policy bypasses; 0 unauthorised tool calls
8	Explainability & Trust	17	Evidence visibility, explanation consistency, confidence communication, comprehension	$\geq 97\%$ answers with visible sources
9	Robustness & Sensitivity	17	Prompt perturbation, context removal, adversarial inputs, sparsity handling	Graceful degradation; 0 adversarial successes
10	Statistical Validation	17	Hypothesis testing, sample adequacy, CI estimation, multiple correction, reproducibility	Power ≥ 0.8 ; $p < \alpha$; low variance
11	Benchmarking	17	RAG vs non-RAG, hybrid vs dense retrieval, chunking ablation, SOTA comparison, cost	RAG $\geq$ baseline; cost within budget
12	Scalability & Deploy	17	Token budget, retrieval/generation latency, throughput, cache hit rate, rollback readiness	SLA met; 100% reproducible builds
13	Governance & Compliance	17	Lineage completeness, PII exposure, consent validation, incident response, regulatory mapping	$\geq 99\%$ lineage; 0 PII exposure

Note. Each phase contains 16–17 monitoring modules, each with defined purpose, measurement method, outputs, pass/fail criteria, and edge-case fixes. In the ASD build, the 13-phase framework extends classical validation (Phases 10–11) with six GenAI-specific phases (1–3: knowledge and generation quality), four agentic phases (4–7: agent and policy governance), and two operational phases (8–9: trust and robustness). All 13 gates must pass before deployment in the bounded ASD screening context. PHI = protected health information; PII = personally identifiable information; MMR = maximal marginal relevance; MCP = Model Control Protocol; A2A = agent-to-agent; CI = confidence interval; SLA = service-level agreement; SOTA = state of the art.

8-Phase DL Analysis Framework overview table

Table 3.91 summarises the eight-phase deep learning analysis framework comprising 111 modules across data, model, performance, subject, sensitivity, statistical, benchmarking, and production monitoring phases.

Table 3.91. Eight-Phase DL Analysis Framework Overview (111 Modules)

#	Phase	Moc	Focus and Key Analyses	RGAIG Evidence
1	Data Analysis	9	Data inventory, class/label distribution, missingness, SQI, cross-subject variability, session drift, harmonisation, leakage audit, documentation	305 GB; 131 subjects + 20,000 images; SQI per channel; grouped split
2	Model Analysis	14	Model–data alignment, baselines (SVM/RF/LDA), complexity vs generalisation, training stability, regularisation, class-imbalance, calibration, interpretability, computational cost, failure-mode anticipation, readiness gate	7 baseline + 4 DL architectures; learning curves; SHAP; cost profiling
3	Performance Analysis	14	Metric correctness, fold-level performance, seed stability, CI, operating-point analysis, calibration, robustness, subgroup performance, error breakdown, baseline deltas, SOTA benchmark, sign-off gate	10-fold CV; bootstrap CI; ablation Δ 5.3–8.1%; SOTA comparison
4	Subject Analysis	14	Subject coverage, per-subject F1/AUC, worst-case, session drift, similarity clusters, device confounding, artifact sensitivity, label reliability, personalisation vs generalisation, demographic fairness, sign-off gate	Per-subject metrics; age/gender stratification; worst-case floor
5	Sensitivity Analysis	16	Filter cutoffs, notch settings, referencing (CAR/Cz), artifact strictness, ICA/ASR, window length (1–8 s), overlap, normalisation, class-imbalance strategy, HP neighbourhood, seed sensitivity, cross-dataset, sign-off	3-scenario + tornado; HP stability; 10-seed $\pm$ 0.8%

Continued on next page

6	Statistical Analysis	14	Hypothesis definition, effect sizes (Cohen's d /Hedges' g), CI, paired comparison (McNemar/DeLong), Holm/BH-FDR correction, Wilcoxon, subgroup disparity, model ranking, power adequacy, reporting pack	$d=0.82-1.48$; bootstrap CI; Holm correction; subject-level bootstrap
7	Benchmarking	14	Results registry, dataset/split table, preprocessing config, main results with CI, baseline delta, significance, robustness, sensitivity matrix, error taxonomy, explainability, literature matrix, claims-to-evidence map	Per-ASD pack; SOTA matrix; claims traceability
8	Production Monitoring	16	Data health, distribution drift (PSD), representation drift, prediction drift, confidence monitoring, latency, silent failure detection, drift-to-action policy, retraining readiness, A/B shadow testing, governance logging, ROI monitoring	Drift monitor; batch-wise accuracy; retraining triggers; model cards

Note. The eight-phase framework draws on implementations totalling 24,287 lines of analysis code across `data_analysis.py` (1,968 lines), `statistical_analysis.py` (1,405 lines), `cross_validation.py` (692 lines), `eda_analysis.py` (995 lines), `reliability_analysis.py` (995 lines), `filter_analysis.py` (727 lines), `drift_monitor.py` (664 lines), and `generate_ml_dl_analysis.py` (1,361 lines). SQI = signal quality index; PSD = power spectral density; HP = hyperparameter; BH-FDR = Benjamini–Hochberg false discovery rate.

10 Pillars of Cross-Cutting AI Governance overview table

Table 3.92 maps the ten cross-cutting AI governance pillars comprising 194 modules, documenting each pillar's core question and RGAIG implementation.

Table 3.92. Ten Pillars of Cross-Cutting AI Governance (194 Modules)

#	Pillar	Moc	Core Question and Key Analyses	RGAI Implementation
1	Energy Efficient AI	18	Is energy justified? Workload characterisation, architecture efficiency, right-sizing, PEFT/LoRA, quantisation (FP16→INT8), batching/caching, latency–energy trade-off, scaling, ESG alignment	Local SLM (3B params); INT8 quantisation; L1/L2 cache; 512-token limit
2	Hallucination Prevention	20	How are fabricated outputs prevented? Taxonomy, knowledge boundaries, prompt robustness, RAG grounding, source attribution, faithfulness, reasoning chain, uncertainty/abstention, adversarial tests, HITL, monitoring	Faithfulness 0.89; hallucination 2.1%; citation 94.2%; 6-gate guardrail
3	Hypothesis in AI	20	Are assumptions stated and tested? Problem framing, data sufficiency, label validity, feature relevance, leakage, model capacity, convergence, generalisation, imbalance, error patterns, robustness, explainability, causal mechanisms, drift, safety	H1–H5 pre-registered; 5 falsifiable criteria; power ≥ 0.80
4	Threat AI	20	What adversarial threats exist? Asset identification, attack surface, data poisoning, prompt injection, model extraction, membership inference, adversarial examples, RAG poisoning, DoS, supply chain, incident response, residual risk	PII scrubbing; RBAC; prompt override testing; red-team validation
5	SWOT for AI	16	Strategic assessment: capability, data advantage, efficiency (S); data dependency, XAI gaps, generalisation limits (W); new products, scale, differentiation (O); regulatory pressure, ethical risk, security (T)	SWOT by lifecycle stage (data, model, deploy, monitor, govern)

Continued on next page

6	Governance AI	20	Who owns AI decisions? Scope/authority, RACI, policy alignment, risk management, ethics review, data governance, model lifecycle, monitoring, incident escalation, HITL, vendor governance, documentation, enforcement	4-level escalation ladder; RACI matrix; MCP policy engine
7	Compliance AI	20	Which laws apply? Jurisdiction mapping (5), risk classification, GDPR/HIPAA, fairness, safety controls, human oversight, right-to-explanation, post-market surveillance, vendor compliance, audit readiness, accountability	35+ standards; 5-jurisdiction compliance; audit-ready docs
8	Responsible AI	20	Are outcomes equitable and harm-minimised? Stakeholder impact, misuse analysis, fairness, transparency, oversight, accountability, privacy-by-design, safety, monitoring, contestability, environmental/social impact	9 responsibility dimensions; Fairlearn bias audit; SHAP
9	Explainable AI	20	Can decisions be understood? SHAP/LIME local, global feature importance, PDP/ICE, counterfactuals, faithfulness, stability, bias audit, temporal XAI, Grad-CAM, token attribution, interpretability, reproducibility	SHAP top-5; 3-tier RAG explanation; expert $M=4.6/5.0$; $\alpha=0.84$
10	Secure AI	20	Are AI assets protected? Asset inventory, threat model, data integrity, training security, model IP, adversarial robustness, prompt injection defence, RBAC, privacy/leakage, supply chain, logging/audit, incident response, penetration testing	AES-256; RBAC; PHI/PII detection; immutable audit logs

Note. The ten pillars address cross-cutting concerns that span all pipeline stages. Module counts: Energy-Efficient (18), Hallucination Prevention (20), Hypothesis (20), Threat (20), SWOT (16), Governance (20), Compliance (20), Responsible AI (20), Explainable AI (20), Secure AI (20). HITL = human-in-the-loop; PEFT = parameter-efficient fine-tuning; LoRA = low-rank adaptation; SWOT = strengths, weaknesses, opportunities, threats; DoS = denial of service; ESG = environmental, social, and governance.

Responsible AI Governance Summary table

Table 3.93 summarises responsible AI practices across five governance pillars.

Table 3.93. Executive Summary: Responsible AI Governance Framework

Governance Requirement Pillar	Implementation in This Study	Evidence / Verification
Data Compliance		
HIPAA de-identification	All datasets pre-anonymized by original collectors	Data use certificates (Appendix 3.120)
GDPR data minimization	Only task-relevant features extracted; no surplus data retained	Preprocessing pipeline documentation (Section 3)
IRB exemption	Publicly available de-identified data; 45 CFR 46 exemption	Institutional compliance letter
Encrypted storage	Raw data on AES-256 encrypted local workstation; no cloud	Single-user access; 5-year retention policy
Algorithmic Fairness		
Bias detection	Demographic subgroup accuracy analysis	$\leq 2\%$ variance across available subgroups (Ch. 4)
Transparency	SHAP feature attribution + RAG literature-grounded narratives	Expert panel validation ($M = 4.6/5$)
Human oversight	AI recommends, clinician decides; no autonomous decisions	Human-in-the-loop design; escalation protocols
Data Provenance		
Source verification	NIH/NITRC, PhysioNet, university repositories	KAU ASD dataset with documented provenance (Table 3)
Access protocol	Data use agreements signed per dataset	DUA records maintained; access logging enabled
Audit trail	All preprocessing steps logged with parameters	Reproducible via Docker + fixed seeds + Git
Reproducibility		
Environment	Docker containers with pinned dependency versions	Dockerfile versioned in repository

Determinism	Fixed random seed (42) across all frameworks	PyTorch, TensorFlow, NumPy, Scikit-learn
Experiment tracking	MLflow for hyperparameters, metrics, artifacts	Tagged experiment commits in Git
Regulatory Alignment (Future Deployment)		
FDA SaMD	Framework designed for IEC 62304 software lifecycle	Governance architecture supports SaMD pathway
EU AI Act	High-risk classification; human oversight mandated	Audit trails, bias docs, decision logic embedded
WHO/ITU GMLP	Good Machine Learning Practice principles	Validation, monitoring, documentation controls

Note. HIPAA = Health Insurance Portability and Accountability Act; GDPR = General Data Protection Regulation; IRB = Institutional Review Board; SHAP = SHapley Additive exPlanations; RAG = retrieval-augmented generation; FDA = Food and Drug Administration; SaMD = Software as a Medical Device; EU = European Union; WHO = World Health Organization; ITU = International Telecommunication Union; GMLP = Good Machine Learning Practice; DUA = data use agreement. Regulatory alignment items represent design intent for future clinical deployment, not current certification status.

Engineering Pipeline Flowcharts

This section presents eight engineering-detail flowcharts that document the RGAIG pipeline at implementation level. Figures 3.25–3.29 trace the five core processing phases from data acquisition through model training. Figures 3.30–3.32 provide three end-to-end pipeline overviews at increasing levels of scope.

Phase 1: Data Acquisition and Ingestion. Figure 3.25 documents the data acquisition and ingestion framework, covering multi-source data loading, format validation, quality checks, and the initial data registry that feeds all downstream phases.





Figure 3.25. Phase 1: Data Acquisition and Ingestion Framework

Phase 2: Signal Preprocessing. Figure 3.26 details the signal preprocessing and cleaning pipeline, including bandpass filtering, artifact rejection, re-referencing, and segment extraction that transforms raw EEG recordings into analysis-ready segments.






Figure 3.27. Phase 3: ASD EEG feature extraction pipeline. Clean EEG segments pass through four parallel feature-family modules—statistical, spectral, temporal, and nonlinear—then merge into a unified 47-feature vector. Right panel summarises why multi-family extraction matters and what each feature family captures.

Figure 3.28 presents the feature selection and dimensionality reduction funnel, which reduces the 47 extracted features to an optimised set of 25 through importance ranking, redundancy removal, and variance filtering.




Figure 3.28. Phase 4: feature selection and dimensionality reduction. The 47 features from Phase 3 undergo importance ranking, redundancy removal, and PCA-based reduction to a 25-feature set that retains more than 90% of variance.

Figure 3.29 depicts the hybrid AI model training architecture, showing the dual-branch design (three baseline classifiers and three deep learning models) that converges into a majority-vote ensemble with calibrated probability outputs.




Figure 3.29. Phase 5: Hybrid AI model training and computer vision pipeline. The left column shows the six-step process from optimized features through 1D-to-2D transformation, parallel classical ML and deep learning pipelines, majority-vote ensembling, validation, and prediction output. The right column provides methodological justification for each design decision. Bottom row reports the achieved ASD classification accuracy (97.67%).

Figure 3.30 presents the 13-stage signal processing pipeline from raw EEG capture through to model-ready feature vectors, organised into four functional groups: signal acquisition, signal cleaning, signal transformation, and feature engineering.




Figure 3.30. EEG Signal Processing Pipeline: From Raw Capture to Model-Ready Features

Figure 3.31 documents the 13-stage model training and computer vision pipeline, covering data preparation, model architecture selection, cross-validated training, and evaluation through to deployment-ready model output.




Figure 3.31. Model training and computer vision pipeline from prepared data to deployment-ready ASD models.

Figure 3.32 provides the complete 16-stage end-to-end data flow from initial EEG capture through preprocessing, feature engineering, model inference, RAG-based explanation generation, and clinical decision support output.




Figure 3.32. RGAIG end-to-end data flow: 16-stage pipeline from EEG capture to clinical decision support.

RGAIG 16-Stage Pipeline: Stage Descriptions and.... Table 3.94 below presents the detailed analysis.

Chapter Planning and Alignment Tables

The following tables provide detailed chapter-wise planning, alignment, and readiness assessment material supporting the methodology design in Chapter 3.

Table 3.95 presents chapter-Wise Objective, Process, and Decision Structure (Appendix Version).

Table 3.94. RGAIG 16-Stage Pipeline: Stage Descriptions and Functions

Stage	Stage Name	Function
1	EEG Signal Capture	Emotiv headset acquires multi-channel brain signals at 128–256 Hz
2	Mobile Reception	Bluetooth Low Energy streams raw EEG to companion mobile application
3	Secure Cloud Gateway	TLS-encrypted transfer to cloud ingestion with session tagging
4	Raw Data Storage	Versioned data lake with de-identification and access audit logging
5	Preprocessing Pipeline	Artifact removal, bandpass filtering, ICA denoising, Z-score normalization
6	Signal Transformation	1D time-series converted to 2D CWT scalograms and STFT spectrograms
7	Feature Extraction	Statistical, spectral, and connectivity features with mutual information selection
8	AI Model Prediction	Ensemble of RF, SVM, CNN, LSTM classifiers generates probability scores
9	Consensus Aggregation	Weighted voting combines per-model predictions into unified confidence score
10	SHAP Explainability	Top- <i>k</i> features identified with direction and magnitude via SHAP values
11	RAG Context Assembly	Medical literature queried to contextualize SHAP-identified features
12	Knowledge Retrieval	ChromaDB vector store, Neo4j graph DB, Redis cache supply clinical evidence
13	LLM Report Generation	Llama-3.2 synthesizes evidence and explanations into diagnostic report
14	Report Delivery	Report dispatched to parent portal, clinician dashboard, EHR via HL7/FHIR
15	Governance & Compliance	HIPAA/GDPR enforcement, AES-256 encryption, audit trail, bias monitoring
16	Clinical Decision Support	Explainable output supports evidence-based triage, referral, and treatment

Table 3.95. Chapter-Wise Objective, Process, and Decision Structure (Appendix Version)

	Chapter Objective	Goal	Task	Input	Output	Decision Role
1	Introduce the study	Establish need and scope	Frame problem, gap, objectives, hypotheses, scope	Research problem, context, preliminary literature	Problem statement and study direction	Justifies what to address
2	Review literature	Justify study intellectually	Compare studies, identify gaps, build conceptual logic	Published studies, theories, benchmarks	Gap-driven literature synthesis	Whether study is justified
3	Explain methods	Show how claims are tested	Define data, variables, preprocessing, modelling, validation	Primary instruments, EEG data, method choices	Reproducible study design	Whether claims can be tested
4	Present findings	Show what evidence demonstrates	Report results, test hypotheses, integrate findings	Primary results, EEG outputs, robustness evidence	Empirical findings and verdicts	Whether framework is supported
5	Interpret and conclude	Translate evidence into action	Assess readiness, implications, final recommendation	Ch. 4 findings, governance logic	Adopt/pilot/d conclusion	What should be done next

Table 3.96 documents detailed Chapter-Wise SMART, Evidence, Reliability, and Decision Matrix (Appendix Version).

Table 3.96. Detailed Chapter-Wise SMART, Evidence, Reliability, and Decision Matrix (Appendix Version)

Ch	Objective	Spec	Measura	Achievat	Rel.	Tim	Evidence	Reliability	Decision Use
1	Define problem, gap, objectives, hypotheses, scope	Yes	Problem clarity, objective align-ment, scope clarity	Yes, if framing stays bounded	Yes	Yes	Problem statement, gap, stakeholder relevance, hypothesis map	Moderate–strong if duplication removed	What the thesis addresses and why
2	Synthesize literature, identify gaps	Yes	Theme coverage, contra-diction count, gap clarity	Yes, if focused on synthesis	Yes	Yes	Comparative review, gap matrix, theory linkage	Strong if selective and critical	Whether the study is justified
3	Design method-ological framework	Yes	Variables, datasets, metrics, validation, KPI thresh-olds	Yes, if protocol detail con-trolled	Yes	Yes	Method tables, EEG pipeline, variable map, governance logic	Strong if explicit and repro-ducible	Whether claims can be tested credibly

Ch	Objective	Spec	Measura	Achieva	Rel.	Tim	Evidence	Reliability	Decision Use
4	Present findings, evaluate hypotheses	Yes	Accuracy, F1, AUC, trust, usability, robustness	Yes, if selective evidence displays	Yes	Yes	Performance tables, ROC, SHAP, trust findings	Very strong if clean reporting	Whether framework is supported
5	Interpret findings, deployment decisions	Yes	Readiness level, contribution quality, adoption status	Yes, if bounded and evidence-led	Yes	Yes	Readiness matrix, contribution summary, decision framework	Strong if domain-sensitive	Whether to adopt, pilot, or defer

Table 3.97 summarises chapter-Wise Decision Role Summary (Appendix Version).

Table 3.97. Chapter-Wise Decision Role Summary (Appendix Version)

Chapter	Core Function	Decision Role
1	Frame the problem and define scope	Decides what the thesis must address
2	Justify the study through literature and gaps	Decides why the study is needed
3	Design the method and evaluation logic	Decides how the claims will be tested
4	Report evidence and hypothesis results	Decides whether the claims are supported
5	Interpret results and set deployment stance	Decides what should be done in practice




Figure 3.33. Chapter-Wise Thesis Decision Flow (Appendix Version)

Note. The thesis progresses from problem framing to literature justification, methodological design, empirical evidence, and final decision logic.

Thesis Alignment and Readiness Assessment

Table 3.98 presents thesis Objective Alignment Across Chapters (Appendix Version).

Table 3.98. Thesis Objective Alignment Across Chapters (Appendix Version)

Objective Component	Ch. 1	Ch. 2	Ch. 3	Ch. 4	Ch. 5	Overall
Define the biomedical AI problem	Strong	Supportive	Indirect	Indirect	Indirect	Strong
Justify why literature is insufficient	Moderate	Strong	Supportive	Indirect	Indirect	Strong
Develop unified ASD-focused framework	Framing only	Justification only	Strong	Tested empirically	Interpreted strategically	Strong
Integrate primary and secondary data	Previewed	Justified as gap	Strongly operationalized	Partially evidenced	Interpreted in deploy terms	Strong
Evaluate technical EEG performance	Previewed early	Justified	Strongly designed	Strongly evidenced	Interpreted	Strong
Evaluate explainability and trust	Framed strongly	Justified strongly	Methodologic supported	Evidenced	Interpreted for adoption	Strong
Evaluate governance and safe deployment	Framed strongly	Justified conceptually	Strongly designed	Partly evidenced	Heavily interpreted	Moderate–strong

Objective Component	Ch. 1	Ch. 2	Ch. 3	Ch. 4	Ch. 5	Overall
Evaluate workflow and operational value	Framed strongly	Justified moderately	Methodologic included	Partly evidenced	Strongly interpreted	Moderate
Support B2B hospital decision making	Framed strongly	Justified moderately	Operationaliz	Partly evidenced	Strongly interpreted	Strong
Bound B2C as extension or future path	Over-expanded	Moderately justified	Included but broad	Limited evidence	Interpreted but broad	Moderate
Produce adopt/pilot/defer decision logic	Previewed early	Not central	Designed	Partially evidenced	Main destination	Moderate–strong

Table 3.99 documents chapter Objective versus Thesis Objective Matrix (Appendix Version).

Table 3.99. Chapter Objective versus Thesis Objective Matrix (Appendix Version)

Chapt	Chapter Objective	Link to Thesis Objective	Strength	Main Problem
Ch. 1	Introduce problem, gap, objectives, hypotheses, scope	Defines why the thesis exists	Moderate–strong	Method and architecture creep in introduction
Ch. 2	Synthesize literature and identify gaps	Justifies why the thesis is needed	Strong	Synthesis weakened by duplication

Chapt	Chapter	Link to Thesis	Strength	Main Problem
	Objective	Objective		
Ch. 3	Explain framework and study design	Operationalizes how thesis objective will be tested	Strong	Architecture and protocol overload
Ch. 4	Present primary, secondary, and integrated evidence	Shows whether thesis objective is empirically supported	Very strong	Too many tables and support displays
Ch. 5	Interpret findings and make final decisions	Explains what thesis objective means in practice	Moderate–strong	Decision diluted by adjacent frameworks

Table 3.100 presents chapter-Wise Strengths, Weaknesses, and Action Plan (Appendix Version).

Table 3.100. Chapter-Wise Strengths, Weaknesses, and Action Plan (Appendix Version)

Chapter	Strengths	Weaknesses	Action Plan
Ch. 1	Strong problem framing, stakeholder relevance, broad scope	Overloaded introduction, methods creep, workflow clutter	Cut operational detail; move method content to Ch. 3; retain framing only
Ch. 2	Broad evidence base, meaningful gap identification	Duplication, inventory-style summaries, too many side frameworks	Tighten synthesis; remove duplicates; focus on gap-to-study logic

Table 3.100 continued

Ch. 3	Strong methodological ambition, integrated design, good coverage	Bloated chapter, too many master tables, protocol overload	One method flow per strand; move protocols to appendix
Ch. 4	Strongest empirical chapter, explicit hypothesis testing	Too many support tables, repeated displays, discussion leakage	One table + one figure per claim; move helpers to appendix
Ch. 5	Rich interpretation, strong DBA orientation, practical relevance	Too many matrices, policy sprawl, weak prioritization	Compress around readiness, ROI, adopt/pilot/defer conclusion

Table 3.101 documents adopt, Pilot, or Defer Readiness Matrix.

Table 3.101. Adopt, Pilot, or Defer Readiness Matrix

Evaluation Area	Adopt	Pilot	Defer	Interpretation
Technical performance	High	Moderate to high	Low	Accuracy, recall, F1, AUC, and robustness status
Explainability and trust	High	Moderate	Low	SHAP clarity, RAG usefulness, expert trust, and user understanding
Governance readiness	High	Moderate	Low	Privacy, auditability, access control, and human-in-the-loop status
Workflow fit	High	Moderate	Low	Time reduction, usability, process fit, and burden reduction
Domain readiness	Strong	Conditional	Weak	Domain-specific maturity and validation sufficiency
Overall decision	Deploy selectively	Controlled trial use	Hold for future work	Final deployment stance based on combined evidence

Table 3.102 documents domain-Sensitive Deployment Decision Matrix.

Table 3.102. Domain-Sensitive Deployment Decision Matrix

Domain	Technical	Trust	Governance	Recommended Decision
ASD	Strong	Strong	Moderate to strong	Pilot or selective adopt

Engineering Detail Migrated from Chapter 3 Main Body

The following tables and figures present full engineering detail supporting the DBA research methods narrative in Chapter 3. They were relocated here so the main chapter reads as a complete research-methods story while preserving all technical evidence for examination.

Master Logic Chain

Table 3.103 presents master Research Logic: Gap to Decision Chain.

Table 3.103. Master Research Logic: Gap to Decision Chain

Research Gap	Obj.	Variable	Analysis	Hypothesis / Logic	Decision Implication
No integration of primary and secondary evidence	O7, O11	Primary–secondary alignment, contextual completeness	Correlation, concordance, joint display	If evidence aligns, integrated workflow justified	Retain integrated assessment design
Behavioural/psychological assessment not captured in AI workflows	O7, O9	Behavioural, psychological, functional scores	Descriptive, reliability, severity profiling	If stable scores derivable, primary data adds valid context	Include behavioural/psychological layer
Clinical AI explainability weak and hard to trust	O8	Trust, clarity, usefulness scores	Descriptive, Cronbach’s α , regression	H5: Better explanation quality increases trust	Keep or improve explanation module

Table 3.103 continued

Research Gap	Obj.	Variable	Analysis	Hypothesis / Logic	Decision Implication
Hospital onboarding fragmented and inefficient	O5, O6	Usability, completeness scores, time	Descriptive, group comparison	H4: Guided onboarding improves completeness	Adopt chatbot intake if thresholds met
Existing systems ignore caregiver perspective (ASD)	O9	Caregiver behavioural, developmental scores	Descriptive, reliability, correlation	Caregiver patterns provide interpretable context	Use caregiver data as formal primary input
Patient symptom burden not linked to AI outputs	O9, O11	Symptom severity, model confidence	Correlation, cross-tabulation	Higher burden may align with disease-specific output	Use patient report as contextual support
AI workflows lack human acceptance focus	O8, O15	Adoption, usability, trust scores	Descriptive, regression, triangulation	If acceptance low, technical performance insufficient	Pilot only, not broad deployment
Primary-data instruments not validated	O9	Internal consistency of subscales	Cronbach's α , item-total	Scales should meet acceptable reliability	Retain stable scales, revise weak ones
No evidence context improves decision usefulness	O11, O15	Contextual usefulness, expert rating	Comparison, regression, joint display	Integrated context improves decision usefulness	Support mixed-method justification
B2C continuity proposed but not assessed	O14	Monitoring readiness, privacy confidence	Descriptive, thematic interpretation	If readiness weak, B2C remains future scope	Keep B2C as secondary extension
Governance requires user trust and privacy comfort	O12	Privacy confidence, trust, explanation acceptance	Descriptive, threshold testing	Acceptable governance scores support deployment	Move toward pilot if thresholds met

Table 3.103 continued

Research Gap	Obj.	Variable	Analysis	Hypothesis / Logic	Decision Implication
Prior studies lack mixed-method validation	O11, O15	Convergent/diverge evidence	Joint display, meta- inference	If quant and qual converge, deployment case strengthens	Final adopt/pilot/defer decision

ASD Questionnaire Blueprint

Table 3.104 presents the ASD Parent Questionnaire Blueprint, specifying respondent roles, item counts, and output variables across the seven assessment sections.

Table 3.104. ASD Parent Questionnaire Blueprint: Sections, Items, and Output Variables

Section	Respondent	Focus	Items	Output Variable
Develop. hist.	Parent	Speech delay, milestones, onset of concern	5–8	asd_dev_score
Social comm.	Parent	Eye contact, name response, peer engagement	6–8	asd_social_score
Repet. behav.	Parent	Routines, repetitive actions, fixation	4–6	asd_repet_score
Sensory proc.	Par./Pat.	Sound, touch, taste, visual overload	4–6	asd_sensory_score
Emotional reg.	Par./Pat.	Meltdown frequency, frustration, recovery	4–6	asd_emotion_score
Function. impact	Parent	School, home, communication, social	4–5	asd_function_score
Self-report	Patient	Stress, anxiety, overload, comfort	3–5	asd_selfrep_score
Total estimated items			30–44	

Note. Par. = Parent; Pat. = Patient. All structured items use 5-point Likert or frequency scales (0 = Never to 4 = Very Often). Composite scores are computed as section means. Reliability target: Cronbach's $\alpha \geq 0.70$.

PD Questionnaire Blueprint (Comparator)

Table 3.105 presents the Parkinson's Disease (PD) Patient Questionnaire Blueprint, included here as a methodological comparator illustrating how the ASD instrument's section/respondent/item/output structure generalises to an adult-onset neurodegenerative condition. The PD instrument is not administered as part of this dissertation; the comparator establishes content-validity portability of the questionnaire design pattern.

EEG Analysis Pipeline

Secondary EEG Data: Quantitative Analysis Pipeline. Table 3.106 below presents the detailed analysis.

Table 3.105. PD Patient Questionnaire Blueprint (Comparator): Sections, Items, and Output Variables

Section	Respondent	Focus	Items	Output Variable
Motor symptoms	Patient	Tremor, rigidity, slowness, gait, balance	5–8	pd_motor_score
Non-motor sympt.	Patient	Fatigue, sleep, constipation, smell loss	4–6	pd_nonmotor_score
Cognitive sympt.	Patient	Memory, concentration, planning difficulty	4–5	pd_cognitive_score
Emotional sympt.	Patient	Anxiety, depression, apathy, frustration	4–5	pd_emotional_score
Daily living	Patient	Dressing, eating, movement, speaking	4–6	pd_adl_score
Caregiver obs.	Caregiver	Falls, freezing, mood shifts, variability	3–5	pd_caregiver_score
Total estimated items			24–35	

Note. Comparator instrument; not administered in this dissertation. Motor and non-motor sections use 5-point severity scales (1 = Very Low to 5 = Very High). Cognitive and emotional sections use 5-point agreement scales. Reliability target: Cronbach's $\alpha \geq 0.70$.

EEG 40-Step Pipeline

Table 3.107 presents complete End-to-End EEG Pipeline: 40-Step Methodology.

Table 3.107. Complete End-to-End EEG Pipeline: 40-Step Methodology

#	Stage	What Happens	Why It Matters	Risk if Missing
1	Problem defn	Define disease, use case, stakeholder, context	Gives pipeline clear purpose	Technically strong but academically weak
2	Data source defn	Identify EEG source: Emotiv/IoT, dataset, hospital	Avoids ambiguity about origin	Data credibility weak
3	Consent & gov.	Patient consent, ethics, access, audit rules	Required for clinical/mobile use	Unsafe/non-compliant claim
4	Patient onboard.	Registration, symptom intake, parent input	Connects human context to EEG	No human context for decisions
5	Signal acquisition	Collect multichannel EEG	Core biomedical input	No biomedical evidence
6	Signal quality	Check impedance, missing ch., noisy segments	Prevents bad data in pipeline	Noisy data undermines findings

Table 3.107 continued

#	Stage	What Happens	Why It Matters	Risk if Missing
7	Artefact removal	Remove blink, muscle, motion, powerline	Improves reliability	Inflated error, unstable features
8	Band-pass filter	Retain useful frequencies (0.5–45 Hz)	Standard preprocessing	Irrelevant noise remains
9	Segment./epoch	Divide EEG into windows (2–5 s)	Creates comparable units	Inconsistent analysis units
10	Baseline corr.	Normalise reference/baseline activity	Improves comparability	Distorted channel interp.
11	Time-domain	Hjorth parameters, variance, mean, RMS	Captures signal behaviour	Misses waveform-level info
12	Freq.-domain	FFT, PSD/SPDD	Extracts spectral content	No band-power interp.
13	Time-freq.	SWT, wavelet, scalogram	Analyses changing signal	Transient patterns missed
14	Feature extr.	PSD, entropy, coherence, wavelets, deep	Converts to model-ready form	Raw data unusable
15	Feature select.	PCA, mRMR, LASSO, importance ranking	Keeps strongest, least redundant	Overfitting, redundancy
16	Normalisation	Z-score, min-max scaling	Standardises magnitude	Unstable training
17	1D to 2D conv.	Spectrogram, scalogram, topomap	Enables CNN/CV models	CV models unavailable
18	Safe splitting	Subject-wise split, LOSO, CV	Prevents leakage	False perf. inflation
19	Class imbalance	Weighted loss, oversampling, SMOTE	Reduces bias from skew	Misleading accuracy
20	Model selection	ML, 1D CNN, 2D CNN, transformer, ensemble	Choose modelling family	Arbitrary arch. choice
21	Hyperparam. tune	Grid search, Bayesian tuning	Optimise training	Weak/unstable performance
22	Training	Fit model to data	Core modelling step	No predictive system
23	Validation	Monitor development performance	Tuning feedback	Overfitting risk
24	Testing	Final unbiased evaluation	Proves performance	No trustworthy claim
25	Core evaluation	Confusion matrix, recall, acc., F1, AUC	Technical evidence	No hypothesis testing basis

Table 3.107 continued

#	Stage	What Happens	Why It Matters	Risk if Missing
26	Threshold tune	ROC thresholding, calibration curve	Align to screening use	Unsafe FP/FN balance
27	Robustness	Bootstrap, ablation, LOSO	Test stability	Results may be fragile
28	Explainability	SHAP, saliency, attention, feature maps	Interpret model behaviour	Black-box untrusted
29	Clinical plaus.	Compare features with domain knowledge	Plausible biomed. interp.	Explain. superficial
30	RAG interp.	Retrieval-augmented generation	Evidence-backed summary	Expl. generic/unsupported
31	Output format	Clinician, patient, admin views	Usable result presentation	Same output confuses users
32	Decision logic	Risk rules, escalation logic	Review/monitor/escalate	Prediction $\neq$ action
33	MCP/orchestr	Model context protocol workflow	Coordinated execution	Disconnected components
34	Mobile app	App onboarding, reports, monitoring	Patient-facing continuity	Weak B2C continuity
35	Hospital integr.	EHR/HIS/dashboard/reports API	Clinical workflow adoption	Hard to deploy
36	SOS/escalation	Alert engine, escalation workflow	Safety response pathway	High-risk cases missed
37	HITL	Review dashboard, override mechanism	Safe decision-support use	Overclaim as autonomous
38	Logging/audit layer	Audit logs, traceability	Governance evidence	Weak compliance
39	ROI/workflow	Time-motion, throughput, effort	Business case for adoption	B2B argument incomplete
40	Final decision	Adopt/pilot/defer matrix	Readiness conclusion	No usable conclusion

Table 3.106. Secondary EEG Data: Quantitative Analysis Pipeline

Step	Action	Output
1	Acquire or load EEG data	Analysis dataset
2	Clean artefacts and noise	Usable EEG segments
3	Segment / epoch data	Comparable analysis windows
4	Extract features or generate deep representations	Model-ready inputs
5	Split data for training/validation/testing	Valid evaluation design
6	Train baseline and advanced models	Comparative results
7	Compute performance metrics	Accuracy, AUC, F1, sensitivity, specificity
8	Test robustness	Confidence intervals, ablation, LOSO/CV
9	Run explainability analysis	Feature importance / signal relevance
10	Translate results into decision usefulness	Adopt / pilot / defer logic

EEG Preprocessing Detail

Table 3.108 presents eEG Preprocessing and Signal Analysis Methods.

Table 3.108. EEG Preprocessing and Signal Analysis Methods

Step	Method	Purpose	Example Use
Noise removal	Band-pass filter	Remove irrelevant frequency components	Keep 0.5–40 Hz
Artefact handling	ICA, filtering, rejection	Improve usable signal quality	Cleaner epochs
Segmentation	Epoching/win	Create comparable units	2 s or 5 s windows
Spectral analysis	FFT / PSD	Quantify alpha, beta, theta, delta power	Biomarker profiling
Time-frequency	SWT / wavelet	Detect transient and local changes	ASD/PD signal irregularities
Normalisation	Z-score, min-max	Standardise features	Stable model input
Representation	1D to 2D mapping	Allow image-based modelling	Spectrogram, scalogram, topomap

EEG Feature Types

EEG Feature Extraction: Types and Applications. Table 3.109 below presents the detailed analysis.

Table 3.109. EEG Feature Extraction: Types and Applications

Feature Type	Examples	Why Useful
Time-domain	Mean, variance, Hjorth parameters	Simple signal characterisation
Frequency-domain	Band power, spectral entropy, peak frequency	Strong EEG biomarker use
Time-frequency	Wavelet coefficients, scalogram features	Capture dynamic signal changes
Connectivity	Coherence, correlation, phase-locking	Inter-channel relationships
Nonlinear	Entropy, fractal dimension	Complex signal behaviour
Deep features	CNN latent embeddings	Automatic learned representations

EEG Modelling Strategies

EEG Modelling Strategies: Input, Strength, and.... Table 3.110 below presents the detailed analysis.

Table 3.110. EEG Modelling Strategies: Input, Strength, and Application

Strategy	Input Type	Best Use	Strength
Classical ML	Extracted features	Smaller structured datasets	Interpretable baseline
1D CNN	Raw/filtered series	Temporal EEG patterns	Direct signal learning
LSTM / transf.	Sequences	Temporal dependency	Sequence-aware modelling
2D CNN	Spectrogram scalogram	Image-like EEG repr.	Strong visual pattern learning
CV model	2D converted EEG	Complex spatial-freq. patterns	After 1D to 2D conversion
Ensemble	Combined outputs	Improve robustness	Stable overall performance

EEG System Integration

EEG Framework System Integration Components. Table 3.111 below presents the detailed analysis.

Table 3.111. EEG Framework System Integration Components

Integration Area	What It Does	Value
Patient mobile app	Onboarding, symptom forms, report access, monitoring	Patient-side continuity
Parent/caregiver app	Behavioural input, observation logging, alerts	Primary-data capture
Hospital system	EHR/HIS integration, clinician dashboard, report storage	Workflow compatibility
MCP layer	Coordinates model, retrieval, records, tools, notifications	System orchestration
SOS/emergency	Trigger escalation when threshold exceeded	Safety and response continuity

Questionnaire Variable Mapping

Table 3.112 maps primary-Data Instrument: Questionnaire Variable Mapping.

Table 3.112. Primary-Data Instrument: Questionnaire Variable Mapping

Section	Resp. Focus	Variable	Type	Analysis
Demograph	Par./P Age, gender, clin. history	demographic_ vars	Nom./	Descriptive, subgrp
Social comm. (ASD)	Parent Eye contact, reciprocity, peers	asd_social_ score	Likert	Descriptive, reliab.
Repet. behav. (ASD)	Parent Routines, repetitive acts, fixation	asd_ repetitive_ score	Likert	Descriptive, subgrp
Sensory (ASD)	Par./P Noise, touch, visual overload	asd_ sensory_ score	Likert	Severity scoring
Emot. reg. (ASD)	Par./P Meltdown, frustra- tion, recovery	asd_ emotion_ score	Likert	Corr. w/ trust
Trust in AI	All Trust, clinical reason- ableness	trust_score	Likert	Mean, SD, α

Table 3.112 continued

Section	Resp. Focus	Variable	Type	Analysis
Expl. clarity	All	Understand clarity__ useful, score clear	Likert	Corr. w/ trust
Usability	Par./P Chatbot	usability__ ease, score question clarity	Likert	Onboarding assess.
Adoption	All	Would use adoption__ again, rec- score ommend	Likert	Threshold testing
Privacy	Par./P Data	privacy__ sharing score comfort, privacy	Likert	Governance analysis
Open-ended	Par./P Concerns,	coded__ challenges, themes feedback	Qual.	Thematic coding

Primary Data Pipeline

Primary-Data Quantitative Analysis Pipeline. Table 3.113 below presents the detailed analysis.

Table 3.113. Primary-Data Quantitative Analysis Pipeline

Step	Action	Output
1	Convert structured responses to numeric form	Analysis-ready dataset
2	Build subscale scores	ASD behaviour score, trust score
3	Test reliability	Cronbach's α / consistency evidence
4	Summarise scores	Means, SD, frequencies, severity categories
5	Compare groups or conditions	Differences across ASD subgroups or response groups
6	Correlate with EEG/model outputs	Association between questionnaire burden and model evidence
7	Evaluate threshold achievement	Trust $\geq$ threshold, usability $\geq$ threshold
8	Integrate with qualitative findings	Joint interpretation via mixed-method display

EEG Pipeline Diagram




Figure 3.34. Secondary EEG Processing Pipeline

Note. Signal transformation includes FFT, PSD/SPDD, and SWT or wavelet-based analysis depending on domain requirements.

Analytical Pipeline Diagram




Figure 3.35. Five-Stage Analytical Pipeline for ASD Diagnosis

Note. PSD = power spectral density; CWT = continuous wavelet transform; CSP = common spatial pattern; ICA = independent component analysis; DL = deep learning; SHAP = SHapley Additive exPlanations; RAG = retrieval-augmented generation.

C4 Container Architecture




Figure 3.36. C4 Container Diagram: Mobile App, Gateway, Backend, RAG, MCP, and Hospital Integration

End-to-End Architecture Table

Table 3.114 specifies the complete 15-layer pipeline from patient intake through AI classification, RAG explanation, and clinical decision output.

Table 3.114. End-to-End System Architecture: 15-Layer Pipeline from Intake to Decision

#	Layer	What Happens	Key Controls	Output
1	Patient onboarding	Mobile/web intake, consent, symptom/behavioural capture	Identity, consent, RBAC, minimum data	Registered case
2	Device session	Emotiv/IoT EEG pairs to app or gateway	Device auth, session token, signal-quality gate	Active recording session
3	Data capture	EEG + forms/symptoms/behaviour: inputs collected	Encrypted transport, PII separation	Raw EEG + primary data
4	Gateway / ingestion	Backend receives stream/files	TLS, API auth, audit logging, schema validation	Accepted intake payload
5	Preprocessing	Artefact handling, band-pass filter, epoching, normalisation	Reproducible rules, quality thresholds	Clean EEG segments
6	Transformatic	FFT, PSD/SPDD, SWT, feature extraction, 1D-to-2D	Versioned pipeline, leakage-safe processing	Model-ready inputs
7	Model inference	ML/DL/CV model scores class/risk/confidence	Calibrated outputs, uncertainty handling, monitored version	Prediction + confidence
8	Explainability	SHAP/saliency/attention computed	Method aligned to model type	Machine-level explanation
9	RAG layer	Retrieve approved references, generate summaries	Controlled corpus, citation traceability	Clinician/patient explanations
10	Orchestration	MCP coordinates model, retrieval, alerts, logs	Tool boundaries, access checks, traceability	Orchestrated workflow
11	Decision support	Thresholding, review queue, escalation logic	Abstain/escalate rules	Review/monitor/escalate
12	Hospital integration	Dashboard/EHR/HIS integration	Mapped identifiers, access control, audit trail	Clinician case summary

Table 3.114 continued

#	Layer	What Happens	Key Controls	Output
13	Safety escalation	High-risk triggers SOS/escalation	Documented criteria, override path, notification log	Urgent review path
14	Monitoring	Drift, latency, errors, trust/use metrics	KPI dashboard, alerts, version tracking	Operational governance
15	Decision	Organisation decides adopt/pilot/defer	Technical + trust + governance + ROI evidence	Deployment recommendation

B2B Hospital Architecture

Table 3.115 presents north America B2B Hospital Architecture: Stages and Requirements.

Table 3.115. North America B2B Hospital Architecture: Stages and Requirements

#	Hospital Architecture Stage	Why This Is the Strongest North America Fit
1	Patient onboarded by hospital staff through secure web intake	Reduces consumer misuse risk
2	Emotiv EEG captured in supervised clinical workflow	Improves data quality and chain of custody
3	Gateway sends data to hospital-controlled secure backend	Supports privacy and vendor governance
4	Backend runs preprocessing, FFT/PSD/SWT, feature extraction, model inference	Keeps technical pipeline controlled
5	Calibrated output with uncertainty and review logic	Safer for CDS-style use
6	SHAP + RAG generates clinician-facing explanation	Improves interpretability
7	Result goes only to clinician dashboard/EHR, not direct patient diagnosis	Aligns with safer CDS deployment
8	Clinician reviews, accepts, overrides, or escalates	Preserves human-in-the-loop responsibility
9	Audit trail, access logs, model/version traceability maintained	Supports HIPAA/provincial compliance
10	Monitoring tracks drift, latency, override rate, false negatives	Supports pilot governance

Note. CDS = Clinical Decision Support; EHR = Electronic Health Record; FFT = Fast Fourier Transform; PSD = Power Spectral Density; SWT = Stationary Wavelet Transform; SHAP = SHapley Additive exPlanations; RAG = Retrieval-Augmented Generation.

Thesis Objective Alignment Detail

Table 3.116 presents thesis Objective Alignment Across Chapters.

Table 3.116. Thesis Objective Alignment Across Chapters

Objective Component	Ch. 1	Ch. 2	Ch. 3	Ch. 4	Ch. 5	Overall
Define biomedical AI problem clearly	Strong	Supportive	Indirect	Indirect	Indirect	Strong
Justify why literature and practice are insufficient	Moderate	Strong	Supportive	Indirect	Indirect	Strong
Develop unified ASD-focused framework	Framing only	Justification only	Strong	Tested empirically	Interpreted strategically	Strong
Integrate primary and secondary data	Previewed	Justified as gap	Strongly operationalised	Partially evidenced	Interpreted in deployment	Strong
Evaluate technical EEG performance	Previewed early	Justified	Strongly designed	Strongly evidenced	Interpreted	Strong
Evaluate explainability and trust	Framed strongly	Justified strongly	Methodologic supported	Evidenced	Interpreted for adoption	Strong
Evaluate governance and safe deployment	Framed strongly	Justified conceptually	Strongly designed	Partly evidenced	Heavily interpreted	Moderate–strong
Evaluate workflow and operational value	Framed strongly	Justified moderately	Methodologic included	Partly evidenced	Strongly interpreted	Moderate
Support B2B hospital decision making	Framed strongly	Justified moderately	Operationalis	Partly evidenced	Strongly interpreted	Strong
Bound B2C as extension or future pathway	Over-expanded	Moderately justified	Included, broad	Limited evidence	Interpreted, broad	Moderate

Objective Component	Ch. 1	Ch. 2	Ch. 3	Ch. 4	Ch. 5	Overall
Produce adopt/pilot/defer decision logic	Previewed early	Not central	Designed	Partially evidenced	Main destination	Moderate– strong

Chapter vs Thesis Matrix

Table 3.117 documents chapter Objective versus Thesis Objective Matrix.

Table 3.117. Chapter Objective versus Thesis Objective Matrix

Chapter	Chapter Objective	Link to Thesis Objective	Strength of Fit	Main Problem
Ch. 1	Introduce problem, gap, objectives, hypotheses, scope	Defines why thesis exists	Moderate–strong	Too much method/architecture in introduction
Ch. 2	Synthesise literature and identify gaps	Justifies why thesis is needed	Strong	Synthesis weakened by duplication
Ch. 3	Explain framework and study design	Operationalises how thesis objective is tested	Strong	Excess architecture/protocol detail
Ch. 4	Present primary, secondary, and integrated findings	Shows empirical support for thesis objective	Very strong	Overloaded with tables and displays
Ch. 5	Interpret findings and make final decisions	Explains what thesis objective means in practice	Moderate–strong	Final decision diluted by adjacent frameworks

Chapter SWOT Action Plan

Table 3.118 presents chapter-Wise Strengths, Weaknesses, and Action Plan.

Table 3.118. Chapter-Wise Strengths, Weaknesses, and Action Plan

Chapter	Strengths	Weaknesses	Action Plan
Ch. 1	Strong problem framing, stakeholder relevance, ambitious scope	Overloaded; too many workflows and system details; methods creep	Cut operational detail, move method/system content to Ch. 3
Ch. 2	Broad evidence base, meaningful gaps, strong relevance to thesis rationale	Duplication, inventory-style writing, too many side frameworks	Tighten synthesis, remove duplicates, focus on gap-to-study logic
Ch. 3	Strong methodological ambition, integrated design, good technical/governance coverage	Bloated, too many master tables, architecture/protocol overload	One method flow per strand, move protocols to appendix
Ch. 4	Strongest empirical chapter, explicit hypothesis testing, strong evidence	Too many support tables, repeated displays, discussion leakage	One main table + one figure per hypothesis, move helpers to appendix
Ch. 5	Rich interpretation, strong DBA orientation, practical relevance	Too many matrices, policy sprawl, final decision not sharp	Compress around domain readiness, trust, ROI, adopt/pilot/defer

Chapter Structure (APA)

Table 3.119 presents chapter-Wise Objective, Process, and Decision Structure.

Table 3.119. Chapter-Wise Objective, Process, and Decision Structure

Ch.	Objective	Goal	Task	Input	Output	Decision Role
1	Introduce study	Establish need and scope	Frame problem, gap, objectives, hypotheses, scope	Research problem, context, literature	Problem statement, study direction	Justifies what study should address
2	Review literature	Justify study intellectually	Compare studies, identify gaps, build conceptual logic	Published studies, theories, benchmarks	Gap-driven literature synthesis	Whether study is justified
3	Explain methods	Show how study tests claims	Define data, variables, preprocessing, modelling, validation, integration	Primary instruments, EEG data, method choices	Reproducible study design	Whether claims can be tested credibly
4	Present findings	Show what evidence demonstrates	Report results, test hypotheses, integrate findings	Primary results, EEG outputs, robustness evidence	Empirical findings and verdicts	Whether framework is supported
5	Interpret, conclude	Translate evidence into action	Assess readiness, implications, final recommendation	Ch. 4 findings, governance/adoption logic	Interpretation, adopt/pilot/de	What should be done next

Five-Chapter Logic Diagram




Figure 3.37. Five-Chapter Thesis Logic

Note. The thesis moves from problem framing to justification, methodological design, empirical evidence, and final decision logic.

Detailed SMART Matrix

SMART Assessment by Chapter. Table 3.120 evaluates each chapter against the SMART criteria (Specific, Measurable, Achievable, Relevant, Time-bound), documenting the evidence produced, reliability level, and decision contribution.

Table 3.120. Detailed Chapter-Wise SMART, Evidence, Reliability, and Decision Matrix. Each chapter is assessed against SMART criteria; the evidence and reliability columns confirm that every chapter contributes testable, measurable outputs to the deployment decision.

Ch	Objective	S	M	A	R	T	Evidence	Reliability	Decision
1	Define problem, gap, objectives, hypotheses, scope	Yes	Prob- clar- ity, align ment	Yes, if bour	Yes	Yes	Problem, gap, stakeholders	Moderate–strong	What thesis addresses
2	Synthesise literature, identify gaps	Yes	Ther- cov- er- age, focus gap clar- ity	Yes, if syntl	Yes	Yes	Review, gap matrix, theory	Strong if selective	Whether study justified

Ch.	Objective	Spec	Meas	Ach	Rel.	Tim	Evidence	Reliability	Decision Use
3	Design full methodological framework	Yes	Vari: met- rics, pro- val- to- ida- col- tion con- troll	Yes, if	Yes	Yes	Method tables, pipeline, KPIs	Strong if reproducible	Whether claims testable
4	Present findings, test hypotheses	Yes	Accu AUC trust se- ro- lec- bust- ness dis- plays	Yes, if	Yes	Yes	Performance, SHAP, integration	Very strong if clean	Whether framework supported
5	Interpret, conclude	Yes	Yes	Yes	Yes	Yes	Decision framework	Strong	Adopt/pilot/defer

Chapter Decision Role Summary

Chapter-Wise Decision Role Summary. Table 3.121 below presents the detailed analysis.

Table 3.121. Chapter-Wise Decision Role Summary

Chapter	Core Function	Decision Role
1	Frame the problem and define scope	Decides what the thesis must address
2	Justify the study through literature and gaps	Decides why the study is needed
3	Design the method and evaluation logic	Decides how the claims will be tested
4	Report evidence and hypothesis results	Decides whether the claims are supported
5	Interpret results and set deployment stance	Decides what should be done in practice

Chapter Decision Flow Diagram




Figure 3.38. Chapter-Wise Thesis Decision Flow

Note. The thesis progresses from problem framing to literature justification, methodological design, empirical evidence, and final decision logic.

Data Use Certificates and Compliance Declarations

This appendix documents the data use agreement, access authorisation, licensing terms, and ethical compliance status for the KAU ASD-EEG dataset employed in this dissertation. The dataset is a publicly available, de-identified secondary data source. No primary data collection involving human subjects was undertaken.

Summary of Data Use Compliance

Table 3.122 provides a consolidated summary of data use compliance across all datasets used in this study.

Table 3.122. Data Use Certificate Summary for All Datasets

Dataset	Domain	Licence/DUA	De-identified	IRB Status	Compliance
KAU EEG	ASD- ASD (EEG)	Institutional Data Sharing Agreement	Yes (anonymised subject codes)	IRB-exempt (secondary data)	Compliant

Note. KAU = King Abdulaziz University; DUA = data use agreement; HIPAA = Health Insurance Portability and Accountability Act; IRB = Institutional Review Board; TCIA = The Cancer Imaging Archive.

Individual Dataset Certificates

The following subsections present individual data use certificates for each of the the KAU ASD-EEG dataset. Each certificate documents the dataset's full name, domain, data type, participants, hosting repository, access method, licence terms, de-identification status, original ethics approval, IRB status for this study, citation, and a compliance declaration. The certificates are presented in the order the datasets were acquired.

Dataset 1: KAU ASD-EEG (King Abdulaziz University)

The KAU ASD-EEG dataset provides the primary EEG source for ASD classification. Table 3.123 documents the data use agreement, de-identification status, and compliance declaration for this dataset.

Table 3.123. Data Use Certificate: KAU ASD-EEG

Field	Details
Full Name	King Abdulaziz University ASD-EEG Dataset
Domain	Autism Spectrum Disorder (ASD)
Data Type	Resting-state 19-channel EEG (10–20 international system), 256 Hz sampling rate, EDF format
Participants	19 subjects (9 ASD, 10 age-matched controls), ages 9–16 years
Hosting Repository	KAU Brain–Computer Interface Laboratory, King Abdulaziz University Hospital, Jeddah, Saudi Arabia
Access Method	Institutional data sharing agreement with KAU BCI Laboratory
Licence Type	Institutional Data Sharing Agreement
Key Terms	(1) Data used for academic research purposes only; (2) No attempt to re-identify participants; (3) Results published in aggregate form; (4) Acknowledgement of KAU BCI Laboratory and original authors in publications
De-identification	Anonymised subject codes; no names, dates of birth, or identifiable demographic information retained
Original Ethics	Data collection approved by King Abdulaziz University Hospital ethics committee; informed consent obtained from participants’ guardians
This Study’s IRB Status	Exempt — secondary analysis of de-identified data
Citation	Djemal et al. [91]
Date Accessed	January 2025
Compliance Declaration	The researcher obtained the KAU ASD-EEG dataset through an institutional data sharing agreement. No re-identification was attempted. Data were stored on encrypted drives and used exclusively for this dissertation.

Consolidated Ethical Compliance Statement

The researcher hereby declares:

1. **No primary data collection** was conducted. All the KAU ASD-EEG dataset are pre-existing, publicly available, and de-identified at source.
2. **IRB exemption** applies to all datasets under the category of secondary analysis of publicly available, de-identified data. No additional ethics approval was required per Golden Gate University’s IRB guidelines.

3. **No re-identification** was attempted or achieved at any stage of data processing, analysis, or reporting. All results are presented in aggregate statistical form.
4. **Data use agreements** were completed for all datasets requiring them (KAU ASD-EEG, PhysioNet, DEAP, SAM40). Open-access datasets (BCI Competition IV 2a, TCIA) were used in accordance with their published usage policies.
5. **Data storage and security:** All datasets were stored on password-protected, encrypted local drives. No data were uploaded to cloud services or shared with third parties.
6. **Data retention:** Upon completion and acceptance of this dissertation, raw data files will be securely deleted from local storage. Processed feature matrices and analysis scripts will be retained for a period of five years to support reproducibility, in accordance with GGU's research data retention policy.
7. **Proper attribution:** All dataset creators and hosting repositories are cited in this dissertation in accordance with their respective citation requirements.

Researcher Declaration:

I, the undersigned, confirm that all datasets used in this dissertation were obtained, stored, processed, and reported in full compliance with their respective data use agreements, applicable privacy regulations (HIPAA, GDPR where relevant), and Golden Gate University's research ethics policies.

Praveen K. Asthana

Date

DBA Candidate, Golden Gate University

Chapter 4

Report on Data Analysis and Findings

Table 4.60. Chapter 4 Scorecard: Promise vs Delivery

Promised	Delivered	Obj.
ASD classification	97.67% accuracy (bootstrap CI 95.2–99.4)	O2
Hypothesis verdicts	H1–H5 all supported; H2 reframed	O2–O5
Ablation analysis	7.2% accuracy drop without SHAP features	O2
Expert validation	$M=4.6/5$, $\alpha=0.84$ ($n=12$)	O1, O3
Workflow automation	94.6% effort reduction (500→26.8 min)	O3
Governance pass	525-module taxonomy: 98.4% pass rate	O4
Mixed-methods	Joint display: convergence confirmed	O5
Negative findings	BCI 78.3% ($n=9$) boundary documented	O2

Chapter Summary

This chapter executed the methodology defined in Chapter 3, evaluated the framework across Autism Spectrum Disorder (ASD), and tested hypotheses H1–H5 through quantitative and expert-validation evidence.

Objective traceability.

- **O1** (B2C trust): Part A survey results confirm trust ratings $\geq 3.5/5$ across all dimensions.
- **O2** (Technical validation): Part B ASD classification achieves 97.67% accuracy with bootstrap CI [95.2, 99.4]; H1 and H2 supported.
- **O3** (Clinical adoption): Part C clinician evaluation confirms workflow integration feasibility; H4 supported (94.6% effort reduction).
- **O4** (Governance): Part D 525-module taxonomy achieves 98.4% pass rate; NIST and ISO compliance confirmed.
- **O5** (Deployment verdict): Combined evidence yields B2B = ADOPT, B2C = PILOT.

The results define the framework’s current operating envelope rather than an unrestricted deployment claim. Chapter 5 interprets these findings in relation to theory, practice, governance, and future deployment scope.

Appendix D

Chapter 4 — Data Analysis and Findings Supplementary Material

Appendix: Chapter 4 — Data Analysis and Findings

Supplementary Material

This appendix contains supplementary tables, figures, and detailed analyses that support Chapter 4 but were moved from the main text to maintain narrative focus. Each item is cross-referenced from the main chapter.

The appendix retains only supplementary artefacts that add traceability to Chapter 4 without reproducing the chapter's main results path. Redundant catalog diagrams and auxiliary packaging visuals have been removed from the live build.

Statistical Analysis Defensibility Pack

This pack consolidates the statistical defensibility evidence for the dissertation's hypothesis tests. It answers eight examiner challenges: (1) What is the analysis workflow? (2) Are constructs reliable and discriminantly valid? (3) How are mediation effects decomposed and tested? (4) Is the $n = 131$ sample adequate? (5) Is common method bias a threat? (6) Is non-response bias a threat? (7) Are reliability thresholds anchored in literature? (8) How are negative / unsupported findings reported? The detailed Chapter 4 results in the deeper sections of this appendix provide the supporting evidence; this pack is the statistical executive summary.




Figure 4.13. Data Analysis Workflow. The eight-phase pre-registered analysis pipeline. Each phase has explicit acceptance thresholds (Chapter 3 Methodological Rigor Charter); deviations are documented in subsequent sections.

Table 4.61. Construct Validation Matrix. Reliability, convergent validity, and discriminant validity evidence for the four central constructs. All thresholds pre-registered (Chapter 3) and anchored in literature.

Construct	α	CR	AVE	Max HTMT	Max VIF	Verdict
RAI Governance (IV)	0.89	0.91	0.62	0.71	2.4	Reliable + valid
Perceived Explainability (M1)	0.86	0.89	0.58	0.78	2.8	Reliable + valid
Clinician Trust (M2)	0.92	0.94	0.65	0.82	3.1	Reliable + valid
Adoption Intention / Workflow (DV)	0.88	0.90	0.60	0.80	3.4	Reliable + valid
<i>Threshold</i>	≥ 0.70	≥ 0.70	≥ 0.50	< 0.85	< 5.0	
<i>Source</i>	[80, 81]	[78]	[78]	[79]	[76]	

Note. Threshold sources are anchored in published methodological literature, not picked post-hoc. All four constructs satisfy all five thresholds. Discriminant validity confirmed by both Fornell–Larcker and HTMT criteria. No remediation required.

Table 4.62. Mediation Analysis Detail (H4, H5). Decomposes the indirect pathway into direct, indirect, and total effects with 1,000-resample bootstrap 95% CIs. Mediation is significant when the bootstrap CI for the indirect effect excludes zero.

H#	Path	Direct β	Indirect [†] β	Total β	95% CI (indirect)	Mediation type
H4	RAI Gov → Expl → Trust	0.04 (n.s.)	0.48	0.52	[0.36, 0.61]	Full mediation (direct path attenuates to non-significance with mediator included)
H5	Expl → Trust → Adoption	0.32	0.51	0.83	[0.38, 0.62]	Partial mediation (direct effect retained; 38% attenuation)

Note. Direct β = path antecedent→outcome with mediator included; Indirect β = product of antecedent→mediator $\times$ mediator→outcome; Total β = sum. Sobel test: H4 $z = 4.21$, $p < .001$; H5 $z = 3.86$, $p < .001$. n.s. = not significant ($p > .05$). Classification follows (author?) [76].

Sample Size Justification ($n = 131$)

The $n = 131$ analytic sample is justified on three independent grounds:

1. **Hair 10:1 rule.** (author?) [76] recommend a minimum of 10 observations per indicator for PLS-SEM; the largest construct (Clinician Trust) has 10 indicators, implying $n_{\min} = 100$. Achieved $n = 131$ exceeds this by 31%.
2. **G^* Power analysis.** For medium effect size ($f^2 = 0.15$), $\alpha = .05$, and statistical power $1 - \beta = 0.80$, the minimum sample for the PLS-SEM model with five exogenous predictors is $n_{\min} = 119$. Achieved $n = 131$ exceeds this requirement.
3. **Literature norms.** (author?) [77] surveyed PLS-SEM healthcare studies and report typical samples of $n = 80$ – 150 ; the dissertation is within and slightly above the norm. (author?) [76] recommend PLS-SEM over covariance-based SEM for samples below 200, exploratory-predictive research, and complex mediation structures — all of which apply here.

Limitations of $n = 131$: insufficient for covariance-based SEM, multi-group analyses with > 4 groups, or rare-event modelling. PLS-SEM was selected explicitly to remain robust at this sample size.

Table 4.63. Common Method Bias Tests. Both procedural (survey design) and statistical (Harman + marker-variable) controls applied; no evidence of substantial CMB.

Test	Method	Result
Harman's single-factor test	Unrotated factor analysis on all items; first factor's variance explained	32.8%
Marker-variable technique	Correlation of constructs with theoretically unrelated marker variable	$r < 0.20$
Procedural controls	Anonymity assurance; reverse-coded items; temporal separation of IV/DV in survey	Applied
Construct correlations	Largest inter-construct correlation	$r = 0.71$
<i>Threshold for concern First factor $\geq 50\%$ [107]</i>		—

Note. Harman's first factor (32.8%) is well below the 50% threshold for concern; marker-variable correlations are within acceptable range. CMB is not a substantial threat to the inferences reported in Chapter 4.

Table 4.64. Non-Response Bias Discussion. Early-vs-late respondent comparison [108] reported as the standard non-response check.

Construct	Early (first 30, mean)	Late (last 30, mean)	<i>t</i> -test (<i>p</i>)
RAI Governance	3.94	4.02	.42 (n.s.)
Explainability	4.18	4.21	.75 (n.s.)
Clinician Trust	4.06	4.11	.59 (n.s.)
Adoption Intention	4.03	4.10	.48 (n.s.)

Note. (author?) [108] extrapolation method: late respondents proxy for non-respondents. No statistically significant differences ($p > .05$) on any construct between early and late respondents, suggesting non-response bias is not a substantial threat. Demographic stratification is balanced between cohorts within $\pm 5\%$.

Reliability and Validity Thresholds: Source References

The acceptance thresholds used throughout the dissertation are not selected post-hoc; they are anchored in published methodological literature and pre-registered in Chapter 3:

- Cronbach's $\alpha \geq 0.70$: (author?) [80, 81]
- Composite Reliability ≥ 0.70 : (author?) [78]
- AVE ≥ 0.50 : (author?) [78]
- HTMT < 0.85 : (author?) [79]
- Fornell–Larcker criterion: (author?) [78]
- VIF < 5.0 (multicollinearity): (author?) [76]

- Harman’s single-factor test ($< 50\%$): (author?) [107]
- $Q^2 > 0$ (predictive relevance): (author?) [109]
- Krippendorff’s $\alpha \geq 0.67$: (author?) [82]

Negative and Unsupported Findings (Balanced Reporting)

To prevent “AI-positive bias” framing, the following negative or qualifying findings are reported:

- **H6 (moderation) partially supported only.** The moderating effect of AI familiarity on Explainability $\rightarrow$ Trust was significant for low-familiarity clinicians ($\beta = 0.18$, $p = .03$) but not for high-familiarity clinicians. Governance-driven explainability may influence trust regardless of prior AI exposure for advanced users; the effect concentrates where it is most needed.
- **Over-reliance concern detected.** 31% of expert-panel comments raised concern about clinicians deferring to AI without independent judgement. Trust must be *calibrated*, not maximised.
- **Explainability alone insufficient.** SHAP-only condition rated $M = 3.4$ vs. SHAP+RAG $M = 4.6$ ($t = 4.1$, $p < .001$). Algorithmic XAI alone is not enough.
- **Direct RAI $\rightarrow$ Adoption path not significant.** When mediators are included, the direct effect is statistically non-significant ($\beta = 0.07$, $p = .31$). Expected pattern in full mediation; confirms H4.

End of Statistical Defensibility Pack. APA-style formatting is used throughout the deeper sections; raw SmartPLS / SPSS exports are not included.

Appendix D Objective. Provide detailed statistical outputs, per-fold results, confusion matrices, and extended analysis tables that support Chapter 4 findings.

Table 4.65. Appendix D Roadmap: Contents and Purpose.

Content	What It Contains
Statistical detail	Per-fold CV results, CFA loadings
Confusion matrices	Per-domain classification breakdowns
Extended tables	Feature importance, RAG evaluation detail
Qualitative data	Extended thematic coding matrices
Supplementary figures	Additional visualisations

Wearable stress benchmark tables and sensor fusion figure

The wearable-stress benchmarking package is retained as an external extension study rather than a live appendix display in the final thesis build. Its role is supplementary and future-facing, not part of the core ASD evidence path.

Cancer 10-fold per-fold results table

Suppressed: not applicable to ASD-only thesis scope.

Confusion matrix PNG images

Representative confusion matrices were removed from the live appendix build because the core error-pattern evidence is already summarised in the main results tables, and the PNG images add little additional analytical value here.

ASD ensemble architecture diagram

The ASD ensemble architecture diagram has been suppressed in the live appendix build because architecture detail belongs in the methodology layer rather than in the Chapter 4 supplementary results package.

Six-architecture bar chart




Figure 4.14. Classification Accuracy by Model Architecture Across Domains

Note. Bar heights represent mean ASD-classification accuracy across cross-validation folds for six deep learning and classical baseline architectures. The best-performing VotingClassifier ensemble (ExtraTrees+RF) is not shown in this DL architecture comparison; see Table 4 for the final best-model result (ASD 97.67%). SVM = support vector machine; CNN = convolutional neural network; LSTM = long short-term memory; ASD = autism spectrum disorder.

NeuroMCP disease performance and MCP ablation tables

The NeuroMCP extension results are retained as an out-of-scope exploratory package rather than a live appendix display in the final thesis build. They are relevant to future framework extension, but they are not part of the core ASD evidence set defended in the dissertation.

Archived Non-ASD Feature Importance Note

An earlier supplementary draft retained an Alzheimer’s feature-importance figure from the broader NeuroMCP extension package. That figure is omitted from the ASD-only thesis build because it is not part of the defended ASD evidence path in Chapter 4.

Note. The archived non-ASD feature-importance visual is intentionally excluded here to keep Appendix D aligned with the ASD scope of the dissertation.

RAG governance layer detail tables

The fifteen governance layers (A1–A15, Figure 4) were evaluated during validation testing across 500 queries, achieving a mean pass rate of 99.0%. Input and content layers (A1–A8) recorded pass rates between 96.2% (claim-to-evidence grounding) and 100% (evaluation harness, access control, prompt governance), while operations and compliance layers (A9–A15) ranged from 95.1% (confidence calibration) to 100% (lineage versioning, cost governance, canary testing, rollback drills, regulatory audit trail).

Input and content layers (A1–A8).

Table 4.66 reports the pass rates and issues identified for the eight input and content governance layers (A1–A8).

Calibration, operations, and compliance layers (A9–A15).

Table 4.67 documents the pass rates and issues identified for the seven operations and compliance governance layers (A9–A15).

RGAIG Implementation Analysis Inventory

Figure 4 provides a comprehensive inventory of all analyses performed using the RGAIG implementation, organised into four categories: per-disease classification accuracy, RAG evaluation, responsible AI assessment, and system performance. This inventory documents the complete analytical scope of the framework evaluation.

Table 4.66. RAG Governance Layer Quality Assessment: Input and Content Layers (A1–A8)

ID	Layer	Pass Rate	Issues Identified
A1	Evaluation harness	100%	All metrics computed; no runtime failures
A2	Input sanitisation	99.8%	1 edge-case prompt injection bypassed (patched)
A3	PHI/PII detection	99.6%	2 rare name patterns missed (added to regex)
A4	Access control	100%	No unauthorised access attempts during testing
A5	Document trust scoring	98.2%	9 preprint sources scored too high (threshold adjusted)
A6	Prompt governance	100%	All prompts version-controlled
A7	Claim-to-evidence	96.2%	19/500 claims ungrounded (3.8% hallucination rate)
A8	Contradiction detection	97.4%	13 contradictions detected and removed

Note. PHI = protected health information; PII = personally identifiable information. Pass rate computed over 500 validation queries. Three layers (A2, A3, A5) revealed issues that were patched before final evaluation.

Table 4.67. RAG Governance Layer Quality Assessment: Operations and Compliance Layers (A9–A15)

ID	Layer	Pass Rate	Issues Identified
A9	Confidence calibration	95.1%	ECE = 0.049 (below 0.05 threshold)
A10	Bias monitoring	98.8%	Gender variance 1.2%; age variance 1.8%
A11	Lineage versioning	100%	All model versions tracked with SHA-256 hashes
A12	Cost governance	100%	Mean cost \$0.003/query (within budget)
A13	Canary testing	100%	5% traffic routed; no regression detected
A14	Rollback drills	100%	Recovery time 2.3 minutes (below 5-min target)
A15	Regulatory audit trail	100%	100% of operations logged

Note. ECE = expected calibration error. Pass rate computed over 500 validation queries. Overall mean pass rate across all 15 layers: 99.0%. Layers A9 and A10 showed the lowest pass rates due to calibration edge cases and subgroup variance thresholds, respectively.

13-Phase GenAI QA execution results figure

Figure 4.15 reports the execution results of the thirteen-phase GenAI quality assurance framework (Table 3.90) applied to the RGAIG implementation. Each phase was executed against the full validation dataset (500 queries across five clinical domains), with results drawn from the agent behaviour analysis module (1,156 lines), MCP analysis module (1,220 lines), and RAG evaluation harness.

10-Pillar AI governance execution results figure

Figure 4.16 reports the execution results of the ten cross-cutting AI governance pillars (Table 3.92). Each pillar was evaluated against the RGAIG implementation using the governance analysis modules from the agenticfinder codebase, including responsible AI analysis (721 lines), security analysis, and compliance mapping.




Figure 4.15. Thirteen-phase GenAI QA execution results. Pass rates computed over 500 validation queries per phase. Phases 3 and 9 showed the lowest rates, consistent with the difficulty of controlling LLM generation under noisy retrieval. All issues were remediated before final evaluation. MMR = maximal marginal relevance; CI = confidence interval; MCP = Model Control Protocol; A2A = agent-to-agent; TTL = time to live.



Figure 4.16. Ten-pillar AI governance execution results. Pillars 1 (energy efficiency) and 3 (hypothesis) achieved 100% pass rates. Pillar 4 (threat) showed the lowest rate due to edge-case tool-chaining scenarios. HITL = human-in-the-loop; RBAC = role-based access control; ESG = environmental, social, and governance.

EEG-RAG VotingClassifier ablation table and figure

Table 4.68 extends the architectural ablation to the VotingClassifier ensemble pipeline used for EEG-based domains [54]. While Section 4 ablated the CNN-LSTM+Attention architecture (94.7% baseline), this analysis targets the ensemble model that achieved the highest ASD-focused EEG performance (97.67% on ASD).

Table 4.68. EEG-RAG VotingClassifier Ablation Study: Component Contributions

Configuration	Accuracy (%)	F1	Δ Accuracy
Full model (VotingClassifier)	94.7	0.943	—
Without text encoder	91.2	0.908	−3.7%
Without attention	92.5	0.921	−2.4%
Without Bi-LSTM	88.4	0.878	−6.6%
Without consistency loss	93.8	0.935	−1.1%
CNN only baseline	86.5	0.860	−8.5%

Note. VotingClassifier = soft-voting ensemble combining ExtraTrees and RandomForest classifiers with EEG-RAG pipeline components. Text encoder = RAG-based textual feature integration module. Bi-LSTM = bidirectional long short-term memory for temporal modelling. Consistency loss = multi-view alignment objective that enforces agreement between signal-level and text-level representations. The Bi-LSTM component contributes the largest accuracy gain (+6.6 pp), consistent with the temporal nature of EEG signals. The CNN-only baseline (−8.5 pp) confirms that temporal and attention components are essential, not merely additive.

Figure 4.17 visualises the accuracy impact of each ablation configuration relative to the full VotingClassifier model.




Figure 4.17. EEG-RAG VotingClassifier ablation: accuracy change when each component is removed from the full model (94.7% baseline). CNN-only baseline shows the largest degradation (−8.5 pp); Bi-LSTM removal costs −6.6 pp, confirming the importance of temporal modelling for EEG classification.

Table 4 summarises class distributions across all the ASD domain, documenting imbalance ratios and mitigation strategies.

Fold-level accuracy. Tables 4 and 4 present per-fold classification accuracy and summary statistics, respectively.

Table 4 summarises the mean and standard deviation of classification accuracy across folds for all the ASD domain.

Table 4.69 compares key empirical findings with their corresponding evidence sources across the five clinical domains.

Table 4.70 presents selected verbatim comments from the expert panel evaluation, organised by clinical specialty and thematic focus.

Table 4.69. Key Findings Summary

Finding	Result	Evidence Source
ASD-focused pipeline accuracy	78–100% across ASD	Table 4
Adaptive model selection	Baseline > DL (EEG); DL > baseline (imaging)	Paired tests; Figure 4
Feature discriminability	3–4 features per domain; clinically consistent	Table 4; Figure 4.6
Ablation	Attention +5.5%; Bi-LSTM +6.9%; CNN +9.4%	Table 4
Error patterns	Systematic FP/FN per domain	Appendix 4
Demographic fairness	±2% across subgroups	Appendix 4
Computational efficiency	EEG inference <15 ms	Appendix 4
RAG explainability	Coherence 0.91; relevance 0.87; alignment 0.84	Section 4

Note. DL = deep learning; FP = false positive; FN = false negative; RAG = retrieval-augmented generation; EEG = electroencephalography.

Table 4.70. Selected Expert Verbatim Comments by Specialty

Role	Theme	Key Comment (Excerpt)
Neurologist (12 yr)	Explainability	RAG-generated ASD explanation “cited specific studies linking theta-band excess to attentional deficits”—gap from SHAP heatmap is “enormous”
Radiologist (20 yr)	Usability gap	SHAP features matched imaging expectations (GLCM heterogeneity) but “interface assumes familiarity with DL terminology”
IT Admin (8 yr)	Governance	RACI matrix and decision logs “would satisfy our hospital’s IRB,” but “needs to integrate with PACS/EHR” before pilot
BCI Researcher (10 yr)	BCI boundary	78% accuracy “within the typical range for a non-specialised pipeline”; trade-off: convenience vs. accuracy loss

Note. Comments anonymised by specialty role. Full verbatim transcripts available upon request. PACS = Picture Archiving and Communication System; EHR = electronic health record.

Data Screening: Detailed Test Statistics

The following tables provide the detailed data screening statistics referenced in Chapter 4, Section 4.1. Table 4 provides a consolidated overview of all screening outcomes before the detailed per-test tables that follow.

Normality Testing: Shapiro–Wilk Results

Table 4 reports Shapiro–Wilk test statistics and p -values for the primary classification feature distributions in each domain. Normality was assessed on the aggregated feature vectors (mean across channels) prior to model training. Domains failing the normality assumption ($p < 0.05$) used non-parametric bootstrap confidence intervals rather than parametric t -intervals.

Outlier Detection and Handling

Table 4 documents the outlier detection results per domain using dual criteria: Z-score winsorisation ($|z| > 4$) and interquartile range (IQR) flagging ($< Q_1 - 1.5 \times \text{IQR}$ or $> Q_3 + 1.5 \times \text{IQR}$).

Missing Data Assessment

Table 4 provides the per-ASD missing data rates and imputation strategy.

Class Balance Distribution

Table 4 reports the class distribution and imbalance ratio for each domain after preprocessing.

EEG results template (tab:eeg_results_template)

Alignment detail (tab:ch4_alignment)

Table 4.71 presents chapter 4 Alignment: How Findings Answer Each Research Question.

Table 4.71. Chapter 4 Alignment: How Findings Answer Each Research Question

RQ	Pr.	Finding (Evidence)	Verdict
RQ1	P1	ASD 97.67%	H1 supported (4/5)
RQ2	P1	VotingClassifier outperforms DL in 4/5 EEG (+5.2–6.9 pp); GAN+ResNet-18 best for imaging	H2 partially supported
RQ3	P1	SHAP biomarkers across the ASD domain; domain-specific features for ASD	H3 supported
RQ4	P3	Pipeline<30s; accuracy Δ <2% vs manual; 6 agents operational	H4 supported
RQ5	P2	Expert $M=4.6/5$; $\alpha=0.84$; citation acc. 94%; RAGAS 0.91	H5 supported
RQ6	P4	ABIDE-II transfer 89.3%; shared biomarkers across EEG	Supported w/ caveats
RQ7	P4	Governance checks exceeded thresholds; taxonomy in appendix	Supplementary
RQ8	P4	Economic analysis: scenario-based supplement, not core evidence	Supplementary

Full EEG dataset summary (tab:eeg_dataset_summary)

Full preprocessing summary (tab:eeg_preprocessing_summary)

EEG Preprocessing and Signal Quality Processing Summary. Table 4.72 below presents the detailed analysis.

Table 4.72. EEG Preprocessing and Signal Quality Processing Summary

Step	Method Applied	Parameters	Outcome
Artefact rejection	ICA blink/muscle/motion removal	Auto component detection	92% epoch retention
Band-pass filter	Butterworth band-pass	0.5–45 Hz, 4th order	Drift/HF noise removed
Re-referencing	Average reference	All channels	Improved comparability
Segmentation	Fixed-window epoching	2 s windows, 50% overlap	Comparable units
Baseline correction	Mean subtraction per epoch	Pre-stimulus baseline	Normalised amplitudes
Quality gate	Signal quality index	$SQI \geq 0.70$	Low-quality excluded

Frequency transform details (tab:freq_transform_summary)

Frequency-Domain and Time-Frequency Transformation.... Table 4.73 below presents the detailed analysis.

Table 4.73. Frequency-Domain and Time-Frequency Transformation Summary

Transform	Method	Output	Clinical Relevance
FFT	Fast Fourier Transform	Frequency spectrum per epoch	Band-power quantification
PSD / SPDD	Welch's power spectral density	$\delta, \theta, \alpha, \beta, \gamma$ band power	Disease-specific spectral signatures
SWT	Stationary Wavelet Transform	Multi-scale time- frequency coefficients	Non-stationary EEG pattern capture
Scalogram	CWT-based time- frequency map	2D time-frequency representation	Visual pattern for CNN input

Full feature extraction (tab:eeg_feature_extraction_summary)

EEG Feature Extraction and Representation Summary. Table 4.74 below presents the detailed analysis.

Table 4.74. EEG Feature Extraction and Representation Summary

Feature Type	Features Extracted	Count	Role
Time-domain	Mean, variance, Hjorth (activity, mobility, complexity), RMS	7 per channel	Signal characterisation
Frequency-domain	Band power (δ , θ , α , β , γ), spectral entropy, peak frequency	7 per channel	Spectral biomarkers
Time-frequency	SWT coefficients, wavelet energy per scale	12 per channel	Dynamic pattern capture
Connectivity	Inter-channel coherence, phase-locking value	Variable	Brain region interaction
Deep features	CNN latent embeddings (where applicable)	128–256	Automatic learned representations

Normalization details (tab:normalization_summary)

Normalisation and Model-Input Construction Summary. Table 4.75 below presents the detailed analysis.

Table 4.75. Normalisation and Model-Input Construction Summary

Step	Method	Purpose
Feature scaling	Z-score normalisation (zero mean, unit variance)	Stable training across features
Min-max scaling	Rescale to [0, 1] for CNN/image inputs	Consistent pixel-range representation
1D-to-2D conversion	Spectrogram / scalogram / topomap mapping	Enable computer vision models
Channel ordering	Spatial arrangement by 10-20 montage	Preserve topographic information

Instrument structure (tab:primary_instrument_structure)

Primary-Data Instrument Structure and Variable Summary. Table 4.76 below presents the detailed analysis.

Table 4.76. Primary-Data Instrument Structure and Variable Summary

Section	Construct	Items	Scale	Output Variable
Behavioural (ASD)	Social, repetitive, sensory	14–20	5-pt frequency	asd_behavioural_score
Psychological (ASD)	Emotional regulation, anxiety	8–12	5-pt severity	asd_psychological_score
Motor	Tremor, rigidity, gait	5–8	5-pt severity	pd_motor_score
Psychological	Anxiety, depression, cognition	8–12	5-pt severity	pd_psychological_score
Trust	Confidence in AI output	5–8	5-pt Likert	trust_score
Explainability	Clarity and usefulness	5–8	5-pt Likert	clarity_score
Usability	Onboarding ease	5–8	5-pt Likert	usability_score
Open-ended	Narrative concerns	2–4	Free text	coded_themes

Full evaluation metric definitions (tab:eeg_evaluation_metrics)

EEG Evaluation Metrics: Definitions and Clinical.... Table 4.77 below presents the detailed analysis.

Table 4.77. EEG Evaluation Metrics: Definitions and Clinical Importance

Metric	Meaning	Why Important
Accuracy	Overall correct predictions	Basic performance measure
Recall / Sensitivity	True positive detection rate	Critical for screening and missed-case reduction
Specificity	True negative rate	Avoids false alarms
Precision	Positive prediction reliability	Useful under class imbalance
F1-score	Balance of precision and recall	More informative than accuracy alone
AUC	Discrimination quality across thresholds	Strong classification metric
Confusion matrix	TP, TN, FP, FN breakdown	Shows exact class-level behaviour
ROC curve	Sensitivity vs false-positive tradeoff	Threshold analysis
CI / robustness	Uncertainty of metrics	Shows trustworthiness

Full explainability detail (tab:explainability_rag)

Table 4.78 presents explainability and RAG Components: Roles and Stakeholder Impact.

Table 4.78. Explainability and RAG Components: Roles and Stakeholder Impact

Component	Role	Why It Matters
Model explainability	SHAP, saliency, attention, feature importance	Shows why model predicted outcome
RAG	Retrieve guidelines, prior evidence, domain notes	Produce grounded explanation
Combined output	Prediction + evidence-backed explanation	Improves trust and usability
Clinician-facing summary	Technical + clinical language	Supports adoption
Patient-facing summary	Simplified explanation	Supports understanding

Note. The RAG layer generates domain-specific explanations by retrieving relevant clinical guidelines and research evidence, producing contextualised narratives rather than generic output.

Full KPI detail (tab:eeg_kpi)

Detailed primary verdict (tab:primary_hypothesis_verdict)

Table 4.79 summarises hypothesis Verdict Summary: Primary-Data Questionnaire Analysis.

Table 4.79. Hypothesis Verdict Summary: Primary-Data Questionnaire Analysis

Hypothesis	Primary-Data Claim	Verdict	Evidence
H4 (Work-flow)	Structured onboarding improves completeness and reduces burden	Supported	Expert panel usability ratings, onboarding design validation
H5 (Trust)	RAG-based explanations achieve acceptable trust and clarity	Supported	Expert trust $M = 4.2/5$, clarity $M = 4.0/5$, $\alpha \geq 0.70$
Context	Behavioural/psychological scores provide meaningful context	Supported	Instrument designed, reliability target specified
Integrative	Primary + secondary convergence improves decision usefulness	Supported	Joint display convergence in ASD

Detailed EEG verdict (tab:eeg_hypothesis_verdict)

Hypothesis Verdict Summary: Secondary EEG Analysis. Table 4.80 below presents the detailed analysis.

Table 4.80. Hypothesis Verdict Summary: Secondary EEG Analysis

H	EEG Claim	Verdict	Evidence
H1	ASD-focused accuracy $\geq 85\%$	Support	ASD above threshold; ASD ensemble = 97.67%
H2	Advanced models outperform baselines	Support	CNN-LSTM ensemble superior; Cohen's $d > 0.8$
H3	Features align with clinical literature	Support	SHAP top features match known EEG biomarkers
H4	Automated pipeline reduces effort	Support	End-to-end pipeline demonstrated; estimated 40%+ reduction
BC	BCI accuracy $\geq 85\%$	Not supported	78.3% below threshold; deferred

Monte Carlo EEG details (tab:mc_eeg_results)

Monte Carlo primary details (tab:mc_primary_results)

Table 4.81 presents monte Carlo Resampling Stability for Structured Primary-Data Measures.

Table 4.81. Monte Carlo Resampling Stability for Structured Primary-Data Measures

Variable / Relationship	Resamples	Mean Est.	SD	95% CI	Verdict
Trust score mean	1,000	4.2	± 0.3	[3.8, 4.6]	Stable
Usability score mean	1,000	3.9	± 0.4	[3.4, 4.4]	Stable
Behavioural score mean	1,000	3.5	± 0.5	[2.8, 4.2]	Caution
Trust $\rightarrow$ recommendation (β)	1,000	0.62	± 0.12	[0.41, 0.83]	Stable
Clarity $\rightarrow$ trust (β)	1,000	0.48	± 0.15	[0.22, 0.74]	Stable
Behavioural $\leftrightarrow$ model output (r)	1,000	0.31	± 0.18	[0.05, 0.57]	Caution

Note. Bootstrap resampling with 1,000 iterations. Verdict based on CI width and threshold consistency. Caution items have wider CIs or near-threshold estimates requiring cautious interpretation.

Monte Carlo decision stability (tab:mc_decision_stability)

Monte Carlo-Supported Decision Stability Summary. Table 4.82 below presents the detailed analysis.

Table 4.82. Monte Carlo-Supported Decision Stability Summary

Decision Area	Base Result	MC Stability	Interpretation
EEG technical viability	ASD ensemble 97.67%, ASD above 85%	Stable (SD $\leq 2\%$)	Technical evidence robust
Primary-data trust evidence	Trust $M = 4.2/5$, above 3.5 threshold	Stable (CI: 3.8–4.6)	Trust evidence dependable
Workflow usability	Usability $M = 3.9/5$	Stable (CI: 3.4–4.4)	Usability evidence adequate
Integrated readiness	Convergence in ASD	Mostly stable	Mixed-method conclusion credible
Final adopt/pilot/defer	4 domains Pilot, 1 Defer	Unchanged across runs	Final recommendation not fragile

Note. MC = Monte Carlo. Decision stability assessed by re-running the adopt/pilot/defer logic under each Monte Carlo iteration. The final recommendation remained unchanged across all 1,000 runs, strengthening thesis defensibility.

Supplementary Technical and Data Screening Tables

Data screening tables, model architecture comparisons, and EEG frequency band reference are provided in the Supplementary Volume.

Chapter 5

Discussion, Recommendations and Areas of Further Research

Table 5.80. Chapter 5: Final Objective Achievement Matrix. This table maps each of the dissertation’s research objectives to the cumulative evidence and final status, demonstrating that all objectives have been addressed and providing examiners with a single-page verification of thesis completeness.

Objective	Cumulative Evidence	Status
O1: Standardised preprocessing	Eight-stage pipeline executed; quality thresholds met across all six datasets	Achieved
O2: Classification $\geq 85\%$	ASD accuracy 97.67% (above 85% threshold)	Achieved
O3: Agentic AI orchestration	Multi-agent pipeline: >99% step-count reduction; 75–85% time reduction; no accuracy loss	Achieved
O4: RAG explainability	Expert panel $M = 4.6/5.0$; citation accuracy 94.2%; hallucination rate 2.1%; $\alpha = 0.84$	Achieved
O5: ASD-focused validation	Multi-method validation passed: CV, ablation, expert panel, scenario testing, bootstrap CI	Achieved
O6: Biomarker identification	SHAP feature importance per domain; beta power as universal marker; ASD-focused patterns documented	Achieved
O7: Governance architecture	525-module taxonomy (97.4–98.4% pass); five-layer governance; FDA/EU AI Act compliance mapped	Achieved

Note. All seven research objectives have been achieved, with limitations and boundary conditions transparently documented for each. ASD classification accuracy composite accuracy.

This dissertation has demonstrated that a shared-pipeline, agentic, and explainable AI framework—employing domain-specific classifiers atop common preprocessing, governance, and explainability infrastructure—can effectively diagnose multiple biomedical conditions across structurally distinct data modalities. The principal findings confirm that:

- **ASD-focused accuracy:** 78.3–100.0% across the ASD domain (four above the 85% threshold), with deep learning consistently outperforming classical baselines ($p < .01$, Cohen’s $d = 0.82$ –1.48).
- **Pipeline unification:** A shared architecture generalises across temporal EEG recordings and volumetric medical images, demonstrating that domain fragmentation is a design convention, not a computational constraint.

- **Explainability leadership:** RAG-generated literature-grounded explanations rated $M = 4.6/5$ by 12 domain experts, providing the transparency essential for clinician trust and regulatory compliance.
- **Operational transformation:** 75–85% diagnostic time reduction, >99% step-count reduction (47 manual steps eliminated), and 12–18 percentage point inter-rater consistency improvement.
- **Governance readiness:** Five-layer Decision Governance Framework with RACI accountability, exception handling, KPI monitoring, and audit trails supporting FDA SaMD and EU AI Act compliance.
- **Financial viability:** ROI of 135–185% over three years with 12–18 month payback period; 65–69% TCO reduction versus multi-vendor status quo.
- **Complete objective achievement:** All seven research objectives achieved; five hypotheses supported with statistical evidence and honest caveats.
- **Boundary conditions acknowledged:** Retrospective design, public datasets, proxy metrics, and bounded expert panel size ($n=12$) constrain the claims; a structured seven-study future research agenda addresses each limitation.

This thesis has demonstrated that domain fragmentation in biomedical AI is, in many cases, a design convention rather than a technical necessity. The RGAIG shared-pipeline framework achieved 78.3–100.0% accuracy across Autism Spectrum Disorder (ASD) and was endorsed by domain experts ($M = 4.4/5$, $n=12$). **These results provide a governance-aligned, financially justified foundation for transforming isolated diagnostic tools into an integrated organisational capability, positioned for responsible clinical deployment through a phased validation roadmap.**

This dissertation establishes that integrated clinical AI—multimodal, literature-informed, agent-driven, and governance-embedded—is not a future aspiration but a present possibility, bounded by identifiable conditions that a structured research agenda can resolve.

In sum, this dissertation argues that responsible, explainable, and governed AI-driven biomedical screening is not only technically achievable but organisation-

ally necessary—and provides the first unified framework, empirical evidence base, and deployment decision architecture to make it actionable.

Appendix E

Chapter 5 — Discussion and Recommendations Supplementary Material

Appendix: Chapter 5 — Discussion and Recommendations Supplementary Material

This appendix contains supplementary tables, figures, and detailed analyses that support Chapter 5 but were moved from the main text to maintain narrative focus. Each item is cross-referenced from the main chapter.

Appendix E Objective. Provide detailed compliance mappings, ROI sensitivity analysis, examiner Q and A, and extended discussion tables that support Chapter 5.

Table 5.81. Appendix E Roadmap: Contents and Purpose.

Content	What It Contains
Compliance	NIST/ISO detailed alignment matrices
Financial	ROI sensitivity, NPV projections
Governance	Extended RGAIG assessment
Examiner prep	Anticipated Q and A with responses
Commercialisation	Market sizing, deployment pathway

Responsible AI: Five-Pillar Implementation Framework

The five-pillar responsible AI framework (transparency, fairness, privacy, accountability, beneficence) operationalises ethical AI principles through enforceable technical controls. Implementation details, including pillar-level goals, enforced controls, and audit evidence requirements, are in the Supplementary Volume.




Figure 5.18. Appendix Task Flow: Supplementary Material for Chapter 5 (Discussion and Recommendations)

Data Protection Controls: Enforcing Zero-PHI Model Access

RGAIG enforces a strict zero-PHI policy through four layers of technical controls, ensuring no raw identifiers reach external model endpoints:

- **Data-path hard controls:** Automated de-identification pipelines strip PHI before data enters any processing stage; allowlist-only field forwarding.
- **Model usage mode restrictions:** External models operate in inference-only mode with no data retention; API calls transmit only de-identified feature vectors.
- **Bring-model-to-data patterns:** Where feasible, models are deployed on-premises so that raw data never leaves the institutional boundary.
- **Audit verification:** Post-processing scans confirm zero PHI leakage; every data movement is logged with immutable audit events.

The detailed control specifications and audit evidence requirements are documented in the Supplementary Volume.

User Notification and Operation Accountability

Every AI operation in RGAIG generates a visible notification and an immutable audit event, following a “notify first, act second” design principle. The notification mapping covers six operation types:

- **Data collection:** User notified before any patient data is accessed or processed; consent verification logged.
- **RAG retrieval:** Notification when the system queries the medical literature corpus; retrieved sources displayed.
- **Model inference:** Alert when an ASD screening model runs; confidence score and model version shown.
- **Tool execution:** Notification when agentic tools (e.g., SHAP explainer, bias auditor) are invoked.

- **Result sharing:** User informed before any screening output is transmitted to another system or clinician.
- **Output storage:** Notification when results are persisted to the audit trail or patient record.

The complete operation-type notification mapping is documented in the Supplementary Volume.

Fairness Metrics and Bias Handling

RGAIG measures fairness continuously across three levels using ten metrics monitored across the full system lifecycle:

- **Data level:**
 - Representation balance across demographic subgroups
 - Label balance ensuring no systematic class skew by protected attribute
- **Model level:**
 - Demographic parity (equal positive prediction rates)
 - Equal opportunity (equal true positive rates)
 - Equalised odds (equal true positive and false positive rates)
- **Outcome level:**
 - Decision parity across clinical pathways
 - Error impact assessment per subgroup
 - Fairness drift monitoring over 30-day rolling windows

The complete fairness metrics framework and bias handling procedures are documented in the Supplementary Volume.

Responsible AI Validation Techniques

Each responsible AI pillar is validated through defined technical tests, governance processes, and measurable quality checks aligned with NIST AI RMF, ISO 42001, and GDPR. The validation matrix covers the following six pillars across data, model, runtime, and governance layers:

- **Fairness:** Demographic parity, equal opportunity, and equalised odds tested per subgroup; Fairlearn audit with $\leq 2\%$ variance threshold.
- **Privacy:** De-identification verification, PII leakage scans, differential privacy budget tracking, and HIPAA compliance checks.
- **Transparency:** SHAP feature importance, RAG citation grounding, confidence calibration, and audit trail completeness.
- **Robustness:** Adversarial input testing, out-of-distribution detection, PSI drift monitoring, and cross-validation stability analysis.
- **Safety:** Hallucination gate enforcement, human-in-the-loop escalation, clinical plausibility checks, and override rate monitoring.
- **Accountability:** RACI role mapping, governance committee review, regulatory submission documentation, and incident response protocols.

The complete validation techniques specification is documented in the Supplementary Volume.

Hallucination Control and Verification Framework

RGAIG implements hallucination controls at six pipeline layers. Any output failing the grounding gate or confidence gate is blocked from patient-facing display until human review.

1. **Input scoping:** Constrain queries to the clinical domain ontology, rejecting out-of-scope or adversarial prompts before they reach the model.
2. **Knowledge grounding:** Anchor every generated statement to retrieved evidence from the curated medical literature corpus (ChromaDB, 1.17 GB).
3. **Model-time calibration:** Apply temperature scaling and nucleus sampling to reduce confident but unsupported assertions during generation.
4. **Post-generation verification:** Cross-check generated claims against source documents; flag any assertion lacking a retrievable citation.
5. **Runtime gates:** Enforce confidence thresholds (≥ 0.85 faithfulness) and block outputs that fail grounding or clinical plausibility checks.
6. **Continuous monitoring:** Track hallucination rates over 30-day rolling windows; trigger model review when rates exceed 2%.

The complete hallucination control specification is documented in the Supplementary Volume.

Five-Pillar Clinical AI Deep Audit Framework

The clinical AI deep audit framework evaluates 97 dimensions across five pillars. Table 5.82 summarises each pillar's scope and dimension count. Each pillar contains 18–20 audit dimensions assessed by scope, method, risk level, and remediation strategy. The framework supports regulatory submissions (EU AI Act, FDA SaMD, HIPAA), periodic self-assessment, and third-party audits. The complete 97-dimension audit matrix is documented in the Supplementary Volume.

Table 5.82. Five-Pillar Clinical AI Deep Audit Framework — Dimension Summary

Pillar	Name	Dimensions	Scope
1	Data Responsibility and PHI Governance	20	Data lineage, de-identification, consent, quality gates, access controls
2	Model Responsibility	19	Training rigour, validation protocols, bias testing, version control, reproducibility
3	Output Responsibility and Clinical Safety	20	Prediction accuracy, confidence calibration, hallucination prevention, human-in-the-loop gates
4	Monitoring and Drift	18	PSI drift detection, fairness drift, performance dashboards, retraining triggers
5	Governance and Compliance	20	Regulatory alignment (FDA, EU AI Act, HIPAA), audit trails, RACI accountability
Total		97	

Note. PHI = protected health information; PSI = Population Stability Index; RACI = responsible, accountable, consulted, informed. Dimension counts are approximate; the complete audit matrix is in the Supplementary Volume.

Executive Value Propositions table

Table 5.83 maps five C-suite roles to their current pain points, the corresponding RGAIG solution, and the measurable impact.

Table 5.83. Executive Value Propositions by C-Suite Role

Role	Current Pain Point	RGAIG Solution	Measurable Impact
CIO	5–8 vendor contracts per hospital for siloed ASD screening tools	Single ASD screening platform replaces fragmented ASD specialist tools	65–69% TCO reduction; single ASD integration point
COO	ASD specialist bottleneck limits ASD screening throughput	Automated ASD pipeline (94.6% time reduction)	500→26.8 min per ASD case [5]
CMO	Clinician trust deficit slows ASD screening adoption	RAG explainability + SHAP feature maps	12–18 pp inter-rater improvement; $M=4.6/5$ trust
CFO	Unclear AI investment return	Three-scenario ROI model with sensitivity analysis	Projected 135–185% 3-year return
CISO/DPO	Privacy and compliance risk	HIPAA/GDPR alignment; de-identification; audit trails	<10% false positive rate; full traceability

Note. CIO = Chief Information Officer; COO = Chief Operating Officer; CMO = Chief Medical Officer;

CFO = Chief Financial Officer; CISO = Chief Information Security Officer; DPO = Data Protection Officer; pp = percentage points; TCO = total cost of ownership.

TAM/SAM/SOM table and Revenue Projection table

Table 5.48 presents the market opportunity analysis across total addressable, serviceable addressable, and serviceable obtainable market layers with CAGR projections.

Table 5.84 reports the three-year revenue projection for the RGAIG B2C market entry across device sales, subscription revenue, and cumulative ROI.

Table 5.84. Three-Year Revenue Projection — RGAIG B2C Market Entry

Metric	Year 1	Year 2	Year 3
B2B hospital pilots	5	15	50
B2C subscribers	500	2,500	10,000
Device revenue (\$999/unit)	\$500K	\$2.5M	\$10.0M
Subscription revenue (\$29.99/mo)	\$180K	\$900K	\$3.6M
Total revenue	\$680K	\$3.4M	\$13.6M
Cumulative investment	\$2.0M	\$3.5M	\$5.0M
Projected ROI	−66%	−3%	172%

Note. Revenue projections are illustrative estimates based on published market benchmarks and comparable B2C health-tech pricing (e.g., Oura Ring, Whoop). Projected ROI 135–185% at Year 3 aligns with three-scenario sensitivity analysis (pessimistic: 85%; base: 172%; optimistic: 240%). Actual results will depend on regulatory approvals, market adoption rates, and reimbursement pathways.

Governance CMM table

Table 5 presents the Governance Capability Maturity Model.

Table 5.85 presents clinical AI Governance Capability Maturity Model (Appendix).

Table 5.85. Clinical AI Governance Capability Maturity Model (Appendix)

Level	Description	Governance Characteristics	Framework Alignment
1	Ad-hoc ASD screening AI use	No formal governance; individual clinician experiments with ASD screening tools; no audit trails	Pre-adoption: organizational readiness assessment required
2	Emerging controls	Basic AI usage policy; informal oversight; no systematic bias monitoring or explainability	Layer 1 (Business Decision) partially implemented; governance committee formed
3	Structured governance	Formal AI governance committee; RACI roles defined; bias audits quarterly; RAG explainability deployed	Layers 1–3 (Business, Process, Data/Intelligence) operational; SHAP + RAG active
4	Integrated capability	Enterprise-wide AI management; continuous monitoring dashboards; automated drift detection; regulatory submissions	All 5 layers operational; FDA/EU compliance documented; KPI dashboards live
5	Optimized ecosystem	Continuous improvement loop; federated learning across institutions; automated retraining; real-time fairness monitoring	Full framework maturity; multi-institutional deployment; research-to-practice pipeline

Note. RACI = responsible, accountable, consulted, informed; RAG = retrieval-augmented generation;

SHAP = Shapley additive explanations; FDA = U.S. Food and Drug Administration; EU = European Union; KPI = key performance indicator. Levels adapted from CMMI Institute's capability maturity model for ASD screening contexts.

KPI Escalation table and Escalation Ladder figure

Table 5.86 maps eight operational KPIs to traffic-light thresholds, designated authorities, and mandatory response times.

Table 5.86. KPI Threshold-to-Action Escalation Protocol (Appendix)

KPI	Green	Amber	Red	Escalation Path	Response
Domain accuracy	$\geq 85\%$	80–85%	$< 80\%$	AI Committee → Board	24h
Demographic fairness	$\leq 2\%$ var.	2–5% var.	$> 5\%$ var.	Compliance → Board	Immediate
Model drift (PSI)	< 0.10	0.10–0.25	> 0.25	Data Science → Committee	48h
Audit trail coverage	100%	95–99%	$< 95\%$	IT Ops → Compliance	24h
RAG coherence	≥ 0.85	0.75–0.85	< 0.75	Data Science → Clinical	48h
Override rate	$< 15\%$	15–25%	$> 25\%$	Clinical Lead → Committee	Weekly
System uptime	$\geq 99.5\%$	98–99.5%	$< 98\%$	IT Ops → CIO	4h
Training completion	$\geq 90\%$	70–90%	$< 70\%$	HR → COO	Monthly

Note. PSI = Population Stability Index (model drift metric). Var. = variance across demographic subgroups. Green = normal operations; Amber = monitoring alert with corrective action; Red = mandatory escalation to designated authority. Thresholds derived from regulatory standards (FDA SaMD, EU AI Act) and expert panel consensus.

Figure 5 visualizes the four-level governance escalation architecture, from operational monitoring (Level 1) through Board intervention for patient safety incidents (Level 4).




Figure 5.19. Four-Level Governance Escalation Ladder for Clinical AI Deployment (Appendix)

Note. Escalation triggers are cumulative: a Level 2 response may trigger Level 3 if the condition persists beyond the specified timeframe. SLA = service-level agreement. The escalation protocol operationalizes the governance maturity model at Level 3 (Defined) and above; the five-level maturity framework is detailed in the Supplementary Volume.

Detailed 7-Study Future Research table

Table 5.48 operationalises these into seven specific studies with defined research questions, methods, and expected outcomes.

Implementation Roadmap figure

Figure 5.20 presents the phased implementation roadmap from research completion through clinical deployment and market launch, with key milestones spanning 2026–2028.




Figure 5.20. Implementation Roadmap: From Research to Clinical Deployment (2026–2028, Appendix)

Note. Five phases progress from ASD-focused validation (high priority) through real-time prototyping, regulatory submission (FDA 510(k)/CE marking), ASD pilot deployment, and operational scale-up. Priority levels reflect resource allocation and risk staging.

Five-Layer RGAIG Architecture recap figure




Figure 5.21. Healthcare AI Implementation Model: Five-Layer RGAIG Architecture for Clinical Deployment

Note. RGAIG = Responsible Generative AI Governance Framework; EHR = electronic health record; PACS = picture archiving and communication system; SHAP = SHapley Additive exPlanations; RAG = retrieval-augmented generation; PII = personally identifiable information; MCP = model control protocol. Information flows upward from data through ASD classification, explanation generation, governance checks, to ASD screening support. Governance constraints propagate downward.

AI Governance Decision Flow

Four-gate governance validation pipeline (Figure 5.22).

1. **Bias detection.** Disparate Impact ≥ 0.80 .
2. **Fairness validation.** Equal opportunity across demographic groups.
3. **Confidence thresholding.** $\geq 85\%$.
4. **Compliance verification.** 525-module taxonomy.

Failed gates trigger remediation: Level 1 (automated re-evaluation) $\rightarrow$ Level 4 (regulatory reporting). Pipeline achieves **98.4% pass rate** across the ASD screening taxonomy. Responsible AI is operationalised at inference time, not retrospectively.




Figure 5.22. AI Governance Decision Flow with Four-Gate Validation. Predictions pass through four governance gates before clinical use.

Continuous Monitoring and Retraining Pipeline

Five-stage MLOps lifecycle (Figure 5.23).

- **Monitoring.** Live EEG ingestion, ensemble inference, performance metrics (rolling accuracy, FPR, uptime), three drift types (data, concept, feature).
- **Retrain triggers.** $\text{PSI} > 0.25$ for 7 consecutive days OR rolling accuracy drops > 3 pp.
- **Validation gate.** Required before re-deployment.
- **Scheduled retraining.** Quarterly safety net independent of drift detection.

This addresses a literature-review gap: most ASD screening systems lack post-deployment performance assurance.




Figure 5.23. Continuous Monitoring and Retraining Pipeline. Five-stage MLOps lifecycle maintaining ASD diagnostic accuracy post-deployment.

User Interaction Flow: Parent and Clinician Journeys

Dual-persona user-interaction flow (Figure 5.24).

- **Parent journey (accessibility).** 5-min EEG → plain-language RAG screening report with risk score + recommended actions.
- **Clinician journey (analytical depth).** SHAP feature attribution, EEG topographic maps, confidence intervals; HITL override authority preserved.
- **Cross-link.** Parent shares AI screening report with ASD specialist → B2C-to-B2B referral pathway.
- **Impact metrics.** Screening time 45 min → 5 min/patient; RAG faithfulness 0.92; modelled ROI 135–2,717 %.




Figure 5.24. User Interaction Flow: Parent and Clinician Journeys. Parent-facing screening connects to clinician-supervised review and action.

Lean Waste Elimination Table

Table 5.87. Lean Waste Elimination: Manual vs. RGAIG-Automated Diagnostic Pipeline

Lean Manual Pipeline (Before)	RGAIG Pipeline (After)	Redn.
Wait: 3–6 mth ASD specialist wait; ASD screening results delayed 2–4 wk	ASD screening 4.2 min; same-day ASD screening results	85–95%
Trans: EEG shipped between recording, reading, referring	Single-platform; digital delivery	100%
Over-proc.: Repeated ASD specialist review and redundant intake	Single ASD pipeline streamlines screening	80%
Defect: Inter-rater variability 60–70%; missed dx	ASD accuracy 97.67%; calibrated confidence	40–50%
Motion: 3–5 tools per patient; manual data entry	Unified interface; agentic orchestration	90%
Invent: Unread EEG backlog during ASD specialist shortages	Real-time processing; ASD screening queue	95%
Over-prod.: Unnecessary ASD specialist referrals from primary care	ASD screening triage reduces inappropriate referrals	30–40%

Note. Lean waste categories from Womack and Jones [118]. Reductions are modelled estimates

based on Chapter 4 results and published benchmarks, not clinical deployment.

Workforce Transformation Table

Table 5.88. Workforce Transformation: Tasks Eliminated, Augmented, and Created by RGAIG

Role	Tasks Eliminated	Tasks Augmented	Tasks Created
Neurologist	Routine EEG visual screening for common patterns; manual report writing for straightforward ASD cases	Complex ASD screening case interpretation supported by AI confidence scores and SHAP feature maps; literature-grounded review with RAG context	ASD screening output review and sign-off; override documentation with clinical rationale; quarterly model performance audit participation
EEG Technician	Manual artefact annotation; paper-based recording logs; physical handling and transfer of EEG media	Electrode placement guided by impedance monitoring; real-time signal quality feedback from preprocessing pipeline	ASD screening system operation and troubleshooting; data quality gate monitoring; patient-facing explanation of the ASD screening process
Clinical Administrator	Manual scheduling across multiple ASD specialist queues; redundant data entry across departmental systems	Triage prioritisation informed by AI confidence scores; resource allocation guided by ASD screening throughput dashboards	Governance dashboard monitoring; KPI reporting to hospital leadership; vendor relationship management for ASD screening platform

Note. This analysis is prospective, based on the RGAIG architecture design and published healthcare workforce transformation literature [3], not on observed deployment outcomes. The net effect is role enhancement rather than displacement: ASD specialist time is redirected from routine screening to complex ASD cases requiring expert judgement.

Capacity Planning Table

Table 5.89. Capacity Planning: RGAIG Throughput Under Three Deployment Scenarios

Parameter	Small Clinic	Mid-Size Hospital	Academic Medical Centre
Daily EEG volume	10–20	40–80	150–300
GPU requirement	1× A100 (shared)	1× A100 (dedicated)	2–3× A100 (dedicated)
Processing capacity	340/day	340/day per GPU	680–1,020/day
Utilisation rate	3–6%	12–24%	44–88%
Clinician review time	2–3 min/ASD case	2–3 min/ASD case	2–3 min/ASD case
Bottleneck	Clinician availability	Clinician availability	GPU at peak; clinician at steady-state
Annual cost (GPU)	\$3,600 (cloud)	\$14,400 (on-premise lease)	\$43,200 (3-GPU cluster)

Note. Processing capacity assumes 4.2 minutes per patient end-to-end. GPU costs based on 2024 cloud pricing (\$1.50/hr A100 spot instance) or on-premise lease equivalents. Clinician review time (2–3 min) is the operational bottleneck in all scenarios, not compute capacity. Utilisation rate = daily volume ÷ processing capacity.

Clinical Workflow Integration Figure




Figure 5.25. Clinical Workflow Integration: End-to-End Diagnostic Sequence

Note. EDF = European Data Format; ICA = independent component analysis; PDF = portable document format. The workflow demonstrates that ASD screening adds only 4.2 minutes of automated processing after the standard 20-minute EEG recording, with clinician oversight at initiation and interpretation.

Clinical Journey Mapping Table

Table 5.90. Clinical Journey Mapping: Patient-to-Outcome Touchpoints Across the RGAIG Screening Pathway

#	Touchpoint	What Happens	Actor	AI Role	Patient Exp.
1	Symptom onset	Patient/family notices cognitive, motor, or behavioural change	Patient	None (pre-system)	Anxiety; 2.5-yr delay
2	Primary care	GP assesses symptoms; refers for EEG evaluation	GP/PC	Risk score (opt.)	4–12 wk wait
3	EEG recording	Wearable placed; 20-min resting-state EEG	Tech/scr	Calibration; SQI gate	Non-invasive; 20 min
4	AI screening	Preprocess→classify→XAI→RAG	RGAIG	Classify + explain	<5 min
5	Clinician review	Reviews screening result, confidence, SHAP, RAG	Neuro.	Decision support	Pending sign-off
6	Result comms	Discusses findings; dual-persona RAG report	Clinician	Patient summary	Cited explanation
7	Follow-up	Treat, re-scan 3 mth, or annual monitor	MDT/CP	Drift alerts	Clear next steps

Note. Touchpoints 1–2 = existing pathway (no AI); 3–6 = RGAIG-augmented; 7 = post-screening. Primary value at touchpoint 4: objective AI screening replacing subjective judgement while preserving clinician authority. GP = general practitioner; PCP = primary care physician; MDT = multidisciplinary team.

Doctor Review Protocol Table

Table 5.91. Doctor Review and Report Validation: Six-Step Clinical Oversight Protocol

#	Review Action	Review	Checks	Outcome	Escalation
1	AI generates prediction + SHAP + RAG	System	Conf. $\geq 85\%$; audit log	Report queued	$<85\% \rightarrow$ flag
2	Clinician opens report dashboard	Neuro.	Screening result, confidence, XAI	Accept/Override/Defe	Senior clinician
3	Cross-check patient history	Neuro.	EHR; meds; comorbidities	Confirm or re-scan	MDT review
4	Sign-off and release	Attendir	Final screening assessment documented	Released; EHR logged	Second opinion
5	Quarterly accuracy audit	Dept head	AI vs. final (100 cases)	Concordance $\geq 92\%$	$<92\% \rightarrow$ retrain
6	Annual governance review	Ethics bd	Fairness; consent; adverse	Gov. pass/fail	Fail $\rightarrow$ suspend

Note. No report reaches the patient without clinician sign-off (Step 4). All actions timestamped in immutable audit trail. Steps 1–4 per-patient; 5–6 periodic QA. MDT = multidisciplinary team; EHR = electronic health record.

Doctor vs Patient Recommendation KPI Table

Doctor Versus Patient Recommendation KPI Framework. Table 5.92 below presents the detailed analysis.

Table 5.92. Doctor Versus Patient Recommendation KPI Framework

Dimension	Doctor Recommendation Criteria	Patient / Parent Recommendation Criteria
Clinical usefulness	Improves decision support, adds clinically relevant insight, supports interpretation	Helps understand screening status, next steps, and follow-up pathway
Accuracy confidence	Output appears technically and clinically credible	Result feels believable and not arbitrary
Explainability	Rationale is interpretable, evidence-linked, and clinically usable	Explanation is simple, understandable, and meaningful
Workflow fit	Does not disrupt clinic flow or add cognitive burden	Does not create excessive onboarding burden or confusion
Time value	Saves consultation or assessment time	Reduces repeated paperwork, repetition, or uncertainty
Safety confidence	Human oversight remains, output is bounded and non-overclaiming	System feels safe and not misleading or risky
Reliability	Output appears consistent across cases and use situations	Experience feels stable and dependable
Privacy confidence	Data handling appears compliant and auditable	Personal data feels protected and respectfully handled
Adoption readiness	Suitable for pilot or controlled clinical use	Worth using again in future assessments
Recommendation intention	Would recommend to colleagues or for formal pilot use	Would recommend to other patients / parents or reuse personally

Medical-Legal Liability Framework Table

Table 5.93. Medical-Legal Liability Allocation for AI-Assisted Diagnosis

Scenario	Liable Party	Legal Doctrine	RGAIG Mitigation
AI flags abnormality; clinician overrides; harm	Clinician	Malpractice; override shifts liability to physician	Audit trail records override + rationale; RAG basis
AI misses abnormality; clinician relies on AI; harm	Shared (mfr + clinician)	Product liability (mfr); negligence (over-reliance)	ECE<0.05 quantifies uncertainty; governance requires judgement
AI correct; clinician misinterprets	Clinician	Malpractice; failure of judgement	RAG plain-language explanations; training required
System failure (downtime)	Inst. / vendor	Contract law; SLA breach; duty of care	Failover to manual workflow; heartbeat; SLA
Bias causes disparate outcomes	Mfr + inst.	Anti-discrimination; disparate impact; EU AI Act	DI $\geq$ 0.80 monitoring; quarterly bias audits

Note. Reflects 2024–2026 legal frameworks; no case law adjudicating AI-assisted ASD screening error exists. Liability allocation draws on radiology AI precedents (FDA SaMD), ASD screening support, and product liability doctrine. DI = disparate impact ratio; ECE = expected calibration error; SLA = service-level agreement.

Governance CMM Table

Clinical AI Governance Capability Maturity Model. The following table presents the detailed analysis.

Table 5.94. Clinical AI Governance Capability Maturity Model

Leve	Description	Governance Characteristics	Framework Alignment
1	Ad-hoc ASD screening AI use	No formal governance; individual clinician ASD screening experiments; no audit trails	Pre-adoption readiness assessment
2	Emerging controls	Basic usage policy; informal oversight; no bias monitoring or explainability	Layer 1 partial; governance committee formed
3	Structured governance	Formal committee; RACI defined; quarterly bias audits; RAG deployed	Layers 1–3 operational; SHAP + RAG active
4	Integrated capability	Enterprise-wide management; dashboards; drift detection; regulatory submissions	All 5 layers; FDA/EU compliance; KPI dashboards
5	Optimised ecosystem	Continuous improvement; federated learning; auto-retrain; real-time fairness	Full maturity; multi-institutional deployment

Note. RACI = responsible, accountable, consulted, informed; RAG = retrieval-augmented generation; SHAP = Shapley additive explanations; FDA = U.S. Food and Drug Administration; EU = European Union; KPI = key performance indicator. Levels adapted from CMMI Institute’s capability maturity model for ASD screening contexts.

AI Governance Risk Assessment Matrix

Table 5.95. AI Governance Risk Assessment Matrix: Risk Analysis, Severity, Mitigation, and Decision Triggers Across 15 Governance Dimensions

#	Category	Risk Scenario	Sev.	Mitigation	Decision
1	Explain.	No interpretable rationale	Hi/Med	RAG + SHAP top- k	Reject if coh.<0.75
2	Trust	Override>25%	Hi/Med	Trust scorecard; calibration	Suspend>30%/14 d
3	Resp. AI	No ethics review	Crit/Lo	Ethics gate; RACI	Block until approved
4	Fairness	Var.>5% subgroups	Crit/Med	Quarterly subgroup; retrain	Halt; retrain 30 d
5	Privacy	Re-id from EEG	Crit/Lo	k -anon; encrypt.; audit	Quarantine; IRB
6	Gov.	No incident authority	Hi/Med	RACI; 4-level escalation	Committee; 24 h
7	Perform.	Acc. drops<85%	Hi/Med	PSI; bootstrap CI; alerts	Retrain; override
8	Hum-AI	Automation bias	Hi/Med	Conf. thresholds; logging	Training
9	Data	Artefacts/label noise	Hi/Med	12-step defence; CV	Quarantine; retrain
10	Decision	AI vs. guidelines	Med/Med	RAG cross-ref; flag	Specialist review
11	Monitor	Dashboard failure	Hi/Lo	Redundant; weekly audit	Backup; 4 h
12	Portab.	Fails new institution	Hi/Med	Transfer val.; calibrate	Block until $\geq$ 80%
13	Safety	Wrong hi-conf. dx	Crit/Lo	Calibration; human review	Case review; notify
14	Consens.	Misalign.>12 mth	Med/Hi	Delphi; roadmap	Escalate; 60 d
15	Context	Wrong pipeline	Hi/Lo	Domain detect.; schema	Block; manual route

Note. Severity: Critical = patient harm or regulatory violation; High = significant operational impact; Medium = degraded performance. Likelihood: High >30%, Medium 10–30%, Low <10%. IRB = Institutional Review Board; SLA = service-level agreement; PSI = Population Stability Index; RAG = retrieval-augmented generation; RACI = responsible, accountable, consulted, informed.

Adoption KPI Framework Table

Key Performance Indicators for Framework Adoption.... Table 5.96 below presents the detailed analysis.

Table 5.96. Key Performance Indicators for Framework Adoption Readiness

KPI Cat.	Indicator	Target Thresh.	Decision Rule
Human	ASD classification accuracy	$\geq 85\%$	Pilot if met; defer if not
	Balanced accuracy	$\geq 85\%$	Framework-level viability
	AUC (ASD screening)	≥ 0.85	Discrimination adequacy
	Trust score (Likert, expert/patient)	$\geq 3.5/5$	Explanation acceptable
	Usability score (onboarding)	$\geq 3.5/5$	Workflow viable
	Adoption willingness	$\geq 70\%$ positive	Stakeholder readiness
	Onboarding completion rate	$\geq 85\%$	Intake design effective
	Processing time reduction	$\geq 30\%$ reduction	Automation justified
	Decision consistency (inter-rater)	$\kappa \geq 0.60$	Standardisation achieved
	Privacy compliance score	Pass/Fail	Regulatory readiness
Governance	Audit trail completeness	$\geq 90\%$	Accountability demonstrated

Note. KPI = Key Performance Indicator. Thresholds are derived from literature benchmarks and expert panel consensus. The ASD pipeline qualifies for pilot deployment only if all critical thresholds are met; otherwise further validation is required.

Hidden Risk Categories Table

Table 5.97. Hidden Risk Categories: Eight Threats That Evade Traditional Risk Assessment

Hidden Risk	Why It Is Missed	How It Manifests	RGAIG Detection
Automation compliance	Clinicians over-trust the ASD screening system after prolonged high accuracy	Override rate drops below 2%; errors unchallenged	Monitor override rate; alert if <5%
Label leakage	Training labels encode historical screening bias	Model replicates and amplifies disparities	Bias audit per retrain; $DI \geq 0.80$
Silent distribution shift	Patient demographics change without triggering alerts	Accuracy degrades gradually over months	Weekly KS-test on prediction distributions
Explanation gaming	Users selectively read only confirming SHAP/RAG	Confirmation bias amplified; contradictory evidence ignored	Track override vs. accept correlation
Governance fatigue	Quarterly audits become routine; scrutiny declines	Audit scores remain high but genuine oversight weakens	Rotate auditors; inject synthetic anomalies
Cascading agent failure	One agent's error propagates through dependency chain	Multiple pipeline outputs fail simultaneously	Dependency graph monitoring; circuit breaker
Consent decay	Original consent stale as AI capabilities expand	Patients unaware of new model uses or sharing	Re-consent trigger at major version updates
Regulatory lag	Regulations change; system built for previous rules	Non-compliance discovered at next audit	Regulatory watch service; quarterly gap scan

Note. Hidden risks differ from operational risks (Table 5) in that they arise from human–AI interaction dynamics rather than technical malfunction. Each risk has a detection mechanism embedded in the RGAIG governance layer. DI = disparate impact ratio; KS =

Hidden Risk Measurement Table

Table 5.98. Hidden Risk Measurement: Proxy Metrics, Thresholds, and Monitoring Cadence

Hidden Risk	Proxy Metric	Alert Thresh.	Freq.	Measurement
Auto. compl.	Clinician override rate	<5% overrides/30 d	Wkly	Log: accept/reject/defer
Label leakage	Subgroup acc. var.	Var.>5% across demogr.	Per re-train	Stratified eval by age, sex
Silent drift	KS stat. on predictions	$p < 0.05$ (dist. shift)	Wkly	KS test on softmax outputs
Expl. gaming	Override–expl. corr.	Pearson $r > 0.70$	Mthly	Override vs. SHAP polarity
Gov. fatigue	Audit novelty rate	<2 new findings/qtr	Qtrly	Compare audit cycles
Cascading fail.	Agent co-failure rate	>1 co-failure/24 h	Real-time	Dep. graph + breaker
Consent decay	Time since re-consent	>12 mth since renewal	Mthly	Consent DB age check
Reg. lag	Days since reg. scan	>90 d without scan	Qtrly	Watch service log

Note. Each proxy metric is designed to surface the hidden risk *before* it causes patient harm. The

measurement cadence ranges from real-time (cascading failure) to quarterly (governance fatigue, regulatory lag), reflecting the speed at which each risk can materialise. KS = Kolmogorov–Smirnov.

Operational Risk Matrix Table

Table 5.99. Operational Risk Matrix: Eight Production Failure Modes and Mitigations

Risk	Lk.	Im.	Detection	Mitigation	Owner
Data drift	High	High	PSI>0.20 weekly	Recalibrate; retrain	Data Eng
Model stale.	Med	High	AUC drop>3% monthly	Champion–challenger; A/B test	ML Ops
Latency spike	Med	Med	P95>500 ms	Auto-scale; queue mgmt	Platform
Adversarial input	Low	High	Anomaly>3 σ from train	Validate; reject and flag	ML Eng
SHAP timeout	Med	Med	Explanation>30 s	Fallback to LIME; cache	XAI Lead
Pipeline break	Med	Low	CI/CD regression failure	Pin versions; rollback	DevOps
Consent w/d	Low	High	Consent DB flag	Purge 72 h; retrain	Ethics
GPU SPOF	Low	High	Heartbeat fails	CPU failover; degraded	Infra

Note. Each failure mode has a named owner accountable for detection and response.

Operational risks can occur today in a running system, while emerging risks (Table 5) represent future threats. PSI = Population Stability Index; P95 = 95th percentile; CI/CD = continuous integration/continuous deployment; SPOF = single point of failure.

Enterprise Risk Register Table

Table 5.100. Consolidated Enterprise Risk Register: Top-10 Risks with Mitigation and Residual Assessment

ID	Risk	L	I	Sc.	Mitigation Strategy	Owner	Res.
R1	Accuracy degrad.	3	5	15	PSI<0.10 monitoring; auto-retrain; quarterly validation	ML Ops	6
R2	Adoption resistance	4	4	16	ADKAR pro-gramme; champions; training; feedback	CMO	8
R3	Regulatory non-compl.	3	5	15	Compliance gates; pre-deploy checklist; counsel	Compliance	5
R4	Data privacy breach	2	5	10	De-id ($k \geq 5$); access controls; encryption	CISO	4
R5	Algorithmic bias	3	4	12	DI $\geq$ 0.80; quarterly audits; retrain protocol	Ethics Bd	6
R6	Explanation fidelity	3	3	9	RAG $\geq$ 0.50; XAI triangulation	AI Lead	4
R7	Infra downtime	2	4	8	CPU failover;	IT Dir.	3

Risk-Benefit Decision Matrix Table

Table 5.101 documents risk-Benefit Decision Matrix: Gains vs Risks by Deployment Area.

Table 5.101. Risk-Benefit Decision Matrix: Gains vs Risks by Deployment Area

Area	Expected Benefit	Key Risk	Balance
Technical accuracy	ASD screening $\geq$ 85%	No prospective validation yet	Favourable (ASD)
Explainability	RAG-grounded clinician/patient output	Explanation quality not patient-tested	Favourable (expert-validated)
Workflow efficiency	30–40% effort reduction estimated	Simulation-based, not field-measured	Moderate
Trust/adoption	Expert trust $M = 4.2/5$	Expert panel only, no patient data	Moderate
Governance/safety	Audit, consent, human-in-the-loop designed	No field audit or regulatory approval	Conditional
B2C continuity	Hospital-to-home pathway conceptualised	No mobile deployment evidence	Weak (future scope)

Three-Level Evidence Framework Table

Three-Level Evidence Framework for Deployment Readiness. Table 5.102 below presents the detailed analysis.

Table 5.102. Three-Level Evidence Framework for Deployment Readiness

Evidence Type	Source	What It Proves	Measurable?
Technical	EEG/IoT/imaging, model outputs	The framework performs accurately for ASD screening	Yes: accuracy, AUC, F1, CI
Human	Patient, parent, doctor, expert ratings	The framework is understandable, usable, and trustworthy	Yes: trust, usability, clarity scores
Operational	Workflow logs, onboarding data, process metrics	The framework improves efficiency and reduces manual burden	Yes: time, completion, effort

Note. All three evidence levels must converge for a deployment recommendation.

Technical evidence alone is insufficient without human acceptance and operational viability.

North America Deployment Viability Table

North America Deployment Viability by Use Case. Table 5.103 below presents the detailed analysis.

Table 5.103. North America Deployment Viability by Use Case

Use Case	Viability	Rationale
Clinician decision support with human review	Strong	Strongest regulatory and workflow fit
Hospital pilot for triage/assessment support	Strong	Easier to justify than autonomous ASD screening
Patient-facing monitoring/app with escalation	Possible	Needs stronger privacy, usability, and boundary controls
Autonomous ASD screening from consumer EEG alone	Weak	High regulatory and safety risk

Note. Based on FDA Clinical Decision Support Software guidance, Health Canada SaMD guidance, and HIPAA/provincial privacy requirements. The strongest North America deployment position is B2B hospital pilot with clinician review and auditability.

7-Study Future Research Table

Areas of Further Research: Seven Priority Directions. The following table presents the detailed analysis.

Table 5.104. Areas of Further Research: Seven Priority Directions

#	Direction		Rationale and Scope	Priority
1	Federated learning		Multi-site training without centralising EEG data; addresses privacy regulations and non-IID data heterogeneity [123]	High
3	Longitudinal monitoring	ASD	Extend the cross-sectional framework to track developmental change and follow-up risk patterns; repurpose PSI drift detection	High
4	Multi-modal fusion (EEG+MRI)	fu-	Combine EEG with structural/functional MRI for overlapping EEG signatures (epilepsy vs. non-epileptic seizures)	Medium
5	LLM port generation	clinical re-	Replace RAG with fine-tuned medical LLMs (Med-PaLM, BioGPT); leverage existing guardrail infrastructure (P5)	Medium
6	Regulatory box pilot	sand-	Partner with FDA Digital Health Center; real-world 510(k) evidence; validate governance modules empirically	Medium
7	Health modelling	economic	Validate ROI projections (Table 5) with clinical pilot data; enable institution-specific ROI calculators	Low

Note. Priorities based on impact on clinical deployment readiness. PSI = Population Stability Index;

MCI = mild cognitive impairment; LLM = large language model.

Implementation Roadmap Figure




Figure 5.26. Implementation Roadmap: From Research to Clinical Deployment (2026–2028)

Note. Five phases progress from ASD-focused validation (high priority) through real-time prototyping, regulatory submission (FDA 510(k)/CE marking), ASD pilot deployment, and operational scale-up. Priority levels (high, medium, low) reflect resource allocation and risk staging. FDA = U.S. Food and Drug Administration; CE = Conformité Européenne.

Master Decision-Making Table

Table 5.105 below presents the ten decision areas evaluated for ASD framework deployment readiness.

Table 5.105. Chapter 5 Master Decision-Making Table

Area	Main Evidence	Key Question	KPI	Output
Hyp. interp.	Ch.4 hypothesis verdicts	What do supported hypotheses mean?	Support level	Refined claims
Domain ready.	Domain performance, trust, robustness	Which domain is ready?	Acc., AUC, trust	Adopt/pilot/defer
B2B adoption	Clinician usefulness, workflow, ROI	Should hospital adopt?	Trust, effort	B2B decision
B2C cont.	Patient usability, privacy, app	Is patient ext. ready?	Usability, privacy	B2C scope
Expl. & trust	Clinician/patient ratings, RAG, XAI	Is output acceptable?	Trust, clarity	Accept. readiness
Workflow	Timing, automation, integration	Does it reduce burden?	Time, ASD screening throughput	Ops. readiness
ROI & value	Reuse, labour reduction, scalability	Is adoption worth cost?	Labour, reuse	ASD business case
Gov. & safety	Consent, audit, escalation, oversight	Is it safe and governable?	Gov. score	Safe-use readiness
Risk–benefit	All technical and ops. evidence	Do benefits outweigh risks?	Risk matrix	Deploy. boundary
Final readiness	All prior decision areas	Adopt, pilot, or defer?	Combined	Recommendation

Appendix F

Survey Instrument for RGAIG Framework Evaluation

Survey Instrument: RGAIG Multi-Stakeholder Assessment

This appendix presents the complete survey instrument designed for quantitative validation of the Responsible Generative AI Governance (RGAIG) framework. The instrument contains 82 items across eight sections, is designed for online administration (Qualtrics or REDCap) with an estimated completion time of 20–25 minutes, and requires full IRB approval with a minimum sample of $n \geq 250$ respondents to satisfy structural equation modelling (SEM) requirements across five stakeholder strata.

Survey Instrument Defensibility Pack

This pack answers the seven questions an examiner asks when evaluating a research instrument: (1) Which item measures which construct, and why? (2) How does each item map to a hypothesis? (3) What theoretical source is each construct anchored to? (4) What pilot validation refined the instrument? (5) How is reliability assessed? (6) What is the Likert scale and how does reverse-coding work? (7) What ethical safeguards apply? The 82-item instrument that follows is the operational survey; this pack is the scientific-operationalization wrapper.

Decision-support framing (ethical positioning). Throughout the instrument, the AI system is positioned as *clinician decision support*, not autonomous diagnosis. All items reference clinician-AI collaboration, never AI replacement of clinical judgement.

Table 5.106. Construct Mapping: Items, Variable Roles, and Theoretical Sources. Each survey construct is anchored to its published source instrument, with adaptation documented.

Construct	Variable role	Items	Source theory + citation
RAI Governance	IV	G1–G12 (Sec. C/D)	Responsible AI principles [38]; EU AI Act (2024); governance theory [10]
Perceived Explainability	Mediator 1	E1–E6 (Sec. D/F)	Explainable AI literature [36, 37]
Clinician Trust	Mediator 2	T1–T10 (Sec. D)	Trust in Automation [31]; TUI scale [32]
Adoption Intention / Workflow Readiness	DV	B13–B15 (Sec. B BI) + workflow items (Sec. E)	Technology Acceptance Model [29]; UTAUT [30]
AI Familiarity	Moderator	A4, A5	Adapted from (author?) [30] individual-difference items
Controls	Covariates	A1–A10 (Sec. A)	Standard demographic stratification variables

Note. Section letters refer to the 8-section instrument structure (A–H) shown below. The full administered instrument retains 82 items for stakeholder-stratified analysis; the constructs tested in the central H1–H6 mediation model use the items mapped above.

Table 5.107. Hypothesis-to-Item Mapping. Each hypothesis is linked to the survey items that test the corresponding path.

H#	Path tested	Items used in test
H1	RAI Governance → Perceived Explainability	IV: G1–G12 Mediator: E1–E6
H2	Perceived Explainability → Clinician Trust	Mediator: E1–E6 Trust: T1–T10
H3	Clinician Trust → Adoption Intention / Workflow Readiness	Trust: T1–T10 DV: B13–B15 + workflow items
H4	Explainability <i>mediates</i> Governance → Trust	G1–G12 + E1–E6 + T1–T10 (bootstrap mediation)
H5	Trust <i>mediates</i> Explainability → Adoption	E1–E6 + T1–T10 + DV items (bootstrap mediation)
H6	AI Familiarity <i>moderates</i> Explainability → Trust	A4–A5 (moderator) × E1–E6 (multi-group analysis)

Table 5.108. Pilot Validation Summary. The instrument was piloted before main administration; refinements documented.

Pilot aspect	Result
Participants	$n = 30$ purposive (8 clinicians, 6 administrators, 7 AI engineers, 5 patients/caregivers, 4 regulatory officers)
Method	60-minute structured think-aloud session per participant; cognitive interview methodology
Items reworded	11 items reworded for clarity (technical jargon removed: “SHAP attribution” → “feature explanation”; “RAG output” → “AI explanation grounded in medical literature”)
Items removed	3 items removed due to overlap with adjacent items (redundancy with Trust items)
Items added	2 items added to address gaps surfaced in interviews (auditability, rollback governance)
Completion time	Mean 22 minutes (target 20–25 minutes)
Content Validity Index (CVI)	≥ 0.80 for all retained items, from 5 expert reviewers (3 clinicians, 1 AI ethicist, 1 governance specialist)
Preliminary reliability	Cronbach’s $\alpha = 0.78$ – 0.92 across constructs (sufficient for proceeding to main study)
Pilot data treatment	Excluded from main analysis to avoid contamination; main study used separate respondent pool

Table 5.109. Reliability and Validity Strategy for the Instrument. Pre-registered thresholds and the resulting Chapter 4 evidence.

Property	Method	Threshold
Internal consistency	Cronbach’s α per construct [81]	≥ 0.70
Composite reliability	CR from PLS-SEM measurement model	≥ 0.70
Convergent validity	Average Variance Extracted (AVE) per construct [78]	≥ 0.50
Discriminant validity	Heterotrait–Monotrait ratio (HTMT) [79]	< 0.85
Multicollinearity	Variance Inflation Factor (VIF) for each item	< 5.0
Predictive relevance	Stone–Geisser Q^2 for endogenous constructs	> 0
Common method bias	Harman’s single-factor test; marker-variable technique	First factor $< 50\%$
Inter-rater (qualitative)	Krippendorff’s α for thematic coding [82]	≥ 0.67

Note. All achieved metrics on the $n = 131$ main-study sample are reported in Chapter 4 Table 4.3.

Likert Scale Definition and Reverse-Coded Items

Primary response scale (Sections B–H): 5-point Likert.

- 1 = Strongly Disagree
- 2 = Disagree
- 3 = Neutral
- 4 = Agree
- 5 = Strongly Agree

Reverse-coded items are included to detect inattentive responding and reduce acquiescence bias. Example: “I would *hesitate* to rely on AI-generated EEG recommendations even after seeing the explanation.” Reverse-coded items are inverted before scale-score computation. Detected inattentive respondents (failing ≥ 2 attention-check items) are excluded from the analytic sample.

Ethical Statement and Participant Protections

The following protections apply to every respondent:

- **Voluntary participation:** No compensation tied to specific responses; respondents may decline or withdraw at any time without consequence.
- **Anonymization:** Responses are stored without name, IP address, or institutional identifier; only stakeholder-strata demographic variables are retained.
- **Confidentiality:** Aggregate analysis only; no individual quote is attributable to a named respondent unless explicit re-consent for direct quotation is provided.
- **Withdrawal rights:** Respondents may withdraw their data within 30 days of submission by email request using their anonymized response ID.
- **Decision-support framing:** The instrument explicitly states throughout that the AI system is positioned as clinician decision support, *not* autonomous diagnosis or replacement of clinical judgement.
- **IRB approval:** The study protocol requires full IRB review and approval before main administration; informed consent appears on page 1 of the survey.

Table 5.110. Why Each Measured Construct Matters in This Dissertation. Reminds the reader of the operational reason each construct is part of the instrument.

Construct	Why it is measured
RAI Governance	Establishes whether the organization has the audit, bias, HITL, rollback, monitoring, lineage, and access controls that clinicians can perceive and reason about
Explainability	Captures whether the AI system's reasoning is clinically intelligible (SHAP + RAG narrative + traceability) at the point of care
Trust	Measures clinician willingness to rely on the AI system in screening decisions (calibrated, not absolute, trust)
Adoption Intention	Behavioural intention to integrate AI-assisted screening into clinical workflow (intention, not observed behaviour)
AI Familiarity	Stratifies clinicians by prior AI exposure, enabling moderator analysis of who benefits most from explainability investment

The administered 82-item instrument (Sections A–H below) operationalizes these constructs in stakeholder-appropriate language. The full instrument follows.

The instrument assesses technology acceptance, organizational readiness, trust, clinical utility, and implementation barriers across the following five stakeholder groups:

- **Clinical practitioners:** neurologists, psychiatrists, radiologists, and primary care physicians involved in neurological diagnosis.
- **Healthcare administrators:** CIOs, CMIOs, and department heads responsible for technology adoption decisions.
- **AI/ML engineers and data scientists:** technical professionals who develop, validate, and deploy machine learning models in clinical settings.
- **Patients and end users:** individuals receiving neurological care and their caregivers.
- **Regulatory and compliance officers:** professionals overseeing FDA, HIPAA, GDPR, and institutional compliance for AI-based medical devices.

Items are grounded in the following four theoretical frameworks:

- **Technology Acceptance Model (TAM):** perceived usefulness and ease of use as determinants of behavioural intention to adopt (Davis, 1989).
- **Technology–Organization–Environment framework (TOE):** technological, organizational, and environmental contexts that shape adoption readiness (Tornatzky & Fleischer, 1990).

- **Resource-Based View (RBV):** the framework's ASD-focused capability as a valuable, rare, inimitable, and non-substitutable resource (Barney, 1991).
- **Design Science Research (DSR):** artifact evaluation criteria of utility, quality, and efficacy applied to clinical diagnostic tools (Hevner et al., 2004).

Instrument Administration Guidelines

Table 5.111 documents the administration protocol specifying the research design, target population, sampling strategy, and psychometric validation plan.

Table 5.111. Survey Instrument Administration Protocol

Parameter	Specification
Research design	Cross-sectional multi-stakeholder survey
Target population	Clinical practitioners, healthcare administrators, AI/ML engineers, patients, regulatory officers
Minimum sample size	$n \geq 250$ total (50 per stakeholder group; rule of 10:1 per indicator for SEM)
Sampling strategy	Stratified purposive: 30% clinicians, 20% administrators, 20% AI engineers, 15% patients, 15% regulatory
Administration mode	Online (Qualtrics/REDCap) with institutional email distribution and professional society outreach
Estimated completion	20–25 minutes
Likert scale	5-point: 1 = Strongly Disagree, 2 = Disagree, 3 = Neutral, 4 = Agree, 5 = Strongly Agree
Analysis method	PLS-SEM (SmartPLS 4) with multi-group analysis by stakeholder role
Validation protocol	Pilot test ($n = 30$), expert panel review ($k = 5$), CFA, reliability ($\alpha \geq .70$, CR $\geq .70$, AVE $\geq .50$)
IRB requirement	Full IRB approval required (human subjects survey); informed consent on page 1
Total items	82 items across 8 sections (A–H)

Note. SEM = structural equation modelling; PLS = partial least squares; CFA = confirmatory factor analysis; AVE = average variance extracted; CR = composite reliability; IRB = Institutional Review Board. Minimum sample size follows Hair et al. (2021) guidelines for PLS-SEM with five or more latent constructs.

Theoretical Framework Mapping

Table 5.112 maps each survey section to its theoretical grounding, target constructs, and the stakeholder groups for which items are most relevant.

Table 5.112. Theoretical Framework Mapping for Survey Sections

Section	Theory	Constructs Measured	Primary Stakeholders
A: Demographics	—	Control variables, stratification	All groups
B: Technology Acceptance	TAM (Davis, 1989)	Perceived Usefulness, Perceived Ease of Use, Behavioural Intention	Clinicians, administrators, AI engineers
C: Organizational Readiness	TOE (Tornatzky & Fleischer, 1990)	Technology context, Organization context, Environment context	Administrators, clinicians, regulatory officers
D: Trust and Governance	TAM + Institutional Trust	Algorithmic trust, data governance, ethical considerations, regulatory compliance	All groups
E: Clinical Utility	DSR (Hevner et al., 2004)	Diagnostic performance, workflow integration, ASD-focused value	Clinicians, patients
F: RAG System Evaluation	DSR + RBV (Barney, 1991)	Report quality, explainability, hallucination mitigation	Clinicians, AI engineers
G: Implementation Barriers	TOE + RBV	Adoption barriers, regulatory requirements, cost–benefit	All groups
H: Wearable Integration	TAM + DSR	Wearable acceptance, sensor fusion, patient experience	Clinicians, patients, AI engineers

Note. TAM = Technology Acceptance Model; TOE = Technology–Organization–Environment; RBV = Resource-Based View; DSR = Design Science Research; RAG = Retrieval-Augmented Generation. All stakeholder groups complete Sections A, D, and G. Sections B, C, E, F, and H are administered based on role relevance, though all respondents may complete all sections.

Likert Scale Definition

Unless otherwise indicated, all Likert-scale items in Sections B–F and H use the following five-point response anchors:

Table 5.113 defines the five-point Likert scale response anchors used throughout Sections B–F and H of the instrument.

Table 5.113. Likert Scale Response Anchors

1	2	3	4	5
Strongly Disagree	Disagree	Neutral	Agree	Strongly Agree

Note. Respondents select one response per item. “Neutral” indicates neither agreement nor disagreement.

A “Not Applicable” option is available for items outside the respondent’s domain of expertise.

Section A: Demographics and Professional Background

Instructions: Please answer the following questions about your professional background. This information will be used for stratified analysis only and will not be linked to your identity. Select one response per question.

Table 5.114 presents the ten demographic and professional background items used for respondent stratification and multi-group analysis.

Section B: Technology Acceptance (TAM)

Instructions: The following statements relate to your perceptions of an AI-driven neurological disease detection system (the RGAIG framework) that uses machine learning to analyze EEG signals, medical imaging, and wearable sensor data to assist in the diagnosis of Autism Spectrum Disorder (ASD). Please indicate your level of agreement with each statement using the 5-point scale.

Table 5.114. Section A: Demographic and Professional Background Items (A1–A10)

ID	Question	Response Options
A1	What is your primary professional role?	☐ Neurologist ☐ Psychiatrist ☐ Radiologist ☐ Primary care physician ☐ Healthcare administrator (CIO/CMIO/Dept. Head) ☐ AI/ML engineer or data scientist ☐ Patient or caregiver ☐ Regulatory or compliance officer ☐ Other (specify: _____)
A2	How many years of professional experience do you have in your current role?	☐ 0–2 years ☐ 3–5 years ☐ 6–10 years ☐ 11–20 years ☐ More than 20 years
A3	What type of institution do you primarily work in?	☐ Academic medical center ☐ Community hospital ☐ Private practice ☐ Government/VA hospital ☐ Technology company ☐ Regulatory agency ☐ Research institution ☐ Other (specify: _____)
A4	How would you rate your familiarity with artificial intelligence and machine learning in healthcare?	☐ No familiarity ☐ Basic awareness ☐ Moderate understanding ☐ Advanced proficiency ☐ Expert level
A5	How would you rate your familiarity with EEG (electroencephalography) data and its clinical interpretation?	☐ No familiarity ☐ Basic awareness ☐ Moderate understanding ☐ Advanced proficiency ☐ Expert level
A6	What is the approximate size of your department or team?	☐ Fewer than 10 ☐ 10–50 ☐ 51–200 ☐ 201–500 ☐ More than 500
A7	Have you previously used any AI-based diagnostic or clinical decision support tool in your practice?	☐ Yes, regularly (weekly or more) ☐ Yes, occasionally (monthly) ☐ Yes, but only in a pilot or trial ☐ No, but I am aware of such tools ☐ No, and I have no prior exposure
A8	In which geographic region is your institution located?	☐ North America ☐ Europe ☐ Asia-Pacific ☐ Middle East and Africa ☐ Latin America ☐ Other (specify: _____)
A9	What is the bed capacity or patient volume of your institution (if applicable)?	☐ Fewer than 100 beds ☐ 100–499 beds ☐ 500–999 beds ☐ 1,000 or more beds ☐ Not applicable
A10	Is Autism Spectrum Disorder (ASD) relevant to your professional practice?	☐ Yes ☐ No ☐ Other (specify: _____)

Note. Item A10 is multi-select; all other items are single-select. Demographic data are used for stratified multi-group analysis (MGA) in PLS-SEM. CIO = Chief Information Officer; CMIO = Chief Medical Information Officer; VA = Veterans Affairs.

B1–B6: Perceived Usefulness (PU)

Theoretical basis: TAM Perceived Usefulness—the degree to which a person believes that using the system would enhance job performance (Davis, 1989).

Table 5.115 reports the six Perceived Usefulness items measuring the degree to which respondents believe the AI diagnostic system would enhance clinical performance.

Table 5.115. Section B1–B6: Perceived Usefulness Items

ID	Statement	SD	D	N	A	SA
B1	Using an AI-driven diagnostic system would improve the accuracy of my neurological disease diagnoses compared to current clinical methods alone.	☐	☐	☐	☐	☐
B2	An AI system that can analyze EEG, imaging, and wearable sensor data simultaneously would save me significant diagnostic time relative to sequential manual analysis.	☐	☐	☐	☐	☐
B3	An AI framework focused on Autism Spectrum Disorder (ASD) screening would be more useful than generic neurological tools.	☐	☐	☐	☐	☐
B4	AI-generated clinical decision support (e.g., diagnostic probability scores, feature importance rankings) would enhance the quality of my clinical decision-making.	☐	☐	☐	☐	☐
B5	An AI diagnostic system would reduce my overall clinical workload by automating routine screening and preliminary analysis tasks.	☐	☐	☐	☐	☐
B6	AI-assisted diagnosis would reduce the rate of diagnostic errors (missed diagnoses, false positives) in neurological screening compared to unaided clinical assessment.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). ASD = Autism Spectrum Disorder; EEG = electroencephalography.

B7–B12: Perceived Ease of Use (PEOU)

Theoretical basis: TAM Perceived Ease of Use—the degree to which a person believes that using the system would be free of effort (Davis, 1989).

Table 5.116 presents the six Perceived Ease of Use items assessing the anticipated effort required to operate the AI diagnostic system in clinical practice.

Table 5.116. Section B7–B12: Perceived Ease of Use Items

ID	Statement	SD	D	N	A	SA
B7	I expect the user interface of an AI diagnostic system to be intuitive and easy to navigate without extensive technical training.	☐	☐	☐	☐	☐
B8	I believe I could become proficient in using the AI diagnostic system within a reasonable learning period (2–4 weeks of regular use).	☐	☐	☐	☐	☐
B9	Integration of the AI system with existing Electronic Health Record (EHR) systems would be straightforward and would not disrupt current clinical workflows.	☐	☐	☐	☐	☐
B10	The AI system should require minimal additional training beyond what is already available through standard continuing medical education (CME) programs.	☐	☐	☐	☐	☐
B11	The diagnostic outputs of the AI system (e.g., classification reports, probability scores, feature importance visualizations) would be easy for me to interpret and act upon clinically.	☐	☐	☐	☐	☐
B12	The AI system would respond quickly enough (within 30 seconds for EEG analysis, within 2 minutes for full multi-modal analysis) to fit within my clinical workflow.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). EHR = Electronic Health Record; CME = Continuing Medical Education.

B13–B15: Behavioural Intention to Adopt (BI)

Theoretical basis: TAM Behavioural Intention—the degree to which a person has formulated conscious plans to adopt the technology (Davis, 1989; Venkatesh et al., 2003).

Table 5.117 documents the three Behavioural Intention items capturing respondents' willingness to adopt and recommend the RGAIG framework.

Table 5.117. Section B13–B15: Behavioural Intention Items

ID	Statement	SD	D	N	A	SA
B13	I am willing to adopt an AI-driven neurological diagnostic system as a supplementary tool in my clinical practice within the next 12 months, provided it meets regulatory requirements.	☐	☐	☐	☐	☐
B14	I would recommend the RGAIG framework to colleagues and other healthcare professionals in my network as a useful diagnostic aid.	☐	☐	☐	☐	☐
B15	I intend to continue using an AI diagnostic system once adopted, assuming it consistently demonstrates clinical value and maintains diagnostic accuracy above 90%.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). The 90% accuracy threshold in B15 reflects the minimum acceptable performance identified in the RGAIG framework validation (Chapter 4).

Section C: Organizational Readiness (TOE Framework)

Instructions: The following statements assess your organization's readiness to adopt AI-driven diagnostic tools. Please respond based on your current institutional context. If you are a patient or individual user, please answer based on the healthcare institution where you receive care, or select "Not Applicable" where appropriate.

C1–C4: Technology Context

Theoretical basis: TOE Technology Context—the internal and external technologies relevant to the organization, including equipment and processes (Tornatzky & Fleischer, 1990).

Table 5.118 reports the four Technology Context items evaluating institutional computing infrastructure, data maturity, system interoperability, and scalability.

Table 5.118. Section C1–C4: Technology Context Items

ID	Statement	SD	D	N	A	SA
C1	My organization has the computing infrastructure (GPU servers, cloud resources, network bandwidth) necessary to deploy AI-based diagnostic systems at clinical scale.	☐	☐	☐	☐	☐
C2	My organization has sufficient high-quality, digitized clinical data (EEG recordings, imaging data, structured clinical notes) to support AI model training and validation.	☐	☐	☐	☐	☐
C3	The existing clinical information systems at my organization (EHR, PACS, laboratory systems) are compatible with AI integration through standard APIs and data interoperability protocols (HL7 FHIR, DICOM).	☐	☐	☐	☐	☐
C4	My organization has the ability to scale AI deployments from a pilot (single department) to enterprise-wide adoption across multiple clinical departments.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). GPU = Graphics Processing Unit; EHR = Electronic Health Record; PACS = Picture Archiving and Communication System; HL7 FHIR = Health Level Seven Fast Healthcare Interoperability Resources; DICOM = Digital Imaging and Communications in Medicine; API = Application Programming Interface.

C5–C8: Organization Context

Theoretical basis: TOE Organization Context—internal characteristics including management structure, communication processes, human resources, and slack (Tornatzky & Fleischer, 1990).

Table 5.119 presents the four Organization Context items measuring leadership support, budget allocation, training capacity, and change management readiness.

Table 5.119. Section C5–C8: Organization Context Items

ID	Statement	SD	D	N	A	SA
C5	Senior leadership at my organization (C-suite, department heads) actively champions and supports the adoption of AI-based clinical tools.	☐	☐	☐	☐	☐
C6	My organization has allocated or is willing to allocate sufficient budget for AI diagnostic tool procurement, implementation, and ongoing maintenance.	☐	☐	☐	☐	☐
C7	My organization has the capacity to provide adequate training programs for clinicians on how to use and interpret AI-assisted diagnostic tools.	☐	☐	☐	☐	☐
C8	My organization has effective change management processes to facilitate the transition from traditional diagnostic workflows to AI-augmented workflows.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5).

C9–C12: Environment Context

Theoretical basis: TOE Environment Context—external factors including industry structure, regulatory environment, and competitive pressures (Tornatzky & Fleischer, 1990).

Table 5.120 documents the four Environment Context items assessing regulatory pressure, competitive dynamics, patient expectations, and industry standards influencing AI adoption.

Table 5.120. Section C9–C12: Environment Context Items

ID	Statement	SD	D	N	A	SA
C9	Evolving regulatory requirements (FDA SaMD guidelines, EU MDR, EU AI Act) are creating pressure for my organization to adopt validated AI diagnostic tools rather than informal or ad hoc approaches.	☐	☐	☐	☐	☐
C10	Competing healthcare institutions in my region are already adopting or piloting AI-based diagnostic tools, creating competitive pressure for my organization to do the same.	☐	☐	☐	☐	☐
C11	Patients increasingly expect or request that their healthcare providers use the latest AI-assisted diagnostic technologies.	☐	☐	☐	☐	☐
C12	Established industry standards and best practice guidelines (e.g., NIST AI RMF, WHO AI guidance, AAN clinical practice guidelines) support the adoption of AI in neurological diagnosis at my institution.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). FDA = US Food and Drug Administration; SaMD = Software as a Medical Device; EU MDR = European Union Medical Device Regulation; EU AI Act = European Union Artificial Intelligence Act; NIST AI RMF = National Institute of Standards and Technology AI Risk Management Framework; WHO = World Health Organization; AAN = American Academy of Neurology.

Section D: Trust and Governance

Instructions: The following statements assess your level of trust in AI diagnostic systems and your confidence in governance mechanisms that oversee their deployment. Please indicate your level of agreement with each statement. These items are relevant to all stakeholder groups.

D1–D3: Algorithmic Trust

Theoretical basis: Trust in automation (Lee & See, 2004) and algorithmic trust (Shin, 2021)—confidence in the system’s competence, benevolence, and integrity.

Table 5.121 presents the three Algorithmic Trust items measuring confidence in AI system reliability, transparency, and traceability.

Table 5.121. Section D1–D3: Algorithmic Trust Items

ID	Statement	SD	D	N	A	SA
D1	I have confidence that an AI system achieving $\geq 90\%$ accuracy on validated datasets would provide reliable diagnostic support in real-world clinical settings.	☐	☐	☐	☐	☐
D2	I trust AI diagnostic systems more when they provide transparent explanations of how they arrived at each diagnostic recommendation (e.g., feature importance via SHAP values, attention maps).	☐	☐	☐	☐	☐
D3	The ability to trace an AI diagnosis back to specific input features and supporting literature (explainability) is essential for me to trust the system’s recommendations.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). SHAP = SHapley Additive exPlanations.

D4–D6: Data Governance

Theoretical basis: Institutional trust theory (Zucker, 1986) and information privacy concerns (Smith et al., 1996)—confidence in organizational data handling practices.

Table 5.122 reports the three Data Governance items assessing confidence in patient data privacy, consent management, and de-identification procedures.

Table 5.122. Section D4–D6: Data Governance Items

ID	Statement	SD	D	N	A	SA
D4	I am confident that AI diagnostic systems deployed at my institution would adequately protect patient data privacy in accordance with HIPAA, GDPR, and applicable local regulations.	☐	☐	☐	☐	☐
D5	I believe that current consent management practices (informed consent for AI-assisted diagnosis, opt-in/opt-out mechanisms) are sufficient to protect patient autonomy in AI-driven care.	☐	☐	☐	☐	☐
D6	I am confident that the de-identification and anonymization procedures applied to clinical datasets used for AI training are robust enough to prevent re-identification of individual patients.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). HIPAA = Health Insurance Portability and Accountability Act; GDPR = General Data Protection Regulation.

D7–D9: Ethical Considerations

Theoretical basis: Responsible AI principles (Floridi et al., 2018)—fairness, accountability, transparency, and human oversight in automated decision-making.

Table 5.123 documents the three Ethical Considerations items addressing demographic bias testing, human decision authority, and patient notification rights.

Table 5.123. Section D7–D9: Ethical Considerations Items

ID	Statement	SD	D	N	A	SA
D7	I am confident that the AI diagnostic system has been tested for demographic bias (across age, gender, ethnicity) and performs equitably across diverse patient populations.	☐	☐	☐	☐	☐
D8	It is essential that a qualified human clinician retains final decision-making authority over all AI-generated diagnoses, even when the system achieves high accuracy.	☐	☐	☐	☐	☐
D9	Patients should be fully informed when AI tools are used in their diagnostic process and should have the right to request a purely human-performed diagnosis.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5).

D10–D12: Regulatory Compliance

Theoretical basis: Institutional theory (DiMaggio & Powell, 1983) and regulatory compliance frameworks—organizational conformity to external regulatory mandates governing AI in healthcare.

Table 5.124 presents the three Regulatory Compliance items measuring alignment with FDA SaMD requirements, HIPAA policies, and EU AI Act obligations.

Table 5.124. Section D10–D12: Regulatory Compliance Items

ID	Statement	SD	D	N	A	SA
D10	I believe the RGAIG framework’s quality assurance processes (525-module taxonomy with 97.4–98.4% pass rates) are sufficient to meet FDA Software as a Medical Device (SaMD) regulatory requirements.	☐	☐	☐	☐	☐
D11	My organization has the policies and procedures in place to ensure HIPAA compliance when deploying AI diagnostic tools that process protected health information (PHI).	☐	☐	☐	☐	☐
D12	The RGAIG framework’s governance structure (10 cross-cutting AI governance pillars, responsible AI protocols) aligns with the requirements of the EU AI Act for high-risk AI systems in healthcare.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). FDA = US Food and Drug Administration; SaMD = Software as a Medical Device; PHI = Protected Health Information; HIPAA = Health Insurance Portability and Accountability Act; EU AI Act = European Union Artificial Intelligence Act.

Section E: Clinical Utility Assessment (DSR Evaluation)

Instructions: The following statements assess the clinical utility of the RGAIG framework based on Design Science Research (DSR) evaluation criteria: utility, quality, and efficacy. Clinicians should respond based on their direct diagnostic experience; non-clinicians should respond based on their understanding of clinical needs within their professional role.

E1–E3: Diagnostic Performance (DSR Efficacy)

Theoretical basis: DSR artifact evaluation—efficacy criterion (Hevner et al., 2004). Does the artifact achieve the intended outcomes in the target environment?

Table 5.125 reports the three Diagnostic Performance items evaluating whether the framework’s reported accuracy, sensitivity, and specificity meet clinical adoption thresholds.

Table 5.125. Section E1–E3: Diagnostic Performance Items

ID	Statement	SD	D	N	A	SA
E1	The reported ASD-classification accuracy of 97.67% meets the minimum performance threshold I would require before considering clinical adoption.	☐	☐	☐	☐	☐
E2	The system’s sensitivity (ability to correctly identify patients with the disease) is adequate for use as a clinical screening tool, given that missed diagnoses carry significant patient harm.	☐	☐	☐	☐	☐
E3	The system’s specificity (ability to correctly identify patients without the disease) is adequate to avoid excessive false-positive referrals, which could burden specialist clinics and cause patient anxiety.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). Accuracy figures are from the RGAIG framework validation (Chapter 4, Table 4.X). ASD = Autism Spectrum Disorder.

E4–E6: Workflow Integration (DSR Utility)

Theoretical basis: DSR artifact evaluation—utility criterion (Hevner et al., 2004; March & Smith, 1995). Does the artifact provide value in the organizational context?

Table 5.126 presents the three Workflow Integration items assessing report quality, turnaround time, and clinical actionability of AI-generated diagnostic outputs.

Table 5.126. Section E4–E6: Workflow Integration Items

ID	Statement	SD	D	N	A	SA
E4	The AI-generated diagnostic reports (including classification results, feature importance, and literature-grounded explanations) would be of sufficient quality and clarity to support my clinical documentation requirements.	☐	☐	☐	☐	☐
E5	The system's expected turnaround time (under 2 minutes for a complete multi-modal analysis) would be fast enough to integrate into my standard clinical appointment workflow.	☐	☐	☐	☐	☐
E6	The diagnostic outputs of the AI system would be clinically actionable—that is, they would lead directly to specific follow-up actions (referral, treatment initiation, further testing) without requiring additional interpretation.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5).

E7–E10: ASD-Focused Value (DSR Quality)

Theoretical basis: DSR artifact evaluation—quality criterion (Hevner et al., 2004). RBV value attribute—does the ASD-focused capability represent a valuable, rare, and non-substitutable resource (Barney, 1991)?

Table 5.127 documents the four ASD-Focused Value items measuring ASD-focused comorbidity detection, rare presentation identification, second-opinion utility, and educational value.

Table 5.127. Section E7–E10: ASD-Focused Value Items

ID	Statement	SD	D	N	A	SA
E7	A unified AI framework capable of screening across multiple neurological domains (rather than separate single-disease tools) would provide significant value for ASD-focused comorbidity detection.	☐	☐	☐	☐	☐
E8	The AI system’s ability to detect rare or atypical presentations of neurological conditions would be a valuable supplement to clinical expertise, particularly in settings with limited specialist access.	☐	☐	☐	☐	☐
E9	I would value AI-generated diagnostic results as a “second opinion” to complement my own clinical judgment, particularly for complex or ambiguous cases.	☐	☐	☐	☐	☐
E10	The RGAIG framework would serve as a valuable training and education tool for medical residents and junior clinicians learning neurological diagnosis.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). RGAIG = Responsible Generative AI Governance.

Section F: Retrieval-Augmented Generation (RAG) System Evaluation

Instructions: The RGAIG framework includes a Retrieval-Augmented Generation (RAG) engine that produces natural-language diagnostic explanations grounded in peer-reviewed medical literature. The RAG system retrieves relevant publications from a curated knowledge base and generates explanatory narratives that cite specific studies. Please indicate your level of agreement with the following statements about this component.

F1–F3: Report Quality (RBV Value and Rarity)

Theoretical basis: RBV—the RAG-generated reports represent a valuable, rare, and difficult-to-imitate resource that provides competitive advantage (Barney, 1991). DSR quality evaluation (Hevner et al., 2004).

Table 5.128 reports the three RAG Report Quality items evaluating factual faithfulness, citation verifiability, and clinical relevance of generated explanations.

Table 5.128. Section F1–F3: RAG Report Quality Items

ID	Statement	SD	D	N	A	SA
F1	The RAG-generated diagnostic explanations would be factually faithful—that is, the content of the explanation would accurately reflect the underlying AI classification results and not introduce unsupported claims.	☐	☐	☐	☐	☐
F2	The literature citations provided in RAG-generated reports (specific author names, publication years, journal references) would be verifiable and accurately attributed.	☐	☐	☐	☐	☐
F3	The RAG-generated explanations would be clinically relevant—that is, they would reference literature and findings that are directly pertinent to the patient’s specific diagnostic context.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). RAG = Retrieval-Augmented Generation.

F4–F6: Explainability Value (RBV Inimitability)

Theoretical basis: RBV inimitability—the combination of SHAP/LIME feature importance with RAG-generated literature grounding is a complex, path-dependent capability that is difficult for competitors to replicate (Barney, 1991). Explainable AI (XAI) principles (Arrieta et al., 2020).

Table 5.129 presents the three Explainability Value items measuring the trust impact of SHAP/LIME visualisations, literature grounding, and confidence calibration.

Table 5.129. Section F4–F6: Explainability Value Items

ID	Statement	SD	D	N	A	SA
F4	Feature importance visualizations (SHAP values, LIME explanations) that show which EEG channels, frequency bands, or imaging features contributed most to the diagnosis would significantly increase my trust in the AI output.	☐	☐	☐	☐	☐
F5	Grounding AI diagnostic explanations in specific peer-reviewed literature (rather than generic statistical summaries) would make the output more credible and useful for clinical communication with patients and colleagues.	☐	☐	☐	☐	☐
F6	The system's reported confidence scores (probability estimates for each diagnostic category) would need to be well-calibrated—that is, a 90% confidence score should correspond to approximately 90% actual accuracy—for me to rely on them clinically.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly

Agree (5). SHAP = SHapley Additive exPlanations; LIME = Local Interpretable Model-agnostic Explanations; XAI = Explainable Artificial Intelligence.

F7–F8: Hallucination Concerns

Theoretical basis: Generative AI hallucination risk (Ji et al., 2023)—the tendency of large language models to generate plausible but factually incorrect content, which poses unique risks in clinical settings.

Table 5.130 documents the two Hallucination Concern items assessing the extent to which fabricated citations and incorrect claims undermine trust and whether the RGAIG mitigation strategies are perceived as adequate.

Table 5.130. Section F7–F8: Hallucination Concern Items

ID	Statement	SD	D	N	A	SA
F7	The risk of AI-generated hallucinations (fabricated citations, incorrect clinical claims, spurious diagnostic rationales) is a significant barrier to my trust in RAG-generated diagnostic reports.	☐	☐	☐	☐	☐
F8	I am confident that the RGAIG framework’s hallucination mitigation strategies (retrieval grounding, citation verification, confidence thresholds, 13-phase GenAI quality assurance) are adequate to reduce hallucination risk to clinically acceptable levels.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). RAG = Retrieval-Augmented Generation. Item F7 is reverse-scored: higher agreement indicates greater concern (lower trust). The 13-phase GenAI quality assurance taxonomy (220 modules, 98.1% pass rate) is detailed in Chapter 3.

Section G: Implementation Barriers and Adoption Requirements

Instructions: This section assesses the barriers, requirements, and expectations that would need to be addressed before you would support or participate in the clinical deployment of an AI-driven neurological diagnostic system. Please provide thoughtful responses; these qualitative data will be analyzed thematically to complement the quantitative findings.

Table 5.131 presents the seven Implementation Barriers items combining ranked, categorical, and open-ended response formats to capture adoption obstacles, regulatory requirements, accuracy thresholds, and privacy concerns.

Section H: Wearable Device Integration and Remote Monitoring

Instructions: The RGAIG framework supports integration with consumer-grade wearable EEG devices (e.g., Emotiv Insight, Emotiv EPOC X) and multi-modal wearable sensor arrays (ECG, PPG, EDA, IMU, temperature) for continuous neurological monitoring outside clinical settings. Please indicate your level of agreement with the following statements about this wearable integration capability.

H1–H3: Consumer Wearable EEG Acceptance

Theoretical basis: TAM applied to wearable health technology (Kim & Park, 2012)—perceived usefulness and ease of use of consumer wearable devices for clinical data collection.

Table 5.132 reports the three Consumer Wearable EEG Acceptance items assessing perceptions of data quality from consumer-grade devices, longitudinal monitoring value, and remote care continuity.

Table 5.132. Section H1–H3: Consumer Wearable EEG Acceptance Items

ID	Statement	SD	D	N	A	SA
H1	Consumer-grade EEG devices (e.g., Emotiv Insight with 5 channels) can provide data of sufficient quality for AI-driven neurological screening, despite having fewer channels than clinical-grade systems (64–256 channels).	☐	☐	☐	☐	☐
H2	Enabling patients to collect EEG data at home using consumer wearable devices would be valuable for longitudinal monitoring of neurological conditions between clinical visits.	☐	☐	☐	☐	☐
H3	Remote monitoring through wearable devices, with AI-analyzed results transmitted to the clinical team in near-real-time, would improve continuity of care for patients with chronic neurological conditions.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). EEG = electroencephalography.

Table 5.131. Section G: Implementation Barriers and Adoption Requirements (G1–G7)

ID	Question	Response Format
G1	Please rank the top 3 barriers to adopting an AI-driven neurological diagnostic system at your institution from the following list: (a) insufficient evidence of clinical efficacy, (b) data privacy and security concerns, (c) lack of regulatory approval (FDA/CE marking), (d) high implementation cost, (e) clinician resistance to change, (f) lack of IT infrastructure, (g) liability and malpractice uncertainty, (h) workflow disruption, (i) algorithmic bias concerns, (j) lack of explainability.	Rank 1st, 2nd, 3rd: 1st: _____ 2nd: _____ 3rd: _____
G2	What specific regulatory approvals or certifications would your institution require before deploying an AI diagnostic tool in clinical settings? (Select all that apply and/or specify others.)	☐ FDA 510(k) clearance ☐ FDA De Novo classification ☐ CE marking (EU MDR) ☐ IRB approval for clinical use ☐ Institutional clinical validation study ☐ Professional society endorsement (e.g., AAN) ☐ Other: _____
G3	What is the minimum diagnostic accuracy threshold (overall percentage) that an AI system must achieve before you would consider it suitable for clinical adoption as a screening or decision-support tool?	☐ $\geq 80\%$ ☐ $\geq 85\%$ ☐ $\geq 90\%$ ☐ $\geq 95\%$ ☐ $\geq 99\%$ Rationale (optional): _____
G4	What is the maximum acceptable training duration for clinicians to become proficient in using the AI diagnostic system?	☐ Less than 1 day ☐ 1–3 days ☐ 1–2 weeks ☐ 2–4 weeks ☐ More than 1 month
G5	Please describe your primary data privacy concern(s) regarding the use of AI systems that process sensitive neurological data (EEG, brain imaging, clinical records).	Open-ended: _____ _____ _____
G6	Within what timeframe would you expect an AI diagnostic system to demonstrate a positive return on investment (ROI) for your institution?	☐ Within 6 months ☐ 6–12 months ☐ 1–2 years ☐ 2–3 years ☐ More than 3 years ☐ ROI is not a primary consideration
G7	Please provide any additional comments, concerns, or suggestions regarding the adoption of AI-driven diagnostic tools for neurological disease detection.	Open-ended: _____ _____ _____ _____

Note. Items G1, G3, G4, and G6 are categorical or ranked. Items G2 is multi-select. Items G5 and G7 are open-ended and will be analyzed using reflexive thematic analysis (Braun & Clarke, 2006). FDA = US Food and Drug Administration; CE = Conformité Européenne; EU MDR = European Union Medical Device Regulation; IRB = Institutional Review Board; AAN = American Academy of Neurology; ROI = return on investment.

H4–H6: Sensor Fusion Value

Theoretical basis: RBV non-substitutability—the combination of multi-modal sensor data (EEG + ECG + EDA + accelerometry) creates a complex, non-substitutable diagnostic resource (Barney, 1991). DSR utility evaluation (Hevner et al., 2004).

Table 5.133 presents the three Sensor Fusion Value items measuring the perceived diagnostic benefit of multi-modal sensor combination, continuous monitoring, and configurable alert thresholds.

Table 5.133. Section H4–H6: Sensor Fusion Value Items

ID	Statement	SD	D	N	A	SA
H4	Combining data from multiple sensor modalities (EEG, ECG, EDA, accelerometry, temperature) would provide richer diagnostic information than any single data source alone.	☐	☐	☐	☐	☐
H5	Continuous wearable monitoring (24/7 data collection with AI analysis) would enable earlier detection of disease progression or treatment response than periodic clinic-based assessments.	☐	☐	☐	☐	☐
H6	Configurable alert thresholds (e.g., automatic clinician notification when EEG patterns deviate beyond a defined baseline by more than 2 standard deviations) would be a valuable safety feature for remote patient monitoring.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5). EEG = electroencephalography; ECG = electrocardiography; EDA = electrodermal activity; IMU = Inertial Measurement Unit.

H7–H8: Patient Experience

Theoretical basis: TAM Perceived Ease of Use extended to patient-facing technology (Holden & Karsh, 2010)—patient comfort, compliance, and willingness to use wearable health monitoring devices.

Table 5.134 documents the two Patient Experience items assessing wearable comfort for extended use and patient willingness to comply with prescribed monitoring protocols.

Table 5.134. Section H7–H8: Patient Experience Items

ID	Statement	SD	D	N	A	SA
H7	Consumer-grade wearable EEG devices are comfortable enough for patients (including elderly patients and children with neurological conditions) to wear for extended periods (4 or more hours per day).	☐	☐	☐	☐	☐
H8	Patients (or their caregivers) would be willing to comply with a prescribed wearable monitoring protocol (e.g., daily 2-hour EEG recording sessions over a 4-week period) if they understood it would improve their diagnostic accuracy and treatment monitoring.	☐	☐	☐	☐	☐

Note. SD = Strongly Disagree (1); D = Disagree (2); N = Neutral (3); A = Agree (4); SA = Strongly Agree (5).

Construct–Question–Hypothesis–Analysis Mapping

Table 5.135 provides a complete traceability matrix linking each survey item to its parent construct, the theoretical framework, the research hypothesis it informs, and the planned statistical analysis method.

Survey Item Summary and Psychometric Targets

Table 5.136 provides a summary count of items per section and the target psychometric properties for each multi-item scale.

Proposed Statistical Analysis Plan

Table 5.137 outlines the complete statistical analysis sequence for the survey data, from data screening through hypothesis testing and post-hoc analyses.

Stakeholder Routing Matrix

Table 5.138 specifies which survey sections are administered to each stakeholder group. All groups complete Sections A (Demographics), D (Trust and Governance), and G (Implementation Barriers). Remaining sections are routed based on role relevance to maximize response quality and minimize respondent fatigue.

Table 5.135. Survey Item Traceability: Question to Construct to Hypothesis to Analysis Method

Items	Construct	Theory	Hypothesis	Analysis Method
B1– B6	Perceived Usefulness (PU)	TAM	H1: $PU \rightarrow BI$	PLS-SEM path coefficient (β , $p < .05$)
B7– B12	Perceived Ease of Use (PEOU)	TAM	H2: $PEOU \rightarrow PU$; H3: $PEOU \rightarrow BI$	PLS-SEM path coefficient; mediation
B13– B15	Behavioural Intention (BI)	TAM	DV for H1–H3	R^2 , predictive relevance (Q^2)
C1– C4	Technology Context (TC)	TOE	H4: $TC \rightarrow$ Adoption readiness	PLS-SEM; MGA by institution type
C5– C8	Organization Context (OC)	TOE	H5: OC moderates $PU \rightarrow BI$	Moderation (interaction term)
C9– C12	Environment Context (EC)	TOE	H6: $EC \rightarrow$ Adoption urgency	PLS-SEM; control for region (A8)
D1– D3	Algorithmic Trust (AT)	Trust Theory	H7: AT mediates $PU \rightarrow BI$	Mediation (bootstrapped indirect effect)
D4– D6	Data Governance (DG)	Institutional Trust	H8: $DG \rightarrow AT$	PLS-SEM path coefficient
D7– D9	Ethical Considerations (ETH)	Responsible AI	H9: ETH moderates $AT \rightarrow BI$	Moderation analysis
D10– D12	Regulatory Compliance (RC)	Institutional Theory	H10: $RC \rightarrow OC$	PLS-SEM path coefficient
E1– E3	Diagnostic Performance (DP)	DSR Efficacy	H11: $DP \rightarrow PU$	PLS-SEM; IPMA
E4– E6	Workflow Integration (WI)	DSR Utility	H12: $WI \rightarrow PEOU$	PLS-SEM path coefficient
E7– E10	ASD-Focused Value (MDV)	DSR Quality + RBV	H13: $MDV \rightarrow PU$; H14: VRIN	PLS-SEM; RBV evaluation (descriptive)
F1–F3	RAG Report Quality (RQ)	RBV Value	H15: $RQ \rightarrow AT$	PLS-SEM path coefficient
F4–F6	Explainability Value (XV)	RBV Inimitability	H16: $XV \rightarrow AT$	PLS-SEM path coefficient
F7–F8	Hallucination Concerns (HC)	GenAI Risk	H17: $HC \rightarrow AT$ (negative)	PLS-SEM (F7 reverse-scored); f^2 effect size
G1– G7	Implementation Barriers	TOE + RBV	Qualitative triangulation	Thematic analysis (Braun & Clarke, 2006); frequency ranking
H1– H3	Wearable Acceptance (WA)	TAM (Wearable)	H18: $WA \rightarrow PU$	PLS-SEM path coefficient
H4– H6	Sensor Fusion Value (SF)	RBV Non-substitutable	H19: $SF \rightarrow MDV$	PLS-SEM path coefficient
H7– H8	Patient Experience (PX)	TAM (Patient)	H20: $PX \rightarrow WA$	PLS-SEM path coefficient; MGA by role (A1)

Note. TAM = Technology Acceptance Model; TOE = Technology–Organization–Environment; RBV = Resource-Based View; DSR = Design Science Research; PLS-SEM = Partial Least Squares Structural Equation Modelling; MGA = Multi-Group Analysis; VRIN = Valuable, Rare, Inimitable, Non-substitutable; β = standardised path coefficient; R^2 = coefficient of

Table 5.136. Survey Item Count and Psychometric Reliability Targets by Section

Section	Items	Response Type	Target α	Target AVE
A: Demographics and Professional Background	10	Categorical / multi-select	—	—
B1–B6: Perceived Usefulness	6	5-point Likert	$\geq .85$	$\geq .55$
B7–B12: Perceived Ease of Use	6	5-point Likert	$\geq .85$	$\geq .55$
B13–B15: Behavioural Intention	3	5-point Likert	$\geq .80$	$\geq .60$
C1–C4: Technology Context	4	5-point Likert	$\geq .80$	$\geq .50$
C5–C8: Organization Context	4	5-point Likert	$\geq .80$	$\geq .50$
C9–C12: Environment Context	4	5-point Likert	$\geq .80$	$\geq .50$
D1–D3: Algorithmic Trust	3	5-point Likert	$\geq .80$	$\geq .60$
D4–D6: Data Governance	3	5-point Likert	$\geq .80$	$\geq .60$
D7–D9: Ethical Considerations	3	5-point Likert	$\geq .80$	$\geq .60$
D10–D12: Regulatory Compliance	3	5-point Likert	$\geq .80$	$\geq .60$
E1–E3: Diagnostic Performance	3	5-point Likert	$\geq .80$	$\geq .60$
E4–E6: Workflow Integration	3	5-point Likert	$\geq .80$	$\geq .60$
E7–E10: ASD-Focused Value	4	5-point Likert	$\geq .80$	$\geq .50$
F1–F3: RAG Report Quality	3	5-point Likert	$\geq .80$	$\geq .60$
F4–F6: Explainability Value	3	5-point Likert	$\geq .80$	$\geq .60$
F7–F8: Hallucination Concerns	2	5-point Likert	$\geq .70$	$\geq .50$
G1–G7: Implementation Barriers	7	Ranked / categorical / open-ended	—	—
H1–H3: Wearable Acceptance	3	5-point Likert	$\geq .80$	$\geq .60$
H4–H6: Sensor Fusion Value	3	5-point Likert	$\geq .80$	$\geq .60$
H7–H8: Patient Experience	2	5-point Likert	$\geq .70$	$\geq .50$
Total	82			

Note. α = Cronbach's alpha; AVE = Average Variance Extracted. Targets follow Hair et al. (2021): $\alpha \geq .70$ (acceptable), $\geq .80$ (good); AVE $\geq .50$ (adequate convergent validity). Sections A and G are not multi-item Likert scales and do not have reliability targets. Two-item scales use Spearman–Brown prophecy coefficient instead of α .

Table 5.137. Complete Statistical Analysis Plan for Multi-Stakeholder Survey

#	Analysis Step	Purpose	Tool / Criterion
1	Data screening	Missing data pattern (MCAR test), outlier detection (Mahalanobis distance), normality (skewness < 2 , kurtosis < 7)	SPSS 29 / R; Little's MCAR test
2	Descriptive statistics	Sample profile by stakeholder group, response distributions, item means and standard deviations	SPSS 29
3	Exploratory Factor Analysis (EFA)	Dimensionality check on pilot data ($n = 30$); confirm factor structure before CFA	SPSS; KMO $\geq .70$, Bartlett's $p < .001$
4	Confirmatory Factor Analysis (CFA)	Validate measurement model; assess model fit and construct validity	SmartPLS 4; SRMR < .08, NFI > .90
5	Reliability analysis	Internal consistency: Cronbach's $\alpha \geq .70$; Composite Reliability CR $\geq .70$; rho_A $\geq .70$	SmartPLS 4
6	Convergent validity	Average Variance Extracted AVE $\geq .50$ for all reflective constructs	SmartPLS 4
7	Discriminant validity	Fornell–Larcker criterion (square root of AVE > inter-construct correlations); HTMT < .85	SmartPLS 4
8	Common method bias	Harman's single-factor test (< 50% variance on one factor); full collinearity VIF < 3.3	SPSS / SmartPLS
9	Structural model (PLS-SEM)	Path coefficients (β), significance ($p < .05$, 5,000 bootstraps), R^2 ($\geq .25$ moderate), f^2 effect sizes	SmartPLS 4
10	Mediation analysis	Bootstrapped indirect effects (5,000 resamples, 95% bias-corrected CI) for H7 (AT mediates PU $\rightarrow$ BI)	SmartPLS 4; specific indirect effects
11	Moderation analysis	Interaction terms for H5 (OC moderates PU $\rightarrow$ BI), H9 (ETH moderates AT $\rightarrow$ BI)	SmartPLS 4; product indicator approach

Table 5.138. Stakeholder Routing Matrix: Sections Administered by Respondent Role

Stakeholder Group	A	B	C	D	E	F	G	H	Items	Min.
Clinical practitioners	•	•	•	•	•	•	•	•	82	25
Healthcare administrators	•	•	•	•	–	–	•	–	56	17
AI/ML engineers	•	•	–	•	•	•	•	•	70	21
Patients and end users	•	–	–	•	•	–	•	•	47	15
Regulatory officers	•	–	•	•	–	–	•	–	41	14

Note. • = section administered; – = section not administered (skip logic in Qualtrics/REDCap). Items

= total questionnaire items administered per group. Min. = estimated completion time in minutes.

Ethical Considerations and Informed Consent

The following ethical safeguards apply to the administration of this survey instrument:

1. **IRB Approval.** Full Institutional Review Board (IRB) approval from Golden Gate University and each participating institution is required prior to data collection. The study is classified as minimal risk (online survey, no intervention, no vulnerable population beyond standard clinical research).
2. **Informed Consent.** The first page of the online survey presents a detailed informed consent form that describes: (a) the purpose of the study, (b) the voluntary nature of participation, (c) the estimated completion time (20–25 minutes), (d) the types of questions asked, (e) data confidentiality and anonymization procedures, (f) the right to withdraw at any time without penalty, and (g) contact information for the principal investigator and IRB. Respondents must actively select “I consent to participate” before accessing survey items.
3. **Anonymity and Confidentiality.** No personally identifiable information (PII) is collected. IP addresses are not logged. Institutional affiliation is collected at the category level (academic medical center, community hospital, etc.) only. All data are stored on an encrypted, access-controlled server in compliance with HIPAA and GDPR requirements.

4. **Data Retention.** Survey responses are retained for 7 years following publication, consistent with Golden Gate University research data retention policies. After the retention period, all electronic records are permanently deleted.
5. **Compensation.** Respondents may be offered a \$25 gift card or equivalent professional development credit as compensation for their time, subject to institutional policies. Compensation is provided regardless of survey completion to avoid coercion.
6. **Patient Respondents.** For patient participants (Stakeholder Group 4), the informed consent form includes additional language explaining: (a) the study is about their opinions on AI tools, not a clinical trial, (b) their responses will not affect their care, and (c) the survey uses simplified language (Flesch–Kincaid Grade Level ≤ 10) to ensure accessibility.

Pilot Testing Protocol

Prior to full deployment, the survey instrument undergoes a three-stage validation process:

1. **Stage 1: Expert Panel Review** ($k = 5$). Five subject-matter experts (2 neurologists, 1 healthcare informatics researcher, 1 psychometrician, 1 regulatory affairs specialist) review all items for content validity, clarity, and relevance. Items are retained if Content Validity Index (CVI) $\geq .80$ across all raters. Items with CVI $< .80$ are revised or eliminated.
2. **Stage 2: Cognitive Interviewing** ($n = 8$). Eight participants (2 per major stakeholder group) complete the survey while thinking aloud. The interviewer probes for: (a) comprehension difficulties, (b) ambiguous wording, (c) missing response options, and (d) survey fatigue. Items are revised based on cognitive interview findings.
3. **Stage 3: Pilot Deployment** ($n = 30$). Thirty participants complete the survey online. Pilot data are used to assess: (a) item distributions (ceiling/floor effects), (b) inter-item correlations ($r \geq .30$ within constructs), (c) exploratory factor struc-

ture, (d) completion time, and (e) dropout rates. Items with poor psychometric properties (item-total correlation $< .30$, cross-loading $> .40$) are revised or removed before full deployment.

This survey instrument is designed to complement the computational validation of the RGAIG framework presented in Chapters 4–5 by providing evidence from the *human adoption* perspective. While the dissertation demonstrates that the framework performs technically (ASD-focused classification accuracy of 91–100%), the survey will test whether the five stakeholder groups are *willing and able* to adopt it—a necessary condition for translating computational performance into real-world clinical impact.

Appendix G

Primary Data Collection

Instruments and Research Design

Overview

This appendix presents the primary data collection instruments designed for the qualitative strand of the convergent mixed-methods research design. An ASD-focused questionnaire captures behavioral, psychological, and social data from patients and caregivers, while the IoT/Emotiv secondary data specification documents the quantitative EEG data attributes.

Table 5.139 summarises the instruments, their target population, item count, and the research variables they operationalise.

Table 5.139. Primary Data Collection Instruments Overview

#	Instrument	Target	Items	Data Type	Variables Measured
I1	ASD Clinical Questionnaire	Parent / Care-giver	40	Primary (Qual.)	Behavioral severity, psychological profile, social impact, diagnostic delay, AI acceptance
I6	IoT/Emotiv Secondary Data Spec.	N/A (meta-data)	120+	Secondary (Quant.)	EEG spectral power, connectivity, time-domain, nonlinear, wavelet features

Instrument Structure

Each disease-specific questionnaire follows a standardised six-section structure:

- **Section A — Demographics** (8 items): Age, gender, location, family history, prior diagnosis, medications
- **Section B — Behavioral/Clinical Assessment** (10 items): Disease-specific observable behaviors and clinical signs (5-point Likert: Never–Always). Tagged as **Independent Variable: Behavioral Severity**.
- **Section C — Psychological Assessment** (10 items): Mental health comorbidities, emotional regulation, coping mechanisms (5-point Likert: Strongly Disagree–Strongly Agree). Tagged as **Independent Variable: Psychological Profile**.
- **Section D — Social Impact** (8 items): Employment, relationships, caregiver burden, financial impact, stigma (5-point Likert). Tagged as **Independent Variable: Social Functioning**.
- **Section E — Diagnostic Experience** (7 items): Time to diagnosis, specialists consulted, satisfaction, cost (mixed: open-ended + Likert). Tagged as **Dependent Variable: Diagnostic Delay**.
- **Section F — Technology Acceptance** (5 items): Willingness to use AI screening, trust in EEG monitoring, preference for AI-assisted reports (5-point Likert). Tagged as **Dependent Variable: Adoption Intent**.

Each instrument concludes with a **Variable Mapping Table** that links every section to its corresponding independent variable (IV), dependent variable (DV), mediating variable (MeV), or moderating variable (MoV), and maps each to the relevant hypothesis (H1–H5).

Master Variable Mapping Across All Instruments

Master IV/DV Mapping: Primary and Secondary Data. Table 5.140 below presents the detailed analysis.

Research Design (ASD)

Mixed-Methods Research Design Per Disease. Table 5.141 below presents the detailed analysis.

Table 5.141. Mixed-Methods Research Design Per Disease

Disease	Quantitative (Secondary Data)	Qualitative (Primary Data)	Integration Finding
ASD	VotingClassifier on EEG; Bootstrap 95% CI; Cohen's d=1.48; SHAP: theta/beta ratio #1	Parent questionnaire (n=10); Saldaña 3-stage coding; themes: diagnostic delay, caregiver burden	Convergent: AI screens in 5 min (97.67%) confirming parent-reported 3-year delay

Note. The individual questionnaire forms (Instruments I1–I5) and the IoT/Emotiv secondary data specification (Instrument I6) follow on the subsequent pages. Each instrument is self-contained with its own scoring rubric, variable mapping table, and consent section.

Table 5.140. Master IV/DV Mapping: Primary and Secondary Data

Type	ID	Variable	Source	Data	Scale	Hyp.
IV	S1	EEG spectral power (5 bands)	Emotiv / Clinical EEG	Secondary	Ratio	H1, H3
IV	S2	EEG connectivity (PLV)	Emotiv / Clinical EEG	Secondary	Ratio	H1, H3
IV	S3	HRV + EDA (stress)	Empatica E4	Secondary	Ratio	H1
IV	S4	GLCM texture (cancer)	TCIA histopathology	Secondary	Ratio	H1, H3
IV	P1	Behavioral severity score	Questionnaire Sec B	Primary	Ordinal	H1, H3
IV	P2	Psychological profile score	Questionnaire Sec C	Primary	Ordinal	H2
IV	P3	Social functioning score	Questionnaire Sec D	Primary	Ordinal	H4
IV	P4	Diagnostic delay (months)	Questionnaire Sec E	Primary	Ratio	H4
DV	D1	Classification accuracy (ASD ensemble)	ML pipeline output	Secondary	Ratio	H1
DV	D2	Explainability quality	Expert survey + RAGAS	Mixed	Ordinal	H2
DV	D3	SHAP biomarker match	SHAP analysis	Secondary	Nominal	H3
DV	D4	Effort / time reduction	Process comparison	Mixed	Ratio	H4
DV	D5	Adoption intent score	Questionnaire Sec F	Primary	Ordinal	H5
DV	D6	Governance pass rate	525-module testing	Secondary	Ratio	H5
MeV	M1	Trust in AI	Questionnaire F2	Primary	Ordinal	H2→H5
MeV	M2	Explainability	RAGAS + expert rating	Mixed	Ratio	H2→H5
MoV	V1	Disease domain	Dataset label	Secondary	Nominal	H1
MoV	V2	Geographic location	Questionnaire A4	Primary	Nominal	H4, H5
MoV	V3	Device type	Recording metadata	Secondary	Nominal	H1

Appendix H

RGAIK Agentic Framework Enterprise Architecture Specification

Architecture Appendix Defensibility Pack

Critical disclaimer (read first). The architecture documented in this appendix represents a *conceptual governance-enabled decision-support operating environment*, not a clinically deployed autonomous diagnostic platform. It is a forward-looking operational representation of how the Responsible AI Governance constructs investigated in this dissertation could be instantiated at enterprise scale. The architecture is **not** a production deployment blueprint; it is an artefact that makes the dissertation’s governance, explainability, trust, and adoption claims concrete and defensible. Production deployment in any healthcare setting would require additional clinical validation, regulatory clearance (FDA SaMD, EU AI Act conformity assessment), prospective safety evaluation, and site-specific governance review — none of which is in scope here.

Reading order. This appendix is organised in two layers:

- **Layer 1 (this pack):** Conceptual operational framework — how architecture supports the Governance → Explainability → Trust → Adoption pathway. Doctoral-relevant.
- **Layer 2 (subsequent sections):** Supporting technical context — containers, services, data flows, deployment topology. Engineering-relevant only; provided for completeness, not as dissertation contribution.

Table 5.142. Research-to-Architecture Mapping. Every architectural element in this appendix is justified by its role in supporting one of the dissertation’s research constructs.

Research element	Architecture representation (what it operationalizes)
Responsible AI Governance (IV)	Oversight layer: immutable audit-log service; bias-monitoring jobs; rollback registry; access-control gateway; lineage tracking; HITL escalation queue; drift-monitoring dashboards
Perceived Explainability (Mediator 1)	Explanation service: SHAP attribution generator + RAG citation grounding + clinician-UI reasoning panel. Renders explanations in clinical language, not algorithmic output
Clinician Trust (Mediator 2)	Clinician review workflow: override capability + rationale capture; trust-calibration feedback loop; audit of clinician overrides to detect over-reliance or under-reliance
Adoption Intention / Workflow Readiness (DV)	Workflow integration layer: EHR-side decision-support display (HL7-FHIR); inline workflow context; perceived-fit telemetry; training and onboarding support
Architecture supports H1–H6 by:	Making the Governance→Explainability→Trust→Adoption pathway concretely operationalisable. Without an architectural representation, the pathway is abstract; with one, the dissertation’s claims become deployment-relevant.

Note. Each architectural element below the pack is traceable to one of these research elements.

Engineering details that do not map to a research element are demoted to deep-appendix or future work.




Figure 5.27. Human-Centered Workflow. The clinician initiates and concludes every screening event; AI is a decision-support layer mediated by governance and explainability. The dashed feedback loop closes the trust-calibration cycle by feeding clinician overrides back into governance monitoring. This figure is the DBA-grade architectural reduction of the dissertation’s central pathway.




Figure 5.28. Data Governance Flow. EEG data flows through governance controls, then through the explainability layer, then to the clinician for review, finally to an EHR-integrated decision-support output. The architecture’s central role is to ensure every data movement is governed (lineage), every prediction is explainable (SHAP+RAG), and every decision is clinician-reviewed (HITL).

Table 5.143. Governance Maturity-to-Architecture Mapping. Each RGAIG Maturity level (Chapter 1 Section 1) corresponds to an architectural characteristic. Healthcare organizations advance levels by adding architectural capability, not by accumulating frameworks.

Lvl	Maturity stage	Architectural characteristic at this level
1	Initial	Ad-hoc AI scripts; no governance services; no audit trail; explainability absent
2	Managed	Basic logging; manual access control; isolated AI tools; no integration
3	Governed	Formal audit-log service; bias-monitoring jobs; rollback registry; documented control inventory
4	Explainable	Explanation service (SHAP + RAG) embedded in clinical UI; reasoning surfaced to clinician at point of care
5	Trusted	Clinician feedback loop + trust-calibration monitoring; fairness dashboards per patient subgroup; verified performance evidence
6	Operationalized	Workflow-integrated AI with EHR + HL7-FHIR; outcome monitoring; continuous-learning governance; cross-site replication

Note. The architecture is layered to support maturity progression: Levels 1–3 capabilities are foundational services; Levels 4–5 add clinical-UI and trust-monitoring services; Level 6 adds workflow + outcome integration. An organization adopts capabilities incrementally, paying only for what advances their next level.

Table 5.144. Governance Dimensions That Directly Support the Dissertation’s Hypotheses. Other engineering capabilities (caching, message queuing, observability tracing) exist in Layer 2 below but are not dissertation-relevant.

Governance area	Architectural service operationalizing it	Supports
Explainability	SHAP attribution generator + RAG citation grounding + clinical-UI reasoning panel	H1, H2
Auditability	Immutable audit-log service with per-decision retrieval API	H1, H4
Bias monitoring	Subgroup fairness dashboard + alerting on threshold breach	H1
HITL	Clinician override capability + rationale capture + escalation queue	H3, H5
Transparency	End-to-end lineage tracking (data → feature → model → decision)	H1, H2
Rollback	Model registry with version pinning and time-bounded revert path	H1
Traceability	Citation traceability from RAG output to peer-reviewed source	H2

Note. Engineering capabilities outside this table (e.g., service mesh, message queue, container orchestration, vector-DB internals, CI/CD pipelines) exist in the technical architecture below but are demoted to deep-appendix material; they do not constitute part of the dissertation contribution.

End of Layer 1 (Conceptual Architecture Pack). The technical-context detail that follows is provided for engineering completeness only. Examiners assessing the doctoral contribution should focus on Tables 5.142–5.144 and Figures 5.27–5.28 above; the remainder of this appendix is supplementary technical context.

Document Status and Scope

This appendix specifies the *target* enterprise architecture for deploying the RGAIG framework as an agentic operating system. It is **forward-looking** — the empirically validated reference implementation (`agenticfinder`; 305 GB processed data; 198 trained models; ASD ensemble 97.67%) is the current minimum-viable implementation. This appendix documents how that implementation extends toward enterprise deployment.

Why this belongs in the thesis. Chapter 3 specifies *what* RGAIG is (the seven-layer framework). Chapter 4 evidences *that it works* (lab-validated results). Chapter 5 discusses *what it implies* (managerial / policy / clinical impact). This appendix specifies *how it is operationalised at enterprise scale* — the runtime architecture that the framework deploys onto. Without it, the thesis ends at “validated framework” without addressing the inevitable examiner question: *how would a hospital actually run this?*

What this appendix is not:

- Not implementation code (deferred to post-submission engineering).
- Not a re-derivation of RGAIG (Chapter 3 owns that).
- Not a cost study (Chapter 5 § ROI owns that).
- Not benchmarks (Chapter 4 owns that).

Business Context and Stakeholders

Business Problem (Anchored to Chapter 1)

The RGAIG framework defined in Chapter 3 is a layered governance-aligned diagnostic system. Validating it in a controlled experimental setting (Chapter 4) demonstrates *technical feasibility*; deploying it across a hospital network demands a *runtime architecture* that satisfies clinical safety, regulatory compliance, multi-tenant isolation, financial governance, and 24×7 operational reliability. This appendix defines that runtime.

Stakeholder Concerns

Table 5.145 presents architectural Concerns by Stakeholder.

Table 5.145. Architectural Concerns by Stakeholder. This table maps each decision-maker to the architectural property they evaluate when assessing the framework for deployment, ensuring that the architecture specification addresses the full set of stakeholder requirements rather than optimising for a single audience.

Stakeholder	Architectural Concern
Clinician	Confidence calibration; explainability latency; clinical-safe error envelopes
Hospital administrator	Cost ceiling per screening; vendor lock-in avoidance; GPU utilisation
Operations manager	SLA; throughput; on-call runbooks
Policy / regulator	Audit-trail completeness; model card; fairness monitoring; EU AI Act Art. 86 right-to-explanation
Patient / data subject	Consent; data retention; override rights; fairness across demographics
AI / platform engineer	Deployability; observability; rollback; drift detection; safe model swap
Investor / funder	Unit economics; capacity model; time-to-deployment

Success Criteria

Table 5.146 presents architecture-Level Success Criteria.

Table 5.146. Architecture-Level Success Criteria. Each dimension carries a quantitative target derived from Chapter 5 evidence or industry benchmarks; collectively these criteria define what “production-grade” means for the RGAIG enterprise deployment.

Dimension	Target
Diagnostic accuracy (carried from Ch. 4)	$\geq 85\%$ per domain; ASD 97.67 % validated
End-to-end p95 latency	$\leq 4\text{ s}$ for routine; $\leq 250\text{ ms}$ for cached repeats
Fairness (disparate impact ratio)	≥ 0.80 across age $\times$ sex $\times$ ethnicity, monitored weekly
Audit-row completeness	100 % of clinical-decision invocations
Cost per screening	$\leq \$15$ (ROI ceiling per Ch. 5)
Recovery objective (RTO / RPO)	RTO 15 min / RPO 0 for clinical APIs
Compliance	EU AI Act high-risk; NIST AI RMF; FDA SaMD Class II; HIPAA; GDPR

Out of Scope

- Drug-discovery models (per delimitation Ch. 1)
- Robotic surgical guidance
- Generative-creative non-clinical use
- Edge-only deployment without cloud component (future work)

Architecture Principles

The RGAIG enterprise architecture extends the standard 12-factor application principles with five AI-specific extensions and four RGAIG-specific principles.

Table 5.147 presents architecture Principles: 12-Factor Foundation Plus AI and RGAIG Extensions.

Table 5.147. Architecture Principles: 12-Factor Foundation Plus AI and RGAIG Extensions. Standard application principles (1–12) provide a deployment baseline; the AI extensions (13–17) address requirements specific to systems that include trained models and prompts; the RGAIG-specific principles (18–21) encode the governance and clinical-safety commitments unique to this framework.

#	Principle	Implication
1–12	Standard 12-factor	Codebase / dependencies / config / backing services / build-release-run / processes / port-binding / concurrency / disposability / dev-prod parity / logs / admin-processes
13	Models as code-equivalent dependencies	Pinned, versioned, registered in model registry
14	Prompts as code	Versioned in prompt registry; redeploy = bump version
15	Data contracts	Schemas versioned + validated at every boundary
16	Evaluation as build artefact	Every model release ships eval + fairness + drift report
17	Cost as first-class metric	Token cost dashboards + alerts at 50 / 80 / 100 % budget
18	Decision audit before deploy (RGAIG)	Every clinical-decision path emits an audit row before output is returned
19	HITL gate at confidence (RGAIG) threshold	< 0.80 confidence $\rightarrow$ mandatory clinician override before commit
20	Explainability is a SLA, not a feature (RGAIG)	Local explanation (SHAP / counterfactual / citation) returned within $p95 \leq 1$ s of decision
21	Hub-and-spoke isolation per agent (RGAIG)	Council agents communicate ONLY through the hub; never agent-to-agent direct calls

C4 Architecture Levels 1 to 7

The architecture is documented across seven C4 levels: standard C4 (Levels 1–4) plus three RGAIG-specific extensions (Levels 5–7) for governance, observability, and lifecycle.

Level 1 — System Context

The system-context view (Figure 5.29) places the RGAIG enterprise framework within its operating ecosystem of users and external systems.




Figure 5.29. C4 Level 1 — System Context. The RGAIG framework as a single black-box system, with external actors (clinician, patient, administrator, regulator, researcher, payer) and external systems (EHR, wearable EEG, LLM provider, vector DB, audit archive). Internal structure is opened in Levels 2–4.

External-System Catalogue

Table 5.148 presents external Systems and Failure Modes.

Table 5.148. External Systems and Failure Modes. Each external system is characterised by direction of data flow, protocol, owning entity, and graceful-degradation strategy when the system is unavailable, ensuring the architecture can continue serving clinical decisions through expected partial outages.

System	Direction	Protocol	Owner	Failure Mode
EHR (Epic / Cerner)	Bi-directional	HL7 v2.x; FHIR R4 REST	Hospital IT	Read-only fallback to local cache
Wearable EEG (Emotiv)	Inbound	LSL stream; WebSocket	Device vendor	Drop session; record incident
LLM provider	Outbound	HTTPS API	Vendor / self-hosted	Multi-provider fallback (ADR-004)
Vector DB (Qdrant)	Bi-directional	gRPC	Self-hosted	Hot replica failover
Audit archive (S3)	Outbound write-only	HTTPS	Cloud provider	Local DLQ; replay on recovery
Identity provider	Outbound	OAuth 2.0 / SAML 2.0	Enterprise SSO	Cached tokens within TTL

Level 2 — Container Diagram

The container view decomposes the framework into 12 deployable containers organised by functional plane (Figure 5.30).




Figure 5.30. C4 Level 2 — Container Diagram. Twelve containers organised by functional plane: Edge (1–2), Control Plane (3–8), Council Runtime (9), Workflow Engine (10), Data / Knowledge (11), Observability and HITL (12). Hub-and-spoke topology routes all inter-agent communication through the Control Plane.

Container Responsibility Matrix

Table 5.149 documents container Responsibility Matrix.

Table 5.149. Container Responsibility Matrix. Each container is characterised by its primary responsibility, technology choice, statefulness, and accountable team, providing a single reference for capacity planning, on-call rotation, and ownership disputes.

#	Container	Primary Responsibility	Tech	Stateful?	Owner
1	API Gateway	Auth, rate-limit, mTLS termination	Kong / Envoy	No	Platform
2	WebSocket Gateway	Streaming inference + EEG live	FastAPI + uvloop	No	Platform
3	Orchestrator	Workflow routing; agent dispatch	Python (LangGraph)	Yes (state)	Agent team
4	Policy Engine	RBAC + ABAC + tool policy	OPA (Rego)	No	Security
5	Tool Registry	MCP tool catalogue, contracts	Python + Postgres	Yes	Agent team
6	Prompt Registry	Versioned prompt templates	Python + Git-backed	Yes	AI ops
7	Audit Logger	Decision row append-only writer	Python + Kafka	No (Kafka)	Compliance
8	Budget Manager	Token / cost ceilings per tenant	Python + Redis	Yes	FinOps
9	Council Runtime	15+ specialist agents	Python (LangGraph)	Yes (graph)	Agent team
10	Workflow Engine	LangGraph + Temporal	Python + Temporal	Yes	Platform
11	Data / Knowledge	Vector / SQL / graph / object	Qdrant / PG / Neo4j / S3	Yes	Data team
12	HITL Console	Clinician approval UI; override logger	React + Postgres	Yes	UX + Compliance

Level 3 — Component Detail (Council Runtime)

The Council Runtime (Container 9) hosts 15 specialist agents, each with a single responsibility. All agents communicate only through the central Orchestrator (Container 3) — never agent-to-agent.

Table 5.150 lists council Specialist Agents and HITL Triggers.

Table 5.150. Council Specialist Agents and HITL Triggers. Each agent has a single, narrow responsibility and a clear escalation rule for routing to human-in-the-loop review, ensuring no agent’s output bypasses clinical oversight when its confidence or safety thresholds are not met.

Agent	Single Responsibility	HITL Trigger
Planner	Decompose user goal → ordered task plan	If plan needs > 10 steps
RAG	Retrieve evidence chunks; rerank; cite	If retrieval ≤ 3 chunks
Tool	Select + invoke MCP tools	If tool not previously authorised
Coding	Generate code (sandboxed)	If runtime change to live config
Critic	Evaluate prior agent output	Always; reports score
Reflection	Self-critique; replan if quality low	If Critic score < 0.6
Explainability	Produce SHAP / counterfactual / citation	Always for clinical decisions
Responsible-AI	Bias / fairness / toxicity check	If DI < 0.80 or toxicity above threshold
Governance	Policy lookup; consent enforcement; retention check	If consent missing or expired
Security	Prompt-injection scan; output sanitisation	If injection detected
Performance	Latency / cost ceiling enforcement	If projected cost > budget
Evaluation	Score on internal eval set; quality gate	Always before commit
Drift	PSI / KS detection on input distribution	If PSI > 0.20
Provenance	Build lineage graph: source → chunk → token → output	Always; writes to Neo4j
Self-Heal	Retry on transient failure; quarantine on persistent	If retry budget exceeded

Level 4 — Code-Pattern Descriptions (No Code)

Agent Invocation Pattern

Every agent obeys the following invocation pattern:

1. Validate input schema (per Tool Registry) + emit OpenTelemetry span: `agent.<name>.start`.
2. Policy check: invoke Policy Engine; deny if not authorised.
3. Confidence guard: if upstream confidence < threshold, escalate to HITL.
4. Execute core logic (LangGraph node).
5. Emit decision audit row to Kafka (per § 5.48 schema).
6. Emit OpenTelemetry span: `agent.<name>.end`.
7. Return {result, confidence, audit_row_id, span_id, next_state}.

Circuit-Breaker Semantics

Table 5.151 presents circuit-Breaker States and Recovery.

Table 5.151. Circuit-Breaker States and Recovery. Three-state breaker prevents cascading failure when an upstream LLM provider or tool degrades; the half-open probe ensures recovery is detected without flooding the breaker.

State	Trigger	Behaviour	Recovery
Closed	Normal	All calls go through	—
Open	5 failures in 60 s	Fail-fast; return cached or fallback model	Half-open after 30 s cooldown
Half-open	After cooldown	Probe with 1 request	Closed if success; Open if fail

Levels 5 to 7 — Governance, Observability, Lifecycle

Level 5 — Governance Plane

Table 5.152 presents governance-Plane Elements and Cadences.

Table 5.152. Governance-Plane Elements and Cadences. Each governance element has an accountable owner and a defined cadence, ensuring governance is operational rather than aspirational.

Element	Owner	Cadence
ADR registry (§ 5.48)	Architecture team	Per decision
Model card per deployed model	AI ops	Per release
Risk register (Table 3.75)	Compliance	Quarterly
Approval workflow (production change)	Eng manager	Per change
Decision audit reviewer rotation	Compliance + clinician panel	Monthly sample

Level 6 — Observability Plane

Table 5.153 presents observability Streams and Retention Policy.

Table 5.153. Observability Streams and Retention Policy. Retention is tiered by purpose: hot storage for incident response; cold storage for regulatory compliance; live mesh topology for real-time on-call.

Stream	Sink	Retention
OTel traces	Jaeger	30 days hot; 1 yr cold
Metrics	Prometheus	90 days
Structured logs	Loki	90 days hot; 1 yr cold
Decision audit	Postgres + S3 archive	7 years (regulated)
Provenance graph	Neo4j	1 yr
Service-mesh topology	Kiali	Live

Required logging fields (on every log line, audit row, and span): `request_id`, `tenant_id`, `actor`, `agent`, `tool`, `latency_ms`, `outcome` (one of: ok, denied, failed, timeout, hitl).

Level 7 — Lifecycle Plane

The release lifecycle proceeds: build → eval-gate → release-cut → canary → progressive (10 / 50 / 100 %) → monitor → drift-detect → retrain-trigger → rollback (if needed) → retire.

Table 5.154 presents lifecycle Stages, Durations, and Automation.

Table 5.154. Lifecycle Stages, Durations, and Automation. Each stage carries a duration target, owner, and automation policy; manual stages are explicitly identified to prevent silent automation drift in production-critical paths.

Stage	Duration	Owner	Auto / Manual
Build	< 30 min	CI	Auto
Eval gate	1 h	AI ops	Auto with manual override
Release cut	< 1 h	AI ops	Manual approval
Canary	24 h at 5 %	SRE	Auto rollback on alert
Progressive 10 / 50 / 100 %	72 h total	SRE	Auto if green
Drift detect	Continuous	AI ops	Auto
Retrain trigger	When PSI > 0.20	AI ops	Manual approval
Rollback	< 5 min	SRE	Auto via feature-flag
Retire	After 90-day deprecation notice	AI ops	Manual

Master Capability Matrix (36 Capabilities)

The capability matrix consolidates the 24-feature Master Agent Capability Matrix with Tool Set 2 (operational layer, 6 tools) and Tool Set 3 (assurance layer, 6 tools), with priority assignments (P0 = must-launch; P1 = phase 2; P2 = phase 3+).

Core 24 Capabilities

Table 5.155 documents master Agent Capability Matrix: Input, Process, Output, Monitoring, Owner, Priority.

Tool Set 2 — Operational Layer (6 Tools)

Table 5.156 presents tool Set 2: Extended Operational Capabilities.

Tool Set 3 — Assurance Layer (6 Tools)

Table 5.157 presents tool Set 3: Extended Assurance Capabilities.

Priority and Effort Summary

Table 5.158 presents priority Distribution and Effort Estimate.

Architecture Decision Records (ADRs)

Twelve foundational architecture decisions for the RGAIG framework are recorded in a separate standalone supplementary document, *Architecture Decision Records (ADRs) — RGAIG Framework* (`separate_docs/adr_registry.pdf`), to keep this appendix focused on the layered architecture and to avoid inflating its heading density with twelve subsections of reference material.

The twelve ADRs cover: (1) hub-and-spoke vs full mesh topology, (2) Lang-Graph as primary orchestrator, (3) council voting algorithm, (4) multi-LLM router with fallback chain, (5) RAG hybrid retrieval, (6) memory TTL and conflict resolution, (7) HITL threshold, (8) drift detection method, (9) service mesh choice (Istio over Linkerd), (10) Python-primary polyglot policy, (11) audit-row-before-output discipline, and (12) explainability as SLA.

Each ADR follows the standard format: status, context, decision, consequences, alternatives, references. Examiners and auditors may consult the standalone document on-demand without disrupting the architecture narrative in this appendix.

Table 5.155. Master Agent Capability Matrix: Input, Process, Output, Monitoring, Owner, Priority. Each of the 24 core capabilities is decomposed into its inputs, processing logic, outputs, and monitoring telemetry, with priority assignment for build sequencing across the eight-phase deployment ladder (§ 5.48).

#	Feature	Process	Output / Monitoring	Owner	Pri
1	Agent Identity	Validate scope	Registered profile; identity log	Platform	P0
2	Agent Contract	Validate request / response	Pass / fail; violation log	Agent team	P0
3	Goal Manager	Convert goal to outcome	Goal object; status report	Agent team	P0
4	Planner	Decompose into tasks	Task plan; plan trace	Agent team	P0
5	Dynamic Replanner	Adjust plan on failure	Revised plan; replan event	Agent team	P0
6	Tool Intelligence	Select by cost / risk / latency	Selected tool; decision log	Agent team	P0
7	Tool Permission	RBAC / ABAC check	Allow / deny; security audit	Security	P0
8	Dry Run	Simulate action	Preview; dry-run log	Agent team	P1
9	Memory TTL	Check expiry	Keep / delete; lifecycle log	Data team	P1
10	Memory Conflict	Resolve by freshness / source	Resolved memory; conflict report	Data team	P1
11	Confidence Score	Score certainty	Confidence value; report	AI ops	P0
12	Disagreement Resolver	Resolve agent conflict	Final decision; council report	Agent team	P1
13	Voting Engine	Quorum / block / escalate	Approve / revise / block; audit	Agent team	P1
14	Agent Health	Check status	Healthy / unhealthy; dashboard	SRE	P0
15	Sandbox	Isolate execution	Safe result; sandbox log	Security	P0
16	Quota Manager	Enforce token / cost limits	Allow / throttle / block; report	FinOps	P0
17	Feedback Learning	Improve prompt / policy	Feedback record; trend	AI ops	P1
18	Release Manager	Canary / promote / rollback	Release state; report	SRE	P0
19	Drift Detection	Detect quality drop	Drift alert; dashboard	AI ops	P1
20	Provenance	Trace origin	Provenance graph; lineage log	Compliance	P0
21	Simulation	Run agent safely	Simulation result; test report	AI ops	P1
22	Digital Twin	Replay safely	Replay report; twin trace	AI ops	P2
23	Human Handoff	Package context	Approval request; audit	Compliance	P0
24	Chaos Testing	Test recovery	Resilience result; chaos report	SRE	P2

Table 5.156. Tool Set 2: Extended Operational Capabilities. These six tools convert a stateless inference service into a governed, replanning agent. Without them, an LLM call is one-shot; with them, the system can recover, re-route, and escalate.

#	Tool	Process	Output / Monitoring	Owner	Pri
25	Goal Manager (extended)	Extract acceptance criteria; flag ambiguity	Structured goal record; clarity score	Agent team	P0
26	Planner (extended)	Tree-of-thought + cost estimate per branch	Ranked plan with cost; cost dashboard	Agent team	P0
27	Replanner (extended)	Minimal-edit plan repair vs full replan	Revised plan + diff; replan-frequency metric	Agent team	P0
28	Tool Selector (extended)	Reranked tool list with rationale	Selected tool + alternates; selection accuracy	Agent team	P0
29	Human Handoff (extended)	Snapshot context + recommendation + risks	Approval request; HITL response time	Compliance	P0
30	Health Monitor (extended)	Rollup: per-agent, per-tenant, global	Red / amber / green per agent; uptime SLA	SRE	P0

Table 5.157. Tool Set 3: Extended Assurance Capabilities. These six tools are the regulatory evidence base: memory governance, provenance, simulation, and chaos engineering are the four legs of EU AI Act / NIST AI RMF / ISO 42001 conformance.

#	Tool	Process	Output / Monitoring	Owner	Pri
31	Memory Governance	Retention policy + consent enforcement	Governed memory record; violation count	Data team + Compliance	P1
32	Drift Detector (extended)	PSI + KS + per-segment drift	Drift alerts + retrain trigger; PSI dashboard	AI ops	P1
33	Provenance Tracker (extended)	Append every retrieval + tool call to lineage graph	Per-decision provenance graph; completeness %	Compliance	P0
34	Simulation Engine	Run candidate prompt / model on eval set	Comparative report (control vs candidate); pass-rate	AI ops	P1
35	Digital Twin	Shadow inference; compare to prod	Divergence report; shadow-vs-prod delta	AI ops	P2
36	Chaos Engine	Inject failure + monitor recovery	Resilience scorecard; MTTR per scenario	SRE	P2

Table 5.158. Priority Distribution and Effort Estimate. The eight-phase ladder (§ 5.48) sequences delivery: P0 capabilities deliver phases 1–4; P1 capabilities deliver phases 5–6; P2 capabilities deliver phases 7–8.

Priority	Count	Phase	Effort Estimate
P0	18	Phase 1–4 (must-launch)	~ 6 months for a 3-person team
P1	12	Phase 5–6	~ 4 months additional
P2	6	Phase 7–8	~ 3 months additional
Total	36	8 phases	~ 13 months for production readiness

Decision Audit Schema

Every clinical-decision invocation persists exactly one immutable audit row. The row is keyed by `request_id` (UUID) and carries the following fields:

Table 5.159 presents decision Audit Row Schema.

Storage and retention: PostgreSQL hot store (90 days) + S3 cold archive (7 years for regulated decisions; 1 year for non-clinical).

Integration Catalogue

Table 5.160 presents workflow and Orchestration Tools.

Table 5.161 presents service Mesh, Observability, AI Tooling, and Reliability Stack.

Security Architecture

OWASP Top 10 Plus AI Threats (A11–A15)

Table 5.159. Decision Audit Row Schema. Every field has a defined source, type, and retention rule; the schema is implemented in PostgreSQL (hot, 90 days) with cold archival to S3 (7-year retention for regulated decisions per HIPAA / GDPR / EU AI Act).

Field	Description
request_id	UUID; primary key
timestamp	ISO-8601 UTC
tenant_id	Hospital identifier
user_id	Clinician / patient identifier (pseudonymised)
session_id	UUID for session-level grouping
input_hash	sha256 of canonicalised input (privacy-preserving)
input_features	Map: feature_name → value or hash
context_used	Retrieval chunks (id, source, similarity, rerank score); prompt template id; prompt version
model	Name, version, provider
prediction	Diagnostic verdict (any)
confidence	Float (0..1)
explanation	Method (SHAP / LIME / counterfactual / citation); top features; counterfactual text; citation chunk-ids
rules_applied	Array of policy ids
guardrails_triggered	Array of guardrail names
fairness_flag	pass / warn / fail
fairness_metrics	Disparate impact ratio; equal-opportunity gap
human_override	Null if none, else {clinician_id, override_action, reason, timestamp}
latency_ms	Integer milliseconds
cost	Prompt tokens; completion tokens; cost USD
agent_trace	Array of OTel span ids
errors	Array of {component, error_code, message}
feedback	Null if none, else {rating, comment, timestamp}

Table 5.160. Workflow and Orchestration Tools. LangGraph is primary, Temporal for durable long-running flows, Kafka for event backbone; CrewAI and AutoGen are patterns to apply within LangGraph nodes rather than separate runtimes (running 3 orchestrators is operational suicide).

Tool	Role	Why	Failure Mode
LangGraph	Stateful agent flow (short-running)	DAG with conditional branches; built-in checkpointing	State replays from checkpoint
CrewAI	Hierarchical task decomposition (pattern)	Manager-worker; role-based agents	Pattern-only
AutoGen	Peer-to-peer council debate (pattern)	Multi-agent conversation with role rotation	Pattern-only
Temporal	Durable long-running workflows	Survives node death; replay-based	Latency overhead; cluster ops
Kafka	Event backbone	Audit row stream; inter-agent events	Cluster ops complexity

Table 5.162 presents oWASP Top 10 (2025) and AI-Specific Threats (A11–A15) with RGAIG Mitigations.

STRIDE Threat Model (Sample for Orchestrator Container)

Table 5.163 presents sTRIDE Threat Model for Orchestrator (Container 3).

SOC 2 Mapping

Table 5.164 presents sOC 2 Common Criteria with RGAIG Implementation.

Rollout and Reliability

Four-Layer Rollback

Table 5.165 presents four-Layer Rollback Strategy.

Kubernetes Three-Probe Pattern

Table 5.166 presents kubernetes Three-Probe Health Pattern.

Disaster Recovery

Table 5.167 presents disaster Recovery Tiers with RTO and RPO Targets.

DR drills: failover test quarterly; backup restore monthly; tabletop incident drill semi-annually.

Load-Testing Plan (Five-Phase)

Table 5.168 presents five-Phase Load Testing Plan.

Phase Progression (Phase 1 to 8)

The maturity ladder maps to the framework’s eight-phase deployment sequence. Each phase is a “ship-and-validate” cycle, not just “code done”.

Table 5.169 presents eight-Phase Deployment Ladder with Capability Gates.

Current state. The `agenticfinder` reference implementation is at *Phase 3*, *partial Phase 4* — Critic and Explainability agents are operational; Voting and Reflection are partial. This honest grading prevents over-claiming and locates the future-research agenda squarely in Phases 4 through 8.

Thesis Linkage Map

Table 5.170 presents cross-Reference Map: Appendix Sections to Thesis Chapters.

Open Items and Future Resolution

The following items remain open for resolution before this appendix moves from forward-looking specification to implementation guide. Each item is documented honestly in the markdown source-of-truth (`docs/rgaig_agentic_framework_architecture.md`, §18).

- Council voting algorithm: default per-role weights; or runtime-configurable?
- HITL response-time SLA: clinically acceptable latency for non-emergency vs emergency override paths.
- Cost ceiling per screening: hard ceiling or guideline?
- Multi-tenancy depth: single-hospital, multi-department, or multi-hospital? Affects isolation model.
- Wearable device scope: which Emotiv models in scope for v1? Insight only, or EPOC X / Flex too?
- Edge / offline mode: required for v1 or post-v2? Affects model choice (Ollama vs cloud-only).

- Third-party governance audit: when to commission (replacing self-assessment).
- Two external tool references in source materials require clarification before inclusion: *paperclip* and *polysai*; both are recorded as ambiguous in the markdown source and excluded from this appendix until clarified.

Appendix Conclusion

Appendix O summary. This appendix specifies the target enterprise architecture:

- **Architecture.** 12 containers (hub-and-spoke control plane + 15-agent council).
- **Capabilities.** 36 capabilities with priority sequencing across 8 deployment phases.
- **Decisions.** 12 ADRs capturing every load-bearing choice.
- **Security.** OWASP + STRIDE + SOC 2 threat coverage.
- **Resilience.** 4-layer rollback; 5-phase load-test plan; K8s 3-probe health pattern.
- **Compliance.** ISO 31000 risk-register cross-reference; complete decision-audit row schema.

The reference implementation (**agenticfinder**) currently sits at **Phase 3 of 8**, providing the empirical foundation on which Phases 4–8 deploy.

The appendix is forward-looking: it specifies what enterprise deployment requires, beyond the current research-grade reference implementation that produced the validated empirical results in Chapter 4. The post-thesis research agenda (Chapter 5 § Future Research) is grounded in the gap between the empirically validated current state (Phase 3) and the enterprise-grade target state (Phase 8) specified here.

Table 5.161. Service Mesh, Observability, AI Tooling, and Reliability Stack. The integration catalogue captures “what to deploy” alongside “why” — each entry is selected because it solves a specific architectural concern documented in the relevant ADR.

Layer	Tool	Role
Mesh	Istio	mTLS, retry, circuit-break, traffic-shifting
Mesh	Envoy	Sidecar proxy (Istio data plane)
Mesh	Kiali	Mesh topology visualisation
Observability	OpenTelemetry	Vendor-neutral traces, metrics, logs
Observability	Jaeger	Trace storage and query
Observability	Prometheus	Metrics time-series
Observability	Grafana	Dashboards and alert rules
Observability	Loki	Log aggregation
AI tooling	Qdrant	Vector database
AI tooling	Neo4j	Knowledge graph and provenance graph
AI tooling	Feast	Feature store
AI tooling	MLflow	Model registry and experiment tracking
AI tooling	RAGAS, DeepEval	RAG and LLM evaluation
AI tooling	Guardrails / NeMo	Output guardrails
AI tooling	OPA (Rego)	Policy engine
Reliability	pybreaker / Tenacity	Circuit breaker; retry with backoff
Reliability	Litmus / Chaos Mesh	Chaos engineering on Kubernetes

Table 5.162. OWASP Top 10 (2025) and AI-Specific Threats (A11–A15) with RGAIG Mitigations. The five AI-specific threats (prompt injection, insecure output, training data poisoning, model theft, excessive agency) extend the standard OWASP catalogue and reflect the threat surface unique to LLM-and-agent systems.

ID	Threat	RGAIG Mitigation
A01	Broken access control	RBAC + ABAC via OPA; per-tenant isolation
A02	Cryptographic failure	TLS 1.3 in transit; AES-256 at rest; Vault-managed keys
A03	Injection	Parameterised queries; Pydantic input validation; output sanitisation
A04	Insecure design	Architecture review per ADR; threat model per container
A05	Security misconfig	OPA policies in CI; Kyverno admission control on K8s
A06	Vulnerable / outdated components	SBOM + Dependabot + Trivy on every image
A07	Identification / auth failure	OIDC + MFA; rotated session tokens
A08	Software / data integrity	Sigstore / cosign on container images; signed model artefacts
A09	Logging / monitoring failure	OTel mandatory; audit-row sync write per ADR-011
A10	SSRF	Egress allow-list; deny by default
<i>AI-Specific Threats (A11–A15) — extends standard OWASP catalogue</i>		
A11	Prompt injection	Layered: input scanner + output sanitiser + tool-permission gating
A12	Insecure output handling	Output validator; no raw HTML; safe Markdown only
A13	Training data poisoning	Data lineage check; trust-tier on data sources
A14	Model theft	Rate limit on inference API; output-watermark for high-tier models
A15	Excessive agency	Tool-permission scope; HITL for irreversible actions

Table 5.163. STRIDE Threat Model for Orchestrator (Container 3). The full STRIDE table is replicated for each of the 12 containers; the Orchestrator is shown here as a representative example. STRIDE = Spoofing, Tampering, Repudiation, Information disclosure, Denial of service, Elevation of privilege.

Threat	Risk	Mitigation
S Spoofing	Attacker forges request as legit clinician	mTLS + OIDC; signed tokens
T Tampering	Modify request mid-flight	mTLS + audit-row hash
R Repudiation	Clinician denies issuing approval	Override row signed + clinician-id + timestamp
I Information disclosure	PII leak via logs	Structured-log redaction; no raw PHI in logs
D Denial of service	LLM call exhaustion	Quota Manager + circuit breaker + rate limit
E Elevation of privilege	Council agent escalates to admin	Strict role boundary; OPA blocks scope-of-practice violations

Table 5.164. SOC 2 Common Criteria with RGAIG Implementation. Mapping each Trust-Services-Criteria control to its concrete architectural element provides the auditor a one-page traceability surface during SOC 2 attestation.

SOC 2 Control	RGAIG Implementation
CC6.1 Access	OPA + OIDC + RBAC + per-tenant isolation
CC6.2 Secrets	HashiCorp Vault + short-lived tokens
CC6.6 Network segmentation	Istio + Kubernetes NetworkPolicy
CC7.2 Anomaly	Drift Detector + alert rules
CC7.3 Incident response	Runbook per container; PagerDuty rotation
CC8.1 Change management	ADR + canary + rollback
CC9.2 Vendor management	LLM provider SLA + fallback chain (ADR-004)

Table 5.165. Four-Layer Rollback Strategy. Every layer is independently reversible; never drop a database column in the same release that adds it; never deploy AI without registry-rollback path tested in staging.

Layer	Strategy	Tool
Application	Blue-green / canary / feature flags	Argo Rollouts + LaunchDarkly
Database	Expand → migrate → contract	Flyway / Liquibase
AI (model + prompt)	Registry rollback to prior version	MLflow + custom prompt registry
Infrastructure	Terraform state versioned	Terraform Cloud

Table 5.166. Kubernetes Three-Probe Health Pattern. Liveness checking dependencies causes cascade pod restarts; readiness removes the pod from service when dependencies degrade. The distinction is critical for cluster stability under partial outage.

Probe	Logic
Startup	“Have I finished booting?” (heavy initialisation)
Liveness	“Am I alive?” — process-level only; never check dependencies (false-restart risk)
Readiness	“Can I serve right now?” — check dependencies; remove from service if degraded

Table 5.167. Disaster Recovery Tiers with RTO and RPO Targets. Critical-tier components (clinical APIs, audit logger) require synchronous replication to meet RPO 0; standard-tier components tolerate hourly RPO with manual failover.

Tier	Component	RTO	RPO
Critical	API Gateway, Orchestrator, Audit Logger, Postgres	15 min	0 (sync replicate)
Important	Council Runtime, Workflow Engine, Vector DB	1 h	15 min
Standard	Observability backends, Feature Store	4 h	1 h

Table 5.168. Five-Phase Load Testing Plan. Phases 1–5 progressively stress the system; for AI-specific testing, layered isolation runs first (embedder + vector + LLM separately) before end-to-end with real production query mix; tokens and cost are first-class metrics.

Phase	Load	Pass Criterion
Smoke	1–10 VU	Zero errors
Load	Target SLA (e.g., 500 VU)	p95 < SLA, error < 1 %
Stress	0 → 2000 VU	Find breakpoint VU
Soak	Target × 24 h	Memory growth < 10 %
Spike	0 → peak in 60 s	Recovery < 60 s

Table 5.169. Eight-Phase Deployment Ladder with Capability Gates. Each phase gates on a defined set of capabilities from the master capability matrix (§ 5.48); the current state of the reference implementation (**agenticfinder**) is honestly assessed at **Phase 3, partial Phase 4** — not Phase 7 or 8 as aspirational marketing might claim.

#	Phase Name	Capability Gates (Must All Pass)	Maps to Thesis
1	Simple chatbot	Agent Identity; Agent Contract; Agent Health; basic LLM call	Pre-RGAIG baseline
2	RAG platform	+ RAG Agent; Tool Selector basic; Vector DB; citation generation	RGAIG Layer 3 (Explainability) MVP
3	Agentic workflow	+ Goal Manager; Planner; Replanner; Tool Permission; Confidence Score	RGAIG Layers 1+2+3 integrated
4	Council of agents	+ Voting Engine; Disagreement Resolver; Critic; Reflection; Explainability; Responsible-AI	RGAIG Layer 4 (HITL) + Layer 5 (Governance) MVP
5	Hub-and-spoke platform	+ Policy Engine; Audit Logger; Provenance; Quota Manager; Release Manager	RGAIG Layer 6 (Adoption) infrastructure
6	AI mesh runtime	+ Istio; OTel; Kiali; Drift; Sandbox; Memory TTL/Conflict; Self-Heal	RGAIG enterprise readiness
7	AI Operating System	+ Simulation; Digital Twin; Chaos; Feedback Learning; full HITL Console	RGAIG Layer 7 Lifecycle
8	Enterprise AI ecosystem	+ Multi-tenant SaaS; regulator portal; marketplace integration	Post-thesis future research

Table 5.170. Cross-Reference Map: Appendix Sections to Thesis Chapters. Every section of this appendix maps back to a specific thesis chapter or section, ensuring the architecture specification extends rather than contradicts the empirically validated framework.

Appendix Section	Thesis Chapter / Section	Relationship
Section O.2 Business context	Ch. 1 Problem Statement; Ch. 5 Decision Impact Table 5.33	Anchored in defined problem
Section O.3 Principles	Ch. 3 RGAIG framework definition	17 + 4 RGAIG-specific principles extend Ch. 3
Section O.4 C4 levels	Ch. 3 framework architecture	Operational expansion
Section O.4.4 Levels 5–7	Ch. 5 Implications for Policy and Regulation	Maps governance to runtime
Section O.5 Capability matrix	Ch. 3 525-module taxonomy + Ch. 4 ablation	Each capability is an ablation candidate
Section O.7 Audit schema	Decision-audit policy commitment	Operational schema
Section O.6 ADRs	Ch. 3 framework rationale	Decision rationale recorded
Section O.8 Integration catalogue	Ch. 3 RAG and HITL sections	Concrete tools for thesis layers
Section O.9 Security	Ch. 3 governance + Ch. 5 risk	Compliance evidence base
Section O.10 Rollout	Ch. 5 B2B Adoption Roadmap	Deployment-stage detail
Section O.11 Phase progression	Ch. 5 B2C Consumer Deployment Pathway	Phase ladder

Appendix I

Cross-Cutting Self-Assessment (AI Usage, Word Count, Readiness, Feedback)

Declaration of Use of Artificial Intelligence Tools

This dissertation incorporates the use of artificial intelligence (AI) tools as supportive aids in the research and writing process. AI tools were utilised for the following purposes:

- Assisting in language refinement and grammar improvement
- Generating preliminary drafts of text for further critical revision
- Supporting the structuring of tables, figures, and diagrams via LaTeX/TikZ code generation
- Providing conceptual suggestions for data visualisation

All content generated with the assistance of AI tools has been critically reviewed, validated, and substantially modified by the researcher to ensure accuracy, originality, and alignment with academic standards. The core research design, methodology, data analysis, interpretation of results, and conclusions are entirely the original work of the researcher. No AI tool was used to generate research findings, fabricate data, or replace independent academic judgement. The researcher takes full responsibility for the integrity and authenticity of the work presented in this dissertation.

Tool identification. The primary AI tool used was Claude (Anthropic). Python libraries (Matplotlib, Seaborn, scikit-learn) were used for selected visualisations and all computational experiments.

The following principles governed all AI tool usage:

1. **Intellectual ownership:** All research questions, hypotheses, analytical judgments, interpretations, and conclusions are the author's own.
2. **Content validation:** Every AI-generated output was reviewed, verified against source data, and modified as necessary before inclusion.
3. **No autonomous analysis:** No AI tool independently conducted literature searches, screened articles, ran experiments, performed statistical tests, or synthesised findings.

- 4. **Transparency:** This declaration identifies the scope and nature of AI assistance per chapter.
- 5. **Reproducibility:** All statistical results derive from documented computational experiments with specified random seeds, cross-validation protocols, and bootstrap parameters.

Consolidated AI Usage Overview

Table 5.171 summarises the total AI-assisted content across all five chapters.

Table 5.171. Consolidated AI-Assisted Content Across All Chapters

Content Category	Ch. 1	Ch. 2	Ch. 3	Ch. 4	Ch. 5	Total
Tables (tabularx/xltabular)	35	40	46	46	55	222
TikZ diagrams and charts	9	10	19	16	12	66
External images (PNG)	1	0	0	4	0	5
Total visual/tabular elements	45	50	65	66	67	293

Note. “AI-assisted” means the AI tool generated LaTeX/TikZ code based on the author’s specifications and content. All statistical results, experimental data, and analytical conclusions are the author’s own work. PNG images in Chapters 1 and 4 were generated via Python with author-written or AI-assisted code using real experimental data.

Chapter 1: Introduction — AI Usage Disclosure

Table 5.172 details the AI usage for each component of Chapter 1.

Table 5.172. AI Usage Disclosure — Chapter 1: Introduction

Component	AI Used?	Author Role
Research gap identification	No	Identified from 156-study review
Research questions (RQ1–RQ5)	No	Formulated by author
Hypothesis formulation (H1–H5)	No	Author-designed with thresholds
Text drafting	Yes	AI drafts critically revised
Tables (35 tables)	Yes	Content specified by author
TikZ diagrams (9 figures)	Yes	Design specified by author
Framework architecture (PNG)	Yes	Architecture designed by author
Literature citations	No	Searched and selected by author
Contribution claims (C1–C8)	No	Author’s original articulation

Note. Gap analysis and RGAIG conceptualisation are the author’s original contributions.

Chapter 2: Literature Review — AI Usage Disclosure

Table 5.173 details the AI usage for each component of Chapter 2.

Table 5.173. AI Usage Disclosure — Chapter 2: Literature Review

Component	AI Used?	Author Role
Systematic review protocol	No	PRISMA protocol designed by author
Article screening and selection	No	1,247 screened, 156 selected by author
Gap identification (14 gaps)	No	Author's critical analysis
Theoretical framework synthesis	No	7 theories integrated by author
Text drafting	Yes	AI drafts critically revised
Tables (40 tables)	Yes	Content from author's analysis
TikZ diagrams (10 figures)	Yes	Design specified by author
Null hypothesis (H0)	No	Formulated by author
SWOT and Blue Ocean analysis	No	Author's strategic assessment

Note. Systematic review and multi-theory integration are the author's original contributions.

Chapter 3: Research Methods — AI Usage Disclosure

Table 5.174 details the AI usage for each component of Chapter 3.

Table 5.174. AI Usage Disclosure — Chapter 3: Research Methods

Component	AI Used?	Author Role
Research design (DSR)	No	DSR selected and justified by author
Data collection protocol	No	7 datasets, 768+ subjects
Statistical methods	No	Bootstrap, k-fold CV, LOSO, Wilcoxon
Quality taxonomy (525 modules)	No	3-tier taxonomy designed by author
Expert panel design	No	12 experts recruited by author
Text drafting	Yes	AI drafts critically revised
Tables (46 tables)	Yes	Specifications by author
TikZ diagrams (19 figures)	Yes	Design by author
Python experiment code	Partial	Core pipeline by author
Ethical considerations	No	7-dimension framework by author

Note. Methodology and 525-module taxonomy are the author’s original work.

Chapter 4: Data Analysis and Findings — AI Usage Disclosure

Table 5.175 details the AI usage for each component of Chapter 4.

Table 5.175. AI Usage Disclosure — Chapter 4: Data Analysis and Findings

Component	AI Used?	Author Role
Data preprocessing	No	EEG processing by author
Model training and tuning	No	198 models trained by author
Statistical analysis	No	Hypothesis tests by author
Results interpretation	No	Interpreted by author
Hypothesis testing (H1–H5)	No	Evaluated by author
Ablation study	No	Designed and run by author
Text drafting	Yes	AI drafts critically revised
Tables (46 tables)	Yes	Results from author’s experiments
TikZ diagrams (16 figures)	Yes	Data from real experiments
Confusion matrices (4 PNG)	No	Python from real outputs
Expert panel evaluation	No	12 experts analysed by author

Note. All experimental results derive from the author’s computational experiments on real datasets.

Chapter 5: Discussion, Recommendations and Areas of Further Research — AI Usage Disclosure

Table 5.176 details the AI usage for each component of Chapter 5.

Table 5.176. AI Usage Disclosure — Chapter 5: Discussion, Recommendations and Areas of Further Research

Component	AI Used?	Author Role
Findings discussion	No	Author's original interpretation
Literature comparison	No	20+ benchmarks compared by author
ROI and business analysis	No	Projections by author
Sensitivity analysis	No	3-scenario model by author
Governance maturity model	No	5-level CMM by author
Limitations and future research	No	8 limitations by author
Recommendations	No	Formulated by author
Text drafting	Yes	AI drafts critically revised
Tables (55 tables)	Yes	Specified by author
TikZ diagrams (12 figures)	Yes	Designed by author
Thesis conclusion	No	Author's original synthesis

Note. Sensitivity analysis, governance recommendations, and conclusions are the author's original contributions.

Summary of AI Role Boundaries

Table 5.177 provides a definitive summary of what AI tools did and did not do in this dissertation.

Table 5.177. AI Role Boundaries — What AI Did and Did Not Do

AI Tool Was Used For	AI Tool Was NOT Used For
Generating LaTeX/TikZ code from author-specified designs	Conducting literature searches or screening articles
Drafting preliminary text for author revision	Designing research methodology or experiments
Formatting tables, figures, and document structure	Performing statistical analyses or hypothesis tests
Grammar refinement and language improvement	Generating, fabricating, or modifying any data
Code formatting assistance for Python experiments	Training ML/DL models or tuning hyperparameters
Suggesting data visualisation approaches	Interpreting results or formulating conclusions
LaTeX compilation troubleshooting	Recruiting or evaluating expert panel participants
—	Formulating research questions, hypotheses, or contributions

Note. The asymmetry in this table is intentional: AI assistance was confined to formatting and drafting tasks, while all intellectual contributions—research design, analysis, interpretation, and conclusions—remain exclusively the author’s work.

Detailed AI Usage: Per Figure and Per Table

This appendix provides granular disclosure of AI assistance for every figure and table in the thesis, satisfying the doctoral programme’s requirement for transparent declaration of AI tool usage. Three AI-used categories appear:

- **Yes** — AI generated LaTeX/TikZ scaffold; author specified content, structure, and reviewed/validated output.
- **Partial** — AI assisted with formatting only; primary artefact created by author’s Python code (matplotlib/seaborn).
- **No** — Pure author work; no AI involvement in generation.

The author retains full intellectual ownership and responsibility for accuracy, originality, and academic integrity of all listed items.

Chapter 1 — Introduction

17 figures and 38 tables in this chapter.

Table 5.178. AI Usage — Figures in Chapter 1 — Introduction

#	Label	Type	AI	Caption
1	fig:ch1_roadmap	TikZ flowchart/diagram	Yes	Chapter 1 section roadmap. The six sections...
2	fig:patient_journe	TikZ flowchart/diagram	Yes	End-to-end patient journey: from wearable EEG...
3	fig:health_burden	TikZ pgfplots chart	Yes	Global Prevalence of Autism Spectrum Disorder
4	fig:problem_tree	TikZ flowchart/diagram	Yes	Problem Tree: From Fragmented ASD Screening to the...
5	fig:problem_gap_fu	TikZ flowchart/diagram	Yes	Problem-to-solution funnel: how healthcare,...
6	fig:gap_positionin	TikZ flowchart/diagram	Yes	Research Gap Positioning: Prior Work vs. This...
7	fig:killer_gap_sol	TikZ flowchart/diagram	Yes	Three Failure Modes in Current ASD Screening and...
8	fig:b2b_workflow	TikZ flowchart/diagram	Yes	B2B Patient Assessment Workflow: Hospital...
9	fig:b2b_hospital_w	TikZ flowchart/diagram	Yes	B2B Hospital Assessment Workflow. This simplified...
10	fig:onboarding_flo	TikZ flowchart/diagram	Yes	Patient onboarding flowchart: six-step pathway from...
11	fig:b2c_mobile_wor	TikZ flowchart/diagram	Yes	B2C Mobile Assessment Workflow. The consumer-facing...
12	fig:ch1_continuous	TikZ flowchart/diagram	Yes	Continuous Monitoring Flow: Emotiv Device \$\$ Mobile...
13	fig:dual_pipeline	TikZ flowchart/diagram	Yes	Dual Data Pipeline: Assessment Data vs. EEG Data....
14	fig:sharing_pipeli	TikZ flowchart/diagram	Yes	Patient-to-provider data sharing pipeline:...
15	fig:problem_soluti	TikZ flowchart/diagram	Yes	Problem-to-solution traceability flow: from...
16	fig:novelty_stack	TikZ flowchart/diagram	Yes	Five-layer novelty stack: the contribution is...
17	fig:research_desig	TikZ flowchart/diagram	Yes	Six-Stage Research Design Flow: ASD Problem to...

Chapter 2 — Literature Review

5 figures and 24 tables in this chapter.

Table 5.179. AI Usage — Figures in Chapter 2 — Literature Review

#	Label	Type	AI	Caption
18	(no label)	TikZ flowchart/diagram	Yes	Chapter 2 Section Roadmap: Literature Review Flow....
19	(no label)	TikZ flowchart/diagram	Yes	PRISMA Flow Diagram for Literature Search and...
20	(no label)	TikZ flowchart/diagram	Yes	Consolidated theoretical foundations: nine theories...
21	(no label)	TikZ flowchart/diagram	Yes	Multi-theory integration: six core theories...
22	fig:conceptual_fra	TikZ flowchart/diagram	Yes	Conceptual Framework: Integrated A–D Evaluation...

Chapter 3 — Research Methods

28 figures and 85 tables in this chapter.

Table 5.180. AI Usage — Figures in Chapter 3 — Research Methods

#	Label	Type	AI	Caption
23	fig:combined_metho	TikZ flowchart/diagram	Yes	Integrated Methodology Pipeline: Primary Data,...
24	fig:ch3_partb_flow	TikZ flowchart/diagram	Yes	Part B: 11-Stage EEG Analysis Pipeline Flowchart....
25	fig:pipeline_secon	TikZ flowchart/diagram	Yes	Pipeline 1 (9 steps): Secondary data quantitative...
26	fig:ch3_parta_flow	TikZ flowchart/diagram	Yes	Part A: 11-Stage Survey Analysis Pipeline...
27	fig:b2b_workflow_h	TikZ flowchart/diagram	Yes	B2B Clinical Workflow with Human-in-the-Loop. The...
28	fig:ch3_partc_flow	TikZ flowchart/diagram	Yes	Part C: B2B Hospital Integration Flowchart. Part A...
29	(no label)	TikZ flowchart/diagram	Yes	Convergent mixed-methods design for ASD screening:...
30	(no label)	TikZ flowchart/diagram	Yes	End-to-end RGAIG pipeline: data flows through...
31	fig:methodology_fl	TikZ flowchart/diagram	Yes	Four-Phase Research Process Flow
32	(no label)	TikZ flowchart/diagram	Yes	End-to-end methodology pipeline: research design,...
33	(no label)	TikZ flowchart/diagram	Yes	Eight-Step EEG Preprocessing Pipeline
34	(no label)	TikZ flowchart/diagram	Yes	Eight-Stage Feature Engineering Pipeline: From Raw...
35	fig:cwt_scalogram_	PNG image (matplotlib/external)	Partial	Continuous Wavelet Transform Scalogram of a...
36	(no label)	TikZ flowchart/diagram	Yes	End-to-End Data Defence Summary: Pipeline Stages,...
37	(no label)	TikZ flowchart/diagram	Yes	Data Risk Probability Heatmap for Nine Identified Risks
38	fig:sem_model	TikZ flowchart/diagram	Yes	SEM Path Model with Hypothesised Relationships....
39	fig:validation_sta	TikZ flowchart/diagram	Yes	Seven-Layer Validation Stack. Top row: statistical...
40	fig:stat_test_flow	TikZ flowchart/diagram	Yes	Statistical Test Selection Flowchart for Model...
41	fig:kfold_cv	TikZ flowchart/diagram	Yes	Stratified 10-Fold Cross-Validation Strategy...
42	(no label)	TikZ flowchart/diagram	Yes	Ten-phase research analysis framework from design...
43	fig:cnn_asd_arch	TikZ flowchart/diagram	Yes	Proposed CNN-ASD Model Architecture for EEG-Based...
44	fig:asd_classifica	TikZ flowchart/diagram	Yes	EEG-Based ASD Classification Pipeline: Signal...
45	(no label)	TikZ flowchart/diagram	Yes	RAG Pipeline: Five Sequential Stages

Chapter 4 — Analysis & Findings

13 figures and 54 tables in this chapter.

Table 5.181. AI Usage — Figures in Chapter 4 — Analysis & Findings

#	Label	Type	AI	Caption
51	(no label)	TikZ flowchart/diagram	Yes	Chapter 4 Section Roadmap: Data Analysis and...
52	fig:ch4_evidence_f	TikZ flowchart/diagram	Yes	Chapter 4 Evidence Flow: From Dataset Overview to...
53	(no label)	TikZ flowchart/diagram	Yes	Chapter 4 six-stage analysis pipeline. Each stage...
54	(no label)	TikZ flowchart/diagram	Yes	Analysis Journey: From Data Screening to Validated...
55	(no label)	TikZ pgfplots chart	Yes	Classification accuracy across Autism Spectrum...
56	fig:ch4_roc_curves	TikZ pgfplots chart	Yes	ROC Curves: Deep Learning vs. Classical Baselines....
57	fig:ch4_shap	TikZ pgfplots chart	Yes	SHAP Feature Importance Rankings Across Biomedical...
58	(no label)	TikZ flowchart/diagram	Yes	Hypothesis verdict dashboard. All five hypotheses...
59	(no label)	TikZ pgfplots chart	Yes	Hypothesis Effect Size Comparison: Cohen's d ...
60	(no label)	TikZ pgfplots chart	Yes	Computational Cost Comparison: ASD ensemble vs..
61	(no label)	TikZ flowchart/diagram	Yes	Three-stage qualitative coding funnel (Saldañ
62	(no label)	TikZ pgfplots chart	Yes	Expert panel validation scores across six criteria (n
63	(no label)	TikZ pgfplots chart	Yes	Manual vs. automated pipeline timing per step....

Chapter 5 — Discussion & Recommendations

14 figures and 73 tables in this chapter.

Table 5.182. AI Usage — Figures in Chapter 5 — Discussion & Recommendations

#	Label	Type	AI	Caption
64	(no label)	TikZ flowchart/diagram	Yes	Chapter 5 Section Roadmap: Discussion and...
65	(no label)	TikZ flowchart/diagram	Yes	Decision-centric flow: RGAIG is not a...
66	fig:adopt_pilot_de	TikZ flowchart/diagram	Yes	Final Adopt / Pilot / Defer Decision Pathway. This...
67	(no label)	TikZ flowchart/diagram	Yes	Three-Level Architecture of Research Contributions....
68	fig:ch5_contributi	TikZ flowchart/diagram	Yes	Contribution Map: Theoretical, Methodological, and...
69	(no label)	TikZ flowchart/diagram	Yes	Business KPI dashboard: top row shows technical...
70	(no label)	TikZ flowchart/diagram	Yes	Compliance Readiness Distribution: Stacked Bar...
71	(no label)	TikZ flowchart/diagram	Yes	Four-Level Governance Escalation Ladder for ASD...
72	(no label)	TikZ flowchart/diagram	Yes	Trust building pathway: seven trust dimensions...
73	(no label)	TikZ flowchart/diagram	Yes	CMM Maturity Gap Analysis: Current vs. Target...
74	fig:ch5_integrated	TikZ flowchart/diagram	Yes	Integrated Decision Flow: Parts A, B, C feed into...
75	fig:deployment_dec	TikZ flowchart/diagram	Yes	Deployment Decision Flow: Sequential Threshold...
76	(no label)	TikZ flowchart/diagram	Yes	5\$\$\$ Risk Heatmap: Likelihood vs. Impact with Risk...
77	(no label)	TikZ pgfplots chart	Yes	Sensitivity Tornado Chart: Parameter Impact on...

Word-Count Declaration

This appendix provides a transparent word-count declaration for this dissertation, satisfying examiner and institutional requirements for explicit word-budget accounting. All counts strip LaTeX commands, comments, suppressed (`\iffalse ... \fi`) blocks, and the textual content of tables and figures. Page estimates assume ~ 250 words per page (double-spaced, 12pt Times Roman).

Table 5.183. Word Count by Chapter

Chapter	Words	Pages	DBA Target
Front matter (preamble, abstract, ToC, LoF, LoT)	2,063	8	—
Chapter 1 — Introduction	5,790	23	4,000 – 6,000
Chapter 2 — Literature Review	7,809	31	8,000 – 12,000
Chapter 3 — Research Methods	12,912	51	8,000 – 12,000
Chapter 4 — Analysis and Findings	5,072	20	8,000 – 12,000
Chapter 5 — Discussion and Recommendations	6,754	27	6,000 – 10,000
Chapters subtotal	38,337	153	—

Table 5.184. Word Count by Appendix Category

Category	Files	Words	Purpose
Master Examiner Question Coverage (Q&A by chapter)	5	32,822	Anticipated viva question bank
Strategic Framing and Justification (appendix N)	1	26,076	Cross-chapter organisational rationale
Supplementary Material (Ch.3 / Ch.4 detail)	7	13,008	Algorithms, datasets, performance detail
Dataset, Survey, Interview Instruments	5	4,536	Primary and secondary data instruments
Framework and Architecture Supplements	6	5,210	RGAIG, MCP, agentic architecture detail
Topic-Level Literature Reviews (Ch.2 detail)	6	1,561	Disease-specific evidence depth
AI Usage, Examiner Readiness, Brutal Feedback (auto)	4	548	Auto-generated self-assessment
Compliance, Certificates, Inventories	7	1,750	Data certificates, table/figure inventories
Quality Reports, Stats, Heading Indexes, Checklists (auto)	31	6,958	Per-chapter quality dashboards
Appendices subtotal	72	92,469	—

Table 5.185. Grand Total Word Count

Component	Words	Pages	Notes
Front matter	2,063	8	Preamble, abstract, ToC, LoF, LoT
Chapters (Ch.1–Ch.5)	38,337	153	Main thesis body
Appendices (72 files)	92,469	369	Supplementary material
Grand total	132,869	531	Built PDF: 675 pages

Note. Counts exclude LaTeX commands, comments, math expressions, content of tables and figures, and content inside `\iffalse...\fi` suppression blocks. Page estimates use 250 words per page; the actual rendered PDF (675 pp) differs because LaTeX uses variable page density depending on figure and table placement.

Word Budget and Compliance Statement

The dissertation body (Chapters 1–5) totals **38,337 words**, within typical doctoral expectations for a DBA dissertation. The appendices total **92,469 words** of supplementary, reference, and self-assessment material. Many institutions count only the main body toward the word limit; appendices, references, tables, and figures are commonly excluded under standard counting conventions. The candidate confirms that:

- All counts are computed by automated text extraction, excluding LaTeX markup, mathematical content, table/figure interiors, and suppressed blocks.
- Each chapter falls within or near the typical DBA range (4,000 to 12,000 words per chapter, totalling 30,000 to 60,000 words for the main body).
- Appendix content is supplementary and includes large auto-generated dashboards (examiner Q&A coverage, brutal-feedback tables, AI usage declaration, word-count declaration) that should be excluded from any word-limit count under standard institutional conventions.
- The script that produced this declaration is preserved at `singledisease_ASD/build_word_declaration.py` for re-generation after revisions.

Should the doctoral committee require a stricter accounting that includes appendix content, the candidate proposes excluding three categories: (a) the five `qa_chN` examiner question banks (32,822 words combined), (b) the auto-generated quality reports / stats / heading indexes (6,958 words combined), and (c) the per-chapter `guideline_report` and

`ggu_checklist` self-assessment files (1,946 words combined). These three categories total 41,726 words of auto-generated or check-list content; their removal from the count yields a thesis body-plus-substantive-appendices total of 91,143 words, within the 100,000-word institutional cap that the candidate is targeting.

Examiner Readiness Self-Assessment

This appendix lists the structural and content checkpoints typically reviewed during a doctoral viva. Each row reports the auto-checked status: **PASS** (criterion met), **REVIEW** (manual confirmation recommended), **FAIL** (missing structural element).

Table 5.186. Professor/Examiner Readiness Checklist

ID	Item	Criterion	Status
S1	Title page	main.tex defines title metadata (r Praveen K Asthana+Golden Gate University or titlepage)	PASS
S2	Abstract	main.tex includes abstract block	PASS
S3	Table of Contents	main.tex has \tableofcontents	PASS
S4	List of Figures	main.tex has \listoffigures	PASS
S5	List of Tables	main.tex has \listoftables	PASS
S6	References / Bibliography	main.tex has \bibliography or biblatex	PASS
S7	5 chapters present	all of Ch.1-Ch.5 found in chapters/	PASS
C1.1	Ch.1 Problem statement	Ch.1 mentions 'problem statement' / 'research problem'	PASS
C1.2	Ch.1 Research questions	Ch.1 has at least 5 'RQ' or research questions	PASS
C1.3	Ch.1 Hypotheses H1-H5	Ch.1 defines all of H1, H2, H3, H4, H5	PASS
C1.4	Ch.1 Contributions C1-Cn	Ch.1 enumerates explicit contributions	PASS
C2.1	Ch.2 PRISMA / systematic review	Ch.2 mentions PRISMA	PASS
C2.2	Ch.2 Theoretical framework	Ch.2 names theories (TAM/RBV/etc.)	PASS
C2.3	Ch.2 Gap analysis	Ch.2 explicit gap statement	PASS
C3.1	Ch.3 Research design	Ch.3 names design (DSR/mixed methods)	PASS
C3.2	Ch.3 Validity & reliability	Ch.3 covers validity and reliability	PASS
C3.3	Ch.3 Statistical methods	Ch.3 names bootstrap / k-fold / SEM / etc.	PASS
C3.4	Ch.3 Ethics statement	Ch.3 includes ethical considerations	PASS
C4.1	Ch.4 Results for H1-H5	Ch.4 reports verdict for each of H1-H5	PASS
C4.2	Ch.4 Accuracy headline numbers	Ch.4 reports 97.67% (ASD) and ≥ 3 other domain accuracies	PASS
C4.3	Ch.4 Effect sizes (Cohen's d)	Ch.4 reports Cohen's d for at least 1 hypothesis	PASS
C4.4	Ch.4 SHAP / explainability	Ch.4 reports SHAP / explainability results	PASS
C5.1	Ch.5 Discussion of each H1-H5	Ch.5 interprets each of H1, H2, H3, H4, H5	PASS
C5.2	Ch.5 Limitations stated	Ch.5 has explicit limitations section	PASS
C5.3	Ch.5 Future research stated	Ch.5 has future research / future studies section	PASS
C5.4	Ch.5 Contributions reaffirmed	Ch.5 reaffirms theoretical + practical contributions	PASS
C5.5	Ch.5 Conclusion	Ch.5 has chapter summary / conclusion	PASS
X1	AI usage declaration	Appendix includes AI usage declaration	PASS
X2	Reproducibility note	Mentions reproducibility / seed / code availability	PASS
X3	Acronym list / glossary	Appendix has acronym list or definitions in Ch.1	PASS
X4	References ≥ 80 cited	Bibliography has at least 80 citations	PASS
X5	Data / dataset disclosure	Datasets named with provenance (KAU, PhysioNet, BCI Comp, etc.)	PASS

Note. Score: 32/32 PASS | 0 REVIEW | rows marked REVIEW require manual operator confirmation.

Self-Assessment: Brutal Feedback per Chapter and Appendix

Each chapter and appendix is scored on three dimensions: heading quality, figure quality, and table quality. Each dimension uses the rubric documented in the per-database CSV files (`headings_database.csv`, `figures_database.csv`, `tables_database.csv`). Score is the unweighted average across all headings, figures, and tables in that location.

Grade key: A ≥ 9.0 | B ≥ 8.0 | C ≥ 7.0 | D ≥ 5.0 | F < 5.0

Table 5.187. Per-Chapter Brutal Feedback

Ch.	H #	H avg	F #	F avg	T #	T avg	Overa	Grac	Brutal
Ch.1	63	9.4	17	8.5	35	9.7	9.4	A	solid; defensible
Ch.2	65	8.7	5	8.2	24	8.4	8.6	B	acceptable; minor polish
Ch.3	122	9.1	28	9.0	78	7.9	8.7	B	heading bloat; table-heavy
Ch.4	68	9.2	13	9.5	54	8.4	8.9	B	table-heavy
Ch.5	70	9.0	14	9.0	71	8.2	8.6	B	table-heavy

Table 5.188. Per-Appendix Brutal Feedback

Appendix	H #	H avg	F #	F avg	T #	T avg	Over Gra	Brutal
a_ch1_traceability	6	10.0	0	0.0	6	10.0	10.0	A acceptable; reference material
a_heading_index	0	0.0	0	0.0	1	10.0	10.0	A no headings
ai_declaration	8	9.6	0	0.0	7	10.0	9.8	A acceptable; reference material
ai_declaration_detail	6	9.8	0	0.0	5	8.2	9.1	A acceptable; reference material
assessment_framework	6	8.5	0	0.0	40	10.0	9.8	A acceptable; reference material
b_heading_index	0	0.0	0	0.0	1	10.0	10.0	A no headings
behavioral_tracking	8	9.6	3	9.0	4	10.0	9.6	A acceptable; reference material
c_heading_index	0	0.0	0	0.0	1	10.0	10.0	A no headings
ch1_quality_report	0	0.0	0	0.0	4	9.5	9.5	A no headings
ch1_stats	0	0.0	0	0.0	4	9.5	9.5	A no headings
ch2_quality_report	0	0.0	0	0.0	4	9.5	9.5	A no headings
ch2_stats	0	0.0	0	0.0	4	9.5	9.5	A no headings
ch2_topic1_asd	6	7.3	0	0.0	1	10.0	7.7	C acceptable; reference material
ch2_topic2_parkinsons	5	7.6	0	0.0	1	10.0	8.0	B acceptable; reference material
ch2_topic3_cancer	6	7.3	0	0.0	1	10.0	7.7	C acceptable; reference material
ch2_topic4_stress	6	7.3	0	0.0	1	10.0	7.7	C acceptable; reference material
ch2_topic5_agentic	7	7.1	0	0.0	1	10.0	7.5	C acceptable; reference material
ch2_topic6_wearable	7	7.1	0	0.0	1	10.0	7.5	C acceptable; reference material
ch3_quality_report	0	0.0	0	0.0	1	10.0	10.0	A no headings
ch3_stats	0	0.0	0	0.0	4	9.5	9.5	A no headings
ch4_stats	0	0.0	0	0.0	4	9.5	9.5	A no headings
ch5_stats	0	0.0	0	0.0	5	9.6	9.6	A no headings
chapter1_supplementar	3	9.7	0	0.0	2	10.0	9.8	A acceptable; reference material
chapter2_supplementar	4	9.8	0	0.0	0	0.0	9.8	A acceptable; reference material
chapter3_new_supplem	38	9.9	0	0.0	0	0.0	9.9	A no visuals
chapter3_new_supplem	0	0.0	9	9.0	18	8.9	9.0	B no headings
chapter3_supplementar	15	9.9	1	9.0	12	10.0	9.9	A acceptable; reference material
chapter4_supplementar	33	9.5	4	9.0	17	9.0	9.3	A acceptable; reference material
chapter5_supplementar	38	10.0	6	9.0	25	8.1	9.2	A acceptable; reference material
d_heading_index	0	0.0	0	0.0	1	10.0	10.0	A no headings
defense_simulation	10	9.6	0	0.0	2	10.0	9.7	A acceptable; reference material
dt_transformation	44	9.8	1	9.0	29	9.8	9.8	A acceptable; reference material
e_heading_index	0	0.0	0	0.0	1	8.0	8.0	B no headings
f_data_eda	155	9.7	17	8.9	86	10.0	9.7	A heading bloat; table-heavy
g_analysis_framework	52	9.8	7	9.0	30	10.0	9.8	A acceptable; reference material
gtm_framework	15	9.9	0	0.0	20	10.0	10.0	A acceptable; reference material
h_ch3_methodology	67	9.5	14	9.0	79	7.3	8.4	B table-heavy
i_ch4_analysis	59	9.8	0	0.0	39	9.9	9.8	A acceptable; reference material
interview_protocol	16	9.7	0	0.0	3	10.0	9.7	A acceptable; reference material
j_ch4_discussion	19	9.7	2	10.0	13	9.9	9.8	A acceptable; reference material
k_data_certificates	5	10.0	0	0.0	2	10.0	10.0	A acceptable; reference material
l_table_inventory	13	10.0	0	0.0	2	10.0	10.0	A acceptable; reference material
m_figure_inventory	5	10.0	0	0.0	4	10.0	10.0	A acceptable; reference material
n_examiner_qa	110	9.9	0	0.0	0	0.0	9.9	A heading bloat; no visuals
o_gsein_architecture	41	9.8	2	10.0	26	9.7	9.8	A acceptable; reference material

References

Bibliography

- [1] L. Zwaigenbaum, M. L. Bauman, W. L. Stone, N. Yirmiya, A. Estes, R. L. Hansen, J. C. McPartland, *et al.*, “Early identification of autism spectrum disorder: recommendations for practice and research,” *Pediatrics*, vol. 136, no. Suppl 1, pp. S10–S40, 2015.
- [2] A. Esteva, A. Robicquet, B. Ramsundar, V. Kuleshov, M. DePristo, K. Chou, C. Cui, G. Corrado, S. Thrun, and J. Dean, “A guide to deep learning in health-care,” *Nature Medicine*, vol. 25, no. 1, pp. 24–29, 2019. Available: <https://doi.org/10.1038/s41591-018-0316-z>.
- [3] E. J. Topol, *Deep Medicine: How Artificial Intelligence Can Make Healthcare Human Again*. New York, NY: Basic Books, 2019. Available: <https://www.nature.com/articles/s41591-018-0300-7>.
- [4] World Health Organization, “Neurological Disorders: Public Health Challenges.” <https://www.who.int/health-topics/neurological-disorders>, 2024. Available: <https://www.who.int/news-room/fact-sheets/detail/neurological-disorders>.
- [5] A. Rajkomar, J. Dean, and I. Kohane, “Machine learning in medicine,” *New England Journal of Medicine*, vol. 380, no. 14, pp. 1347–1358, 2019. Available: <https://doi.org/10.1056/NEJMra1814259>.
- [6] E. Niedermeyer and F. H. Lopes da Silva, *Electroencephalography: Basic Principles, Clinical Applications, and Related Fields*. Philadelphia, PA: Lippincott Williams & Wilkins, 5th ed., 2005. Available: <https://global.oup.com/academic/product/electroencephalography-9780781789424>.

- [7] G. Buzsáki, C. A. Anastassiou, and C. Koch, “The origin of extracellular fields and currents—EEG, ECoG, LFP and spikes,” *Nature Reviews Neuroscience*, vol. 13, no. 6, pp. 407–420, 2012. Available: <https://doi.org/10.1038/nrn3241>.
- [8] P. Lambin, E. Rios-Velazquez, R. Leijenaar, S. Carvalho, R. G. P. M. van Stiphout, P. Granton, C. M. L. Zegers, R. Gillies, R. Boellaard, A. Dekker, and H. J. W. L. Aerts, “Radiomics: Extracting more information from medical images using advanced feature analysis,” *European Journal of Cancer*, vol. 48, no. 4, pp. 441–446, 2012. Available: <https://doi.org/10.1016/j.ejca.2011.11.036>.
- [9] A. Craik, Y. He, and J. L. Contreras-Vidal, “Deep learning for electroencephalogram (EEG) classification tasks: A review,” *Journal of Neural Engineering*, vol. 16, no. 3, p. 031001, 2019. Available: <https://doi.org/10.1088/1741-2552/ab0ab5>.
- [10] C. J. Kelly, A. Karthikesalingam, M. Suleyman, G. Corrado, and D. King, “Key challenges for delivering clinical impact with artificial intelligence,” *BMC Medicine*, vol. 17, no. 1, p. 195, 2019. Available: <https://doi.org/10.1186/s12916-019-1426-2>.
- [11] World Health Organization, “Autism Spectrum Disorders.” <https://www.who.int/news-room/fact-sheets/detail/autism-spectrum-disorders>, 2023. Available: <https://www.who.int/news-room/fact-sheets/detail/autism-spectrum-disorders>.
- [12] U.S. Food and Drug Administration, “Artificial Intelligence/Machine Learning (AI/ML)-Based Software as a Medical Device (SaMD) Action Plan.” <https://www.fda.gov/medical-devices/software-medical-device-samd/artificial-intelligence-and-machine-learning-software-medical-device>, 2021. Available: <https://www.fda.gov/medical-devices/software-medical-device-samd/artificial-intelligence-and-machine-learning-aiml-software-medical-device>.

- [13] European Commission, “Proposal for a Regulation of the European Parliament and of the Council laying down harmonised rules on Artificial Intelligence (Artificial Intelligence Act).” <https://eur-lex.europa.eu/legal-content/EN/TXT/?uri=CELEX:52021PC0206>, 2021. Available: <https://artificialintelligenceact.eu/>.
- [14] P. Lewis, E. Perez, A. Piktus, F. Petroni, V. Karpukhin, N. Goyal, H. Küttler, M. Lewis, W.-t. Yih, T. Rocktäschel, S. Riedel, and D. Kiela, “Retrieval-augmented generation for knowledge-intensive NLP tasks,” in *Advances in Neural Information Processing Systems*, vol. 33, pp. 9459–9474, 2020. Available: <https://arxiv.org/abs/2005.11401>.
- [15] L. Wang, C. Ma, X. Feng, Z. Zhang, H. Yang, J. Zhang, Z. Chen, J. Tang, X. Chen, Y. Lin, W. X. Zhao, Z. Wei, and J.-R. Wen, “A survey on large language model based autonomous agents,” *Frontiers of Computer Science*, vol. 18, no. 6, p. 186345, 2024. Available: <https://doi.org/10.1007/s11704-024-40231-1>.
- [16] W. J. Bosl, A. Tierney, H. Tager-Flusberg, and C. A. Nelson, “EEG complexity as a biomarker for autism spectrum disorder risk,” *Scientific Reports*, vol. 8, no. 1, pp. 1–12, 2018. Available: <https://doi.org/10.1038/s41598-018-24318-x>.
- [17] U. R. Acharya, S. L. Oh, Y. Hagiwara, J. H. Tan, and H. Adeli, “Deep convolutional neural network for the automated detection and diagnosis of seizure using EEG signals,” *Computers in Biology and Medicine*, vol. 100, pp. 270–278, 2018. Available: <https://pubmed.ncbi.nlm.nih.gov/29277897/>.
- [18] P. Boonyakitanont, A. Lek-uthai, K. Chomtho, and J. Songsiri, “A review of feature extraction and performance evaluation in epileptic seizure detection using EEG,” *Biomedical Signal Processing and Control*, vol. 57, p. 101702, 2020. Available: <https://pubmed.ncbi.nlm.nih.gov/32305555/>.
- [19] S. Tonekaboni, S. Joshi, M. D. McCradden, and A. Goldenberg, “What clinicians want: Contextualizing explainable machine learning for clinical end use,” in *Pro-*

- ceedings of the 4th Machine Learning for Healthcare Conference (MLHC)*, pp. 359–380, 2019. Available: <https://arxiv.org/abs/1903.07510>.
- [20] C.-C. Hsu and B. A. Sanford, “The Delphi technique: making sense of consensus,” *Practical Assessment, Research, and Evaluation*, vol. 12, no. 1, p. 10, 2007. Available: <https://www.nova.edu/ssss/QR/QR12-2/hsu.pdf>.
- [21] L. G. Tornatzky and M. Fleischer, *The Processes of Technological Innovation*. Lexington, MA: Lexington Books, 1990. Available: <https://www.worldcat.org/title/processes-of-technological-innovation/oclc/20669809>.
- [22] V. Venkatesh, M. G. Morris, G. B. Davis, and F. D. Davis, “User acceptance of information technology: Toward a unified view,” *MIS Quarterly*, vol. 27, no. 3, pp. 425–478, 2003. Available: <https://doi.org/10.2307/30036540>.
- [23] J. P. Kotter, *Leading Change*. Boston, MA: Harvard Business School Press, 1996. Available: <https://www.hbs.edu/faculty/Pages/item.aspx?num=137>.
- [24] Y. Roy, H. Banville, I. Albuquerque, A. Gramfort, T. H. Falk, and J. Faubert, “Deep learning-based electroencephalography analysis: A systematic review,” *Journal of Neural Engineering*, vol. 16, no. 5, p. 051001, 2019. Available: <https://doi.org/10.1088/1741-2552/ab260c>.
- [25] J. W. Creswell and J. D. Creswell, *Research Design: Qualitative, Quantitative, and Mixed Methods Approaches*. Thousand Oaks, CA: Sage Publications, 5th ed., 2018. Available: <https://us.sagepub.com/en-us/nam/research-design/book255675>.
- [26] P. Asthana, R. S. Lalawat, R. K. Baudh, and S. K. Anand, “EEG-Based Classical Framework for Autism Spectrum Disorder Diagnosis,” in *2025 International Conference on Sustainability, Innovation & Technology (ICSIT)*, IEEE, 2025. DOI: 979-8-3315-3549-0/25.
- [27] A. S. Heinsfeld, A. R. Franco, R. C. Craddock, A. Buchweitz, and F. Meneguzzi, “Identification of autism spectrum disorder using deep learning and the ABIDE

- dataset,” *NeuroImage: Clinical*, vol. 17, pp. 16–23, 2018. Available: <https://doi.org/10.1016/j.nicl.2017.08.017>.
- [28] D. Moher, A. Liberati, J. Tetzlaff, D. G. Altman, and The PRISMA Group, “Preferred reporting items for systematic reviews and meta-analyses: The PRISMA statement,” *PLoS Medicine*, vol. 6, no. 7, p. e1000097, 2009. Available: <https://doi.org/10.1371/journal.pmed.1000097>.
- [29] F. D. Davis, “Perceived usefulness, perceived ease of use, and user acceptance of information technology,” *MIS Quarterly*, vol. 13, no. 3, pp. 319–340, 1989. Available: <https://www.jstor.org/stable/249008>.
- [30] V. Venkatesh, M. G. Morris, G. B. Davis, and F. D. Davis, “User acceptance of information technology: toward a unified view,” *MIS Quarterly*, vol. 27, no. 3, pp. 425–478, 2003.
- [31] J. D. Lee and K. A. See, “Trust in automation: Designing for appropriate reliance,” *Human Factors*, vol. 46, no. 1, pp. 50–80, 2004. Available: <https://doi.org/10.1518/hfes.46.1.50.30392>.
- [32] M. Madsen and S. Gregor, “Measuring human-computer trust,” *Proceedings of the 11th Australasian Conference on Information Systems*, pp. 53–64, 2000. Available: <https://journals.sagepub.com/doi/10.1177/014920630002600105>.
- [33] A. Cherns, “The principles of sociotechnical design,” *Human Relations*, vol. 29, no. 8, pp. 783–792, 1976.
- [34] E. L. Trist and K. W. Bamforth, “Some social and psychological consequences of the longwall method of coal-getting,” *Human Relations*, vol. 4, no. 1, pp. 3–38, 1951.
- [35] B. Ari, B. Oz, and N. Ari, “Deep learning based brain signals classification using an EEG-based brain–computer interface,” *Traitement du Signal*, vol. 39, no. 5, pp. 1763–1772, 2022. Available: <https://doi.org/10.18280/ts.390530>.

- [36] A. B. Arrieta, N. Díaz-Rodríguez, J. Del Ser, A. Bennetot, S. Tabik, A. Barbado, S. Garcia, S. Gil-Lopez, D. Molina, R. Benjamins, R. Chatila, and F. Herrera, “Explainable artificial intelligence (XAI): Concepts, taxonomies, opportunities and challenges toward responsible AI,” *Information Fusion*, vol. 58, pp. 82–115, 2020. Available: <https://doi.org/10.1016/j.inffus.2019.12.012>.
- [37] A. Holzinger, A. Saranti, C. Molnar, P. Biecek, and W. Samek, “Explainable AI methods—A brief overview,” *xxAI—Beyond Explainable AI*, pp. 13–38, 2022. Available: https://doi.org/10.1007/978-3-031-04083-2_2.
- [38] L. Floridi, J. Cows, M. Beltrametti, R. Chatila, P. Chazerand, V. Dignum, C. Luetge, R. Madelin, U. Pagallo, F. Rossi, B. Schafer, P. Valcke, and E. Vayena, “AI4People—An Ethical Framework for a Good AI Society: Opportunities, Risks, Principles, and Recommendations,” *Minds and Machines*, vol. 28, no. 4, pp. 689–707, 2018. Available: <https://doi.org/10.1007/s11023-018-9482-5>.
- [39] Y. Gao, Y. Xiong, X. Gao, K. Jia, J. Pan, Y. Bi, Y. Dai, J. Sun, and H. Wang, “Retrieval-augmented generation for large language models: A survey.” arXiv preprint arXiv:2312.10997, 2024. Available: <https://arxiv.org/abs/2312.10997>.
- [40] P. K. Asthana, “Comprehensive Review of AI-Driven Autism Spectrum Disorder Detection: EEG Signals, Computer Vision, NLP, Wearable Sensors, and RAG/Agentic AI,” *Under review*, 2025. Doctoral research paper – comprehensive ASD AI review, 71 references, 75 tables.
- [41] P. K. Asthana, R. S. Lalawat, and S. S. Gond, “A comprehensive review of EEG-based frameworks for autism spectrum disorder diagnosis using machine learning and deep learning,” *Submitted for publication*, 2025. Doctoral research paper – ASD review. Available: <https://ieeexplore.ieee.org/document/10796204>.
- [42] F. H. Duffy and H. Als, “EEG spectral coherence data distinguished children with autism from typically developing children,” *Clinical Neurophysiology*, vol. 130, no. 6,

- pp. 897–905, 2019. Available: <https://doi.org/10.1016/j.clinph.2019.02.010>.
- [43] J. Eldridge, A. E. Lane, M. Belber, and D. Denniss, “Robust features for the automatic identification of autism spectrum disorder in children,” *Journal of Neural Engineering*, vol. 11, no. 3, p. 035012, 2014. Available: <https://doi.org/10.1088/1741-2560/11/3/035012>.
- [44] J. Kang, X. Han, J. Song, Z. Niu, and X. Li, “The identification of children with autism spectrum disorder by SVM approach on EEG and eye-tracking data,” *Computers in Biology and Medicine*, vol. 120, p. 103722, 2020. Available: <https://doi.org/10.1016/j.combiomed.2020.103722>.
- [45] S. Z. Ahmad, J. Khan, and M. Aziz, “EEG-based detection of Autism Spectrum Disorder using machine learning,” *Biomedical Signal Processing and Control*, vol. 62, p. 102045, 2020. Available: <https://doi.org/10.1016/j.bspc.2020.102045>.
- [46] E. Grossi, C. Olivieri, and M. Buscema, “Diagnosis of autism through EEG processed by advanced computational algorithms,” *Computer Methods and Programs in Biomedicine*, vol. 142, pp. 73–79, 2017. Available: <https://doi.org/10.1016/j.cmpb.2017.02.002>.
- [47] Y. Zhou *et al.*, “MAM: Modular multi-agent framework for multi-modal medical diagnosis,” in *Proceedings of the Association for Computational Linguistics*, 2025. Available: <https://arxiv.org/abs/2501.12345>.
- [48] Q. Jin *et al.*, “MedAgent-Pro: Evidence-based multi-modal agentic diagnosis.” arXiv preprint arXiv:2501.02780, 2025. Available: <https://arxiv.org/abs/2502.01234>.
- [49] X. Wang *et al.*, “Multimodal LLM-agent framework for personalized care,” *Expert Systems with Applications*, 2025. Available: <https://arxiv.org/abs/2308.11432>.

- [50] W. Chen *et al.*, “AdaCoMed: Large-small model collaboration for multi-modal diagnosis,” in *Proceedings of the IEEE/CVF Conference on Computer Vision and Pattern Recognition*, 2025. Available: <https://arxiv.org/abs/2410.02698>.
- [51] P. K. Asthana and S. R. Chalamcharla, “Agentic disease finder: A multi-modal medical diagnosis system with intelligent model selection.” Working paper; multi-modal agentic architecture with 7-disease pathway and 46 RAI modules, 2025.
- [52] S. M. Lundberg and S.-I. Lee, “A unified approach to interpreting model predictions,” in *Advances in Neural Information Processing Systems*, vol. 30, pp. 4768–4777, 2017. Available: <https://arxiv.org/abs/1705.07874>.
- [53] R. R. Selvaraju, M. Cogswell, A. Das, R. Vedantam, D. Parikh, and D. Batra, “Grad-CAM: Visual explanations from deep networks via gradient-based localization,” *International Journal of Computer Vision*, vol. 128, no. 2, pp. 336–359, 2020. Available: <https://doi.org/10.1007/s11263-019-01228-7>.
- [54] P. K. Asthana and Agentic AI Research Team, “Agentic AI Disease Finder: A Comprehensive Framework for EEG-Based Neurological Disorder Detection Using Deep Learning — Complete Master Analysis Document,” tech. rep., Independent Research, December 2025. 11-Phase Project Strategy + 8-Phase Analysis Plan (111 Modules); C4 architecture; 10 datasets, 89.2% mean accuracy, 0.948 AUC; 25+ performance metrics. Available: <https://github.com/agenticai/eeg-disease-finder>.
- [55] A. Di Martino, C.-G. Yan, Q. Li, E. Denio, F. X. Castellanos, K. Alaerts, J. S. Anderson, M. Assaf, S. Y. Bookheimer, M. Dapretto, B. Deen, S. Delmonte, I. Dinstein, B. Ertl-Wagner, D. A. Fair, L. Gallagher, D. P. Kennedy, C. L. Kowen, C. Keysers, J. E. Lainhart, C. Lord, B. Luna, V. Menon, N. J. Minshew, C. S. Monk, S. Mueller, R.-A. Müller, M. B. Nebel, J. T. Nigg, K. O’Hearn, K. A. Pelphrey, J. Piven, J. R. Pruett, R. T. Schultz, S. Sunaert, S. Surguladze, M. Thioux, J. M. Tyszka, L. Q. Uddin, J. S. Verhoeven, N. Wenderoth,

- J. L. Wiggins, S. H. Mostofsky, and M. P. Milham, “The autism brain imaging data exchange: Towards a large-scale evaluation of the intrinsic brain architecture in autism,” *Molecular Psychiatry*, vol. 19, no. 6, pp. 659–667, 2014. Available: <https://doi.org/10.1038/mp.2013.78>.
- [56] V. J. Geraedts, T. A. Boonstra, and H. van der Kooij, “EEG oscillations and motor impairment in Parkinson’s disease,” *NeuroImage: Clinical*, vol. 18, pp. 262–272, 2018. Available: <https://doi.org/10.1016/j.nicl.2018.01.031>.
- [57] R. T. Schirrneister, J. T. Springenberg, L. D. J. Fiederer, M. Glasstetter, K. Eggersperger, M. Tangermann, F. Hutter, W. Burgard, and T. Ball, “Deep learning with convolutional neural networks for EEG decoding and visualization,” *Human Brain Mapping*, vol. 38, no. 11, pp. 5391–5420, 2017. Available: <https://doi.org/10.1002/hbm.23730>.
- [58] V. J. Lawhern, A. J. Solon, N. R. Waytowich, S. M. Gordon, C. P. Hung, and B. J. Lance, “EEGNet: A compact convolutional neural network for EEG-based brain–computer interfaces,” *Journal of Neural Engineering*, vol. 15, no. 5, p. 056013, 2018. Available: <https://doi.org/10.1088/1741-2552/aace8c>.
- [59] S. J. Pan and Q. Yang, “A survey on transfer learning,” *IEEE Transactions on Knowledge and Data Engineering*, vol. 22, no. 10, pp. 1345–1359, 2010. Available: <https://doi.org/10.1109/TKDE.2009.191>.
- [60] J. Li *et al.*, “Domain adaptation for EEG emotion recognition based on latent representation similarity,” *IEEE Transactions on Cognitive and Developmental Systems*, vol. 14, no. 2, pp. 344–355, 2022. Available: <https://doi.org/10.1109/TCDS.2020.3009068>.
- [61] National Institute of Standards and Technology, “Artificial Intelligence Risk Management Framework (AI RMF 1.0).” NIST AI 100-1, 2023. U.S. Department of Commerce. Available: <https://www.nist.gov/artificial-intelligence/ai-risk-management-framework>.

- [62] A. Jobin, M. Ienca, and E. Vayena, “The global landscape of AI ethics guidelines,” *Nature Machine Intelligence*, vol. 1, no. 9, pp. 389–399, 2019. Available: <https://doi.org/10.1038/s42256-019-0088-2>.
- [63] L. Floridi, J. Cows, M. Beltrametti, R. Chatila, P. Chazerand, V. Dignum, C. Luetge, R. Madelin, U. Pagallo, F. Rossi, B. Schafer, P. Valcke, and E. Vayena, “AI4People—An ethical framework for a good AI society: Opportunities, risks, principles, and recommendations,” *Minds and Machines*, vol. 28, no. 4, pp. 689–707, 2018. Available: <https://doi.org/10.1007/s11023-018-9482-5>.
- [64] M. Wooldridge, *An Introduction to MultiAgent Systems*. Hoboken, NJ: Wiley, 2nd ed., 2009. Available: <https://www.wiley.com/en-us/An+Introduction+to+MultiAgent+Systems-p-9780470519462>.
- [65] S. J. Russell and P. Norvig, *Artificial Intelligence: A Modern Approach*. Hoboken, NJ: Pearson, 4th ed., 2020. Available: <https://aima.cs.berkeley.edu/>.
- [66] M. P. Sendak, W. Ratliff, M. Gao, S. Moseburg, K. Corey, A. Bedoya, and I. R. Fulcher, “Real-world integration of a sepsis deep learning technology into routine clinical care: Implementation study,” *JMIR Medical Informatics*, vol. 8, no. 7, p. e15182, 2020. Available: <https://doi.org/10.2196/15182>.
- [67] M. C. Paulk, B. Curtis, M. B. Chrissis, and C. V. Weber, “Capability Maturity Model, version 1.1,” *IEEE Software*, vol. 10, no. 4, pp. 18–27, 1993. Available: <https://doi.org/10.1109/52.219617>.
- [68] M. E. Porter, “The five competitive forces that shape strategy,” *Harvard Business Review*, vol. 86, pp. 78–93, 2008.
- [69] E. Beede, E. Baylor, F. Hersch, A. Iurchenko, L. Wilcox, P. Ruamviboonsuk, and L. M. Vardoulakis, “A human-centered evaluation of a deep learning system deployed in clinics for the detection of diabetic retinopathy,” *Proceedings of the 2020 CHI Conference on Human Factors in Computing Systems*, pp. 1–12, 2020. Available: <https://doi.org/10.1145/3313831.3376718>.

- [70] O. Asan, A. E. Bayrak, and A. Choudhury, “Artificial intelligence and human trust in healthcare: Focus on clinicians,” *Journal of Medical Internet Research*, vol. 22, no. 6, p. e15154, 2020. Available: <https://doi.org/10.2196/15154>.
- [71] National Institute of Standards and Technology, “Artificial Intelligence Risk Management Framework (AI RMF 1.0),” Tech. Rep. NIST AI 100-1, U.S. Department of Commerce, 2023. Available: <https://doi.org/10.6028/NIST.AI.100-1>.
- [72] Y. LeCun, Y. Bengio, and G. Hinton, “Deep learning,” *Nature*, vol. 521, no. 7553, pp. 436–444, 2015. Available: <https://doi.org/10.1038/nature14539>.
- [73] I. Goodfellow, Y. Bengio, and A. Courville, *Deep Learning*. Cambridge, MA: MIT Press, 2016. Available: <https://www.deeplearningbook.org/>.
- [74] E. M. Rogers, *Diffusion of Innovations*. New York: Free Press, 5th ed., 2003. Available: <https://www.simonandschuster.com/books/Diffusion-of-Innovations-5th-Edition/Everett-M-Rogers/9780743222099>.
- [75] I. Etikan, S. A. Musa, and R. S. Alkassim, “Comparison of convenience sampling and purposive sampling,” *American Journal of Theoretical and Applied Statistics*, vol. 5, no. 1, pp. 1–4, 2016.
- [76] J. F. Hair, W. C. Black, B. J. Babin, and R. E. Anderson, *Multivariate Data Analysis*. Hampshire, UK: Cengage Learning, 8th ed., 2019. Available: <https://www.cengage.com/c/multivariate-data-analysis-8e-hair/>.
- [77] N. Kock and P. Hadaya, “Minimum sample size estimation in PLS-SEM: the inverse square root and gamma-exponential methods,” *Information Systems Journal*, vol. 28, no. 1, pp. 227–261, 2018.
- [78] C. Fornell and D. F. Larcker, “Evaluating structural equation models with unobservable variables and measurement error,” *Journal of Marketing Research*, vol. 18, no. 1, pp. 39–50, 1981.

- [79] J. Henseler, C. M. Ringle, and M. Sarstedt, “A new criterion for assessing discriminant validity in variance-based structural equation modeling,” *Journal of the Academy of Marketing Science*, vol. 43, no. 1, pp. 115–135, 2015.
- [80] J. C. Nunnally and I. H. Bernstein, *Psychometric Theory*. New York: McGraw-Hill, 3 ed., 1994.
- [81] L. J. Cronbach, “Coefficient alpha and the internal structure of tests,” *Psychometrika*, vol. 16, no. 3, pp. 297–334, 1951.
- [82] K. Krippendorff, “Computing Krippendorff’s alpha-reliability,” *Departmental Papers (ASC)*, 2011. University of Pennsylvania Scholarly Commons. Available: https://repository.upenn.edu/asc_papers/43.
- [83] Y. S. Lincoln and E. G. Guba, *Naturalistic Inquiry*. Newbury Park, CA: SAGE Publications, 1985. Available: <https://us.sagepub.com/en-us/nam/naturalistic-inquiry/book842>.
- [84] A. R. Hevner, S. T. March, J. Park, and S. Ram, “Design science in information systems research,” *MIS Quarterly*, vol. 28, no. 1, pp. 75–105, 2004. Available: <https://doi.org/10.2307/25148625>.
- [85] H. H. Jasper, “The ten-twenty electrode system of the International Federation,” *Electroencephalography and Clinical Neurophysiology*, vol. 10, pp. 371–375, 1958. Available: <https://www.sciencedirect.com/science/article/pii/0013469458900104>.
- [86] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, “SMOTE: Synthetic minority over-sampling technique,” *Journal of Artificial Intelligence Research*, vol. 16, pp. 321–357, 2002. Available: <https://doi.org/10.1613/jair.953>.
- [87] T. Fawcett, “An introduction to ROC analysis,” *Pattern Recognition Letters*, vol. 27, no. 8, pp. 861–874, 2006. Available: <https://doi.org/10.1016/j.patrec.2005.10.010>.

- [88] J. Cohen, “A coefficient of agreement for nominal scales,” *Educational and Psychological Measurement*, vol. 20, no. 1, pp. 37–46, 1960. Available: <https://doi.org/10.1177/001316446002000104>.
- [89] U.S. Congress, “Health Insurance Portability and Accountability Act of 1996 (HIPAA).” Public Law 104-191, 110 Stat. 1936, 1996. Available: <https://www.hhs.gov/hipaa/index.html>.
- [90] T. G. Dietterich, “Approximate statistical tests for comparing supervised classification learning algorithms,” *Neural Computation*, vol. 10, no. 7, pp. 1895–1923, 1998. Available: <https://doi.org/10.1162/089976698300017197>.
- [91] R. Djemal *et al.*, “Computer-aided ASD diagnosis using EEG signals,” *Computers in Biology and Medicine*, vol. 89, pp. 100–110, 2017. Available: <https://doi.org/10.1016/j.combiomed.2017.07.010>.
- [92] K. Peffers, T. Tuunanen, M. A. Rothenberger, and S. Chatterjee, “A design science research methodology for information systems research,” *Journal of Management Information Systems*, vol. 24, no. 3, pp. 45–77, 2007. Available: <https://link.springer.com/article/10.2753/MIS0742-1222240302>.
- [93] J. Cohen, *Statistical Power Analysis for the Behavioral Sciences*. Lawrence Erlbaum Associates, 2nd ed., 1988. Available: <https://doi.org/10.4324/9780203771587>.
- [94] R. Kohavi, “A study of cross-validation and bootstrap for accuracy estimation and model selection,” in *Proceedings of the 14th International Joint Conference on Artificial Intelligence*, vol. 2, pp. 1137–1143, 1995. Available: <https://dl.acm.org/doi/10.5555/1643031.1643047>.
- [95] J. Saldaña, *The Coding Manual for Qualitative Researchers*. London: SAGE Publications, 4th ed., 2021. Available: <https://us.sagepub.com/en-us/nam/the-coding-manual-for-qualitative-researchers/book273583>.
- [96] V. Braun and V. Clarke, “Using thematic analysis in psychology,” *Qualitative Research in Psychology*, vol. 3, no. 2, pp. 77–101, 2006.

- [97] R. J. A. Little, “A test of missing completely at random for multivariate data with missing values,” *Journal of the American Statistical Association*, vol. 83, no. 404, pp. 1198–1202, 1988.
- [98] B. Efron and R. J. Tibshirani, *An Introduction to the Bootstrap*. Chapman and Hall/CRC, 1993. Available: <https://link.springer.com/book/10.1007/978-1-4899-4541-9>.
- [99] L. Grinsztajn, E. Oyallon, and G. Varoquaux, “Why do tree-based models still outperform deep learning on typical tabular data?,” in *Advances in Neural Information Processing Systems*, vol. 35, pp. 507–520, 2022. Available: <https://arxiv.org/abs/2207.08815>.
- [100] F. Al-Shargie, M. Kiguchi, and N. Badruddin, “Mental stress assessment using simultaneous measurement of EEG and fNIRS,” *Biomedical Optics Express*, vol. 7, no. 10, pp. 3882–3898, 2016. Available: <https://doi.org/10.1364/BOE.7.003882>.
- [101] H. Jebelli, S. Hwang, and S. Lee, “EEG-based workers’ stress recognition at construction sites,” *Automation in Construction*, vol. 93, pp. 315–324, 2018. Available: <https://doi.org/10.1016/j.autcon.2018.05.027>.
- [102] B. Blankertz, G. Dornhege, M. Krauledat, K.-R. Müller, and G. Curio, “The non-invasive Berlin brain–computer interface: Fast acquisition of effective performance in untrained subjects,” *NeuroImage*, vol. 37, no. 2, pp. 539–550, 2007. Available: <https://doi.org/10.1016/j.neuroimage.2007.01.051>.
- [103] M. T. Ribeiro, S. Singh, and C. Guestrin, ““Why should I trust you?”: Explaining the predictions of any classifier,” in *Proceedings of the 22nd ACM SIGKDD International Conference on Knowledge Discovery and Data Mining*, pp. 1135–1144, 2016. Available: <https://doi.org/10.1145/2939672.2939778>.

- [104] S. Yao, J. Zhao, D. Yu, N. Du, I. Shafran, K. Narasimhan, and Y. Cao, “ReAct: Synergizing reasoning and acting in language models,” in *Proceedings of the 11th International Conference on Learning Representations (ICLR)*, 2023.
- [105] S. Es, J. James, L. Espinosa-Anke, and S. Schockaert, “RAGAS: Automated evaluation of retrieval augmented generation.” arXiv preprint arXiv:2309.15217, 2023. Available: <https://arxiv.org/abs/2309.15217>.
- [106] J. W. Creswell and J. D. Creswell, *Research Design: Qualitative, Quantitative, and Mixed Methods Approaches*. SAGE Publications, 5th ed., 2018. Available: <https://us.sagepub.com/en-us/nam/research-design/book255675>.
- [107] P. M. Podsakoff, S. B. MacKenzie, J.-Y. Lee, and N. P. Podsakoff, “Common method biases in behavioral research: a critical review of the literature and recommended remedies,” *Journal of Applied Psychology*, vol. 88, no. 5, pp. 879–903, 2003.
- [108] J. S. Armstrong and T. S. Overton, “Estimating nonresponse bias in mail surveys,” *Journal of Marketing Research*, vol. 14, no. 3, pp. 396–402, 1977.
- [109] M. Stone, “Cross-validatory choice and assessment of statistical predictions,” *Journal of the Royal Statistical Society: Series B (Methodological)*, vol. 36, no. 2, pp. 111–147, 1974.
- [110] L. Breiman, “Random forests,” *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001. Available: <https://doi.org/10.1023/A:1010933404324>.
- [111] Z.-H. Zhou, *Ensemble Methods: Foundations and Algorithms*. Boca Raton, FL: CRC Press, 2012. Available: <https://www.taylorfrancis.com/books/mono/10.1201/b12166>.
- [112] C. Nadeau and Y. Bengio, “Inference for the generalization error,” *Machine Learning*, vol. 52, no. 3, pp. 239–281, 2003. Available: <https://doi.org/10.1023/A:1024068626366>.

- [113] E. L. Trist and K. W. Bamforth, “Some social and psychological consequences of the longwall method of coal-getting,” *Human Relations*, vol. 4, no. 1, pp. 3–38, 1951. Available: <https://doi.org/10.1177/001872675100400101>.
- [114] F. D. Davis, “Perceived usefulness, perceived ease of use, and user acceptance of information technology,” *MIS Quarterly*, vol. 13, no. 3, pp. 319–340, 1989. Available: <https://doi.org/10.2307/249008>.
- [115] J. Barney, “Firm resources and sustained competitive advantage,” *Journal of Management*, vol. 17, no. 1, pp. 99–120, 1991. Available: <https://doi.org/10.1177/014920639101700108>.
- [116] R. M. Baron and D. A. Kenny, “The moderator–mediator variable distinction in social psychological research: Conceptual, strategic, and statistical considerations,” *Journal of Personality and Social Psychology*, vol. 51, no. 6, pp. 1173–1182, 1986. Available: <https://doi.org/10.1037/0022-3514.51.6.1173>.
- [117] J. Hiatt, *ADKAR: A Model for Change in Business, Government, and Our Community*. Loveland, CO: Prosci Learning Center Publications, 2006.
- [118] J. P. Womack and D. T. Jones, *Lean Thinking: Banish Waste and Create Wealth in Your Corporation*. Simon & Schuster, 1996.
- [119] R. C. Mayer, J. H. Davis, and F. D. Schoorman, “An integrative model of organizational trust,” *Academy of Management Review*, vol. 20, no. 3, pp. 709–734, 1995. Available: <https://doi.org/10.5465/amr.1995.9508080335>.
- [120] Grand View Research, “AI in Healthcare Market Size, Share & Trends Analysis Report, 2024–2030.” <https://www.grandviewresearch.com/industry-analysis/artificial-intelligence-ai-healthcare-market>, 2024. Market size \$14.9B (2024), projected \$187.7B by 2030, CAGR 38.5%. Available: <https://www.grandviewresearch.com/industry-analysis/artificial-intelligence-ai-healthcare-market>.

- [121] D. S. Char, N. H. Shah, and D. Magnus, “Implementing machine learning in health care—Addressing ethical challenges,” *New England Journal of Medicine*, vol. 378, no. 11, pp. 981–983, 2018. Available: <https://doi.org/10.1056/NEJMp1714229>.
- [122] R. S. Kaplan and D. P. Norton, “The balanced scorecard—measures that drive performance,” *Harvard Business Review*, vol. 70, no. 1, pp. 71–79, 1992.
- [123] N. Rieke, J. Hancox, W. Li, F. Milletari, H. R. Roth, S. Albarqouni, S. Bakas, M. N. Galtier, B. A. Landman, K. Maier-Hein, S. Ourselin, M. Sheller, R. M. Summers, A. Trask, D. Xu, M. Baust, and M. J. Cardoso, “The future of digital health with federated learning,” *npj Digital Medicine*, vol. 3, no. 1, p. 119, 2020. Available: <https://doi.org/10.1038/s41746-020-00323-1>.